

# Chapter 7

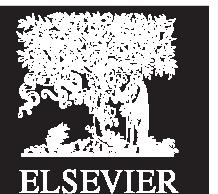
***Digital Design and Computer Architecture: ARM® Edition***

Sarah L. Harris and David Money Harris



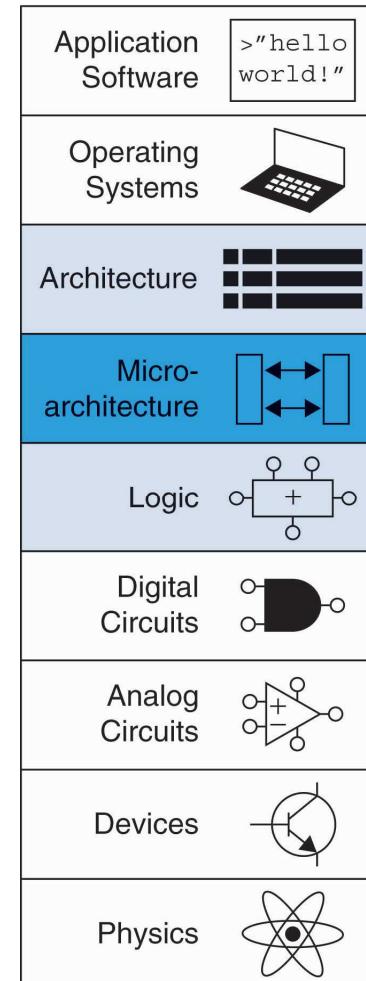
Digital Design and Computer Architecture: ARM® Edition © 2015

Chapter 7 <1>



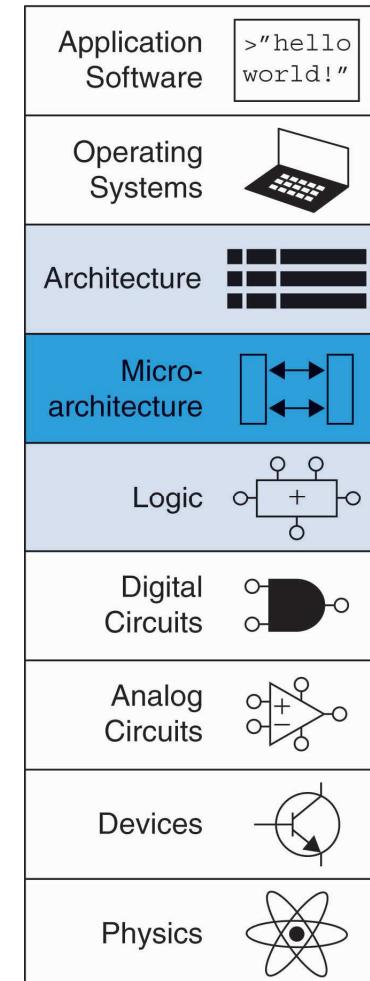
# Chapter 7 :: Topics

- Introduction
- Performance Analysis
- Single-Cycle Processor
- Multicycle Processor
- Pipelined Processor
- Advanced Microarchitecture



# Introduction

- **Microarchitecture:** how to implement an architecture in hardware
- Processor:
  - **Datapath:** functional blocks
  - **Control:** control signals



# Microarchitecture

- Multiple implementations for a single architecture:
  - **Single-cycle:** Each instruction executes in a single cycle
  - **Multicycle:** Each instruction is broken up into series of shorter steps
  - **Pipelined:** Each instruction broken up into series of steps & multiple instructions execute at once



# Processor Performance

- **Program execution time**

**Execution Time = (#instructions)(cycles/instruction)(seconds/cycle)**

- **Definitions:**

- CPI: Cycles/instruction
- clock period: seconds/cycle
- IPC: instructions/cycle = IPC

- **Challenge is to satisfy constraints of:**

- Cost
- Power
- Performance



# ARM Processor

- Consider **subset** of ARM instructions:
  - Data-processing instructions:
    - **ADD, SUB, AND, ORR**
    - with register and immediate Src2, but **no shifts**
  - Memory instructions:
    - **LDR, STR**
    - with **positive immediate offset**
  - Branch instructions:
    - **B**



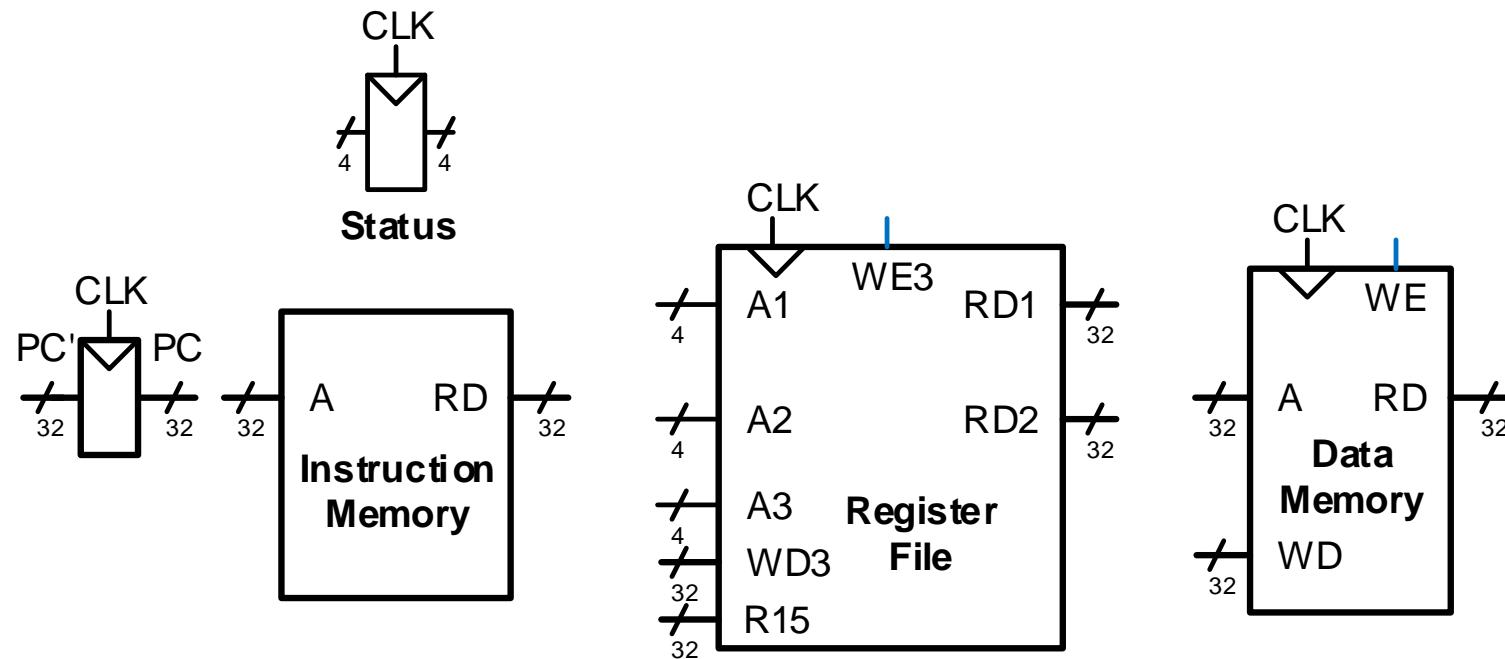
# Architectural State Elements

**Determines everything about a processor:**

- Architectural state:
  - 16 registers (including PC)
  - Status register
- Memory



# ARM Architectural State Elements



# Single-Cycle ARM Processor

- Datapath
- Control



# Single-Cycle ARM Processor

- **Datapath**
- Control

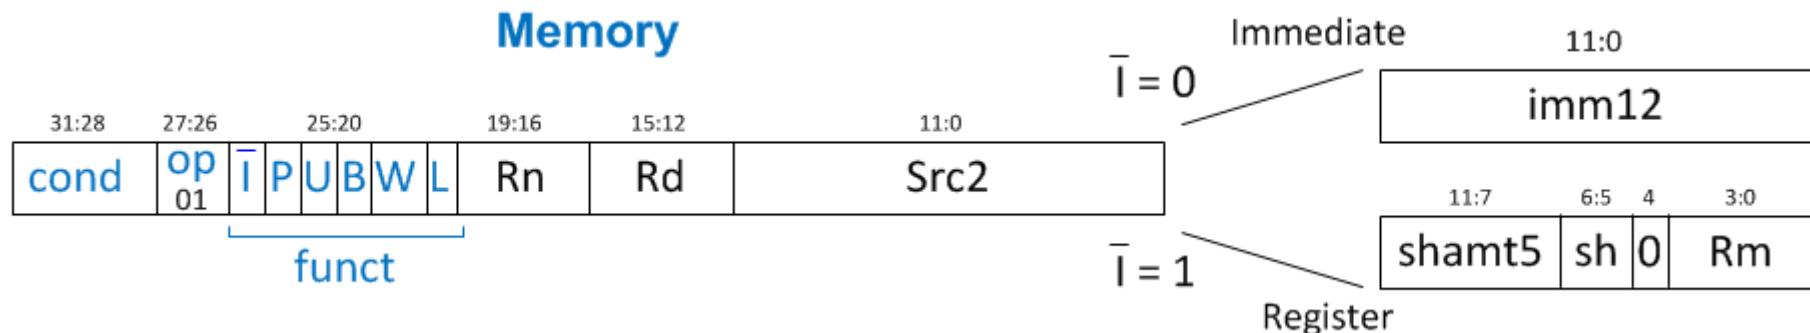


# Single-Cycle ARM Processor

- **Datapath:** start with LDR instruction
- **Example:**

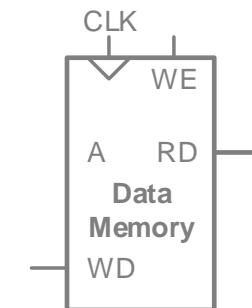
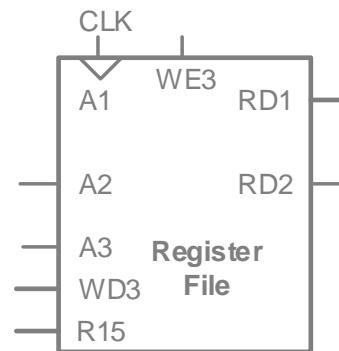
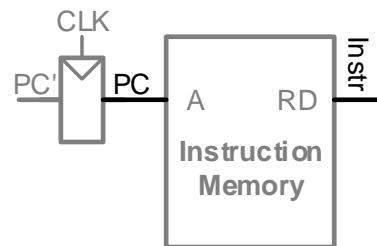
LDR R1, [R2, #5]

LDR Rd, [Rn, imm12]



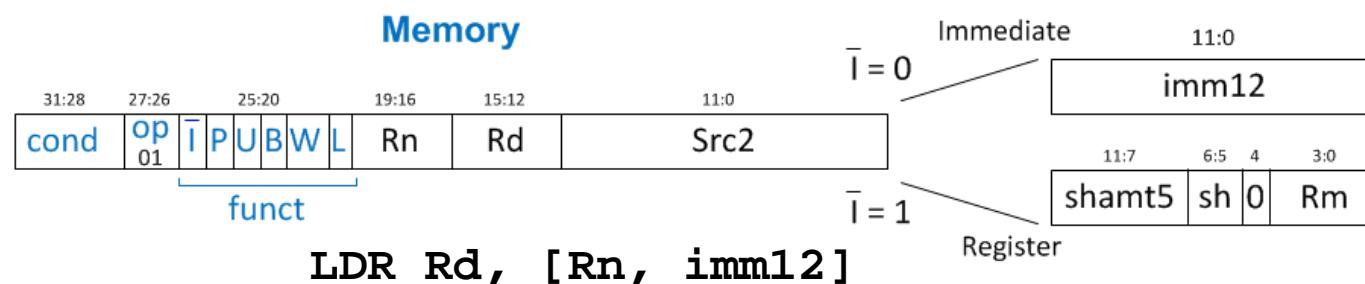
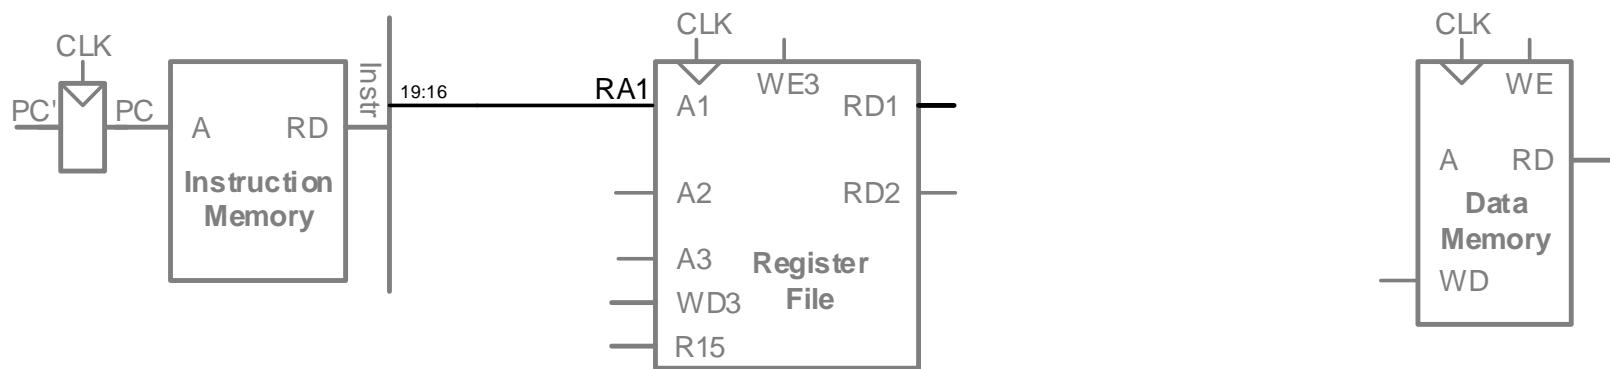
# Single-Cycle Datapath: LDR fetch

## STEP 1: Fetch instruction



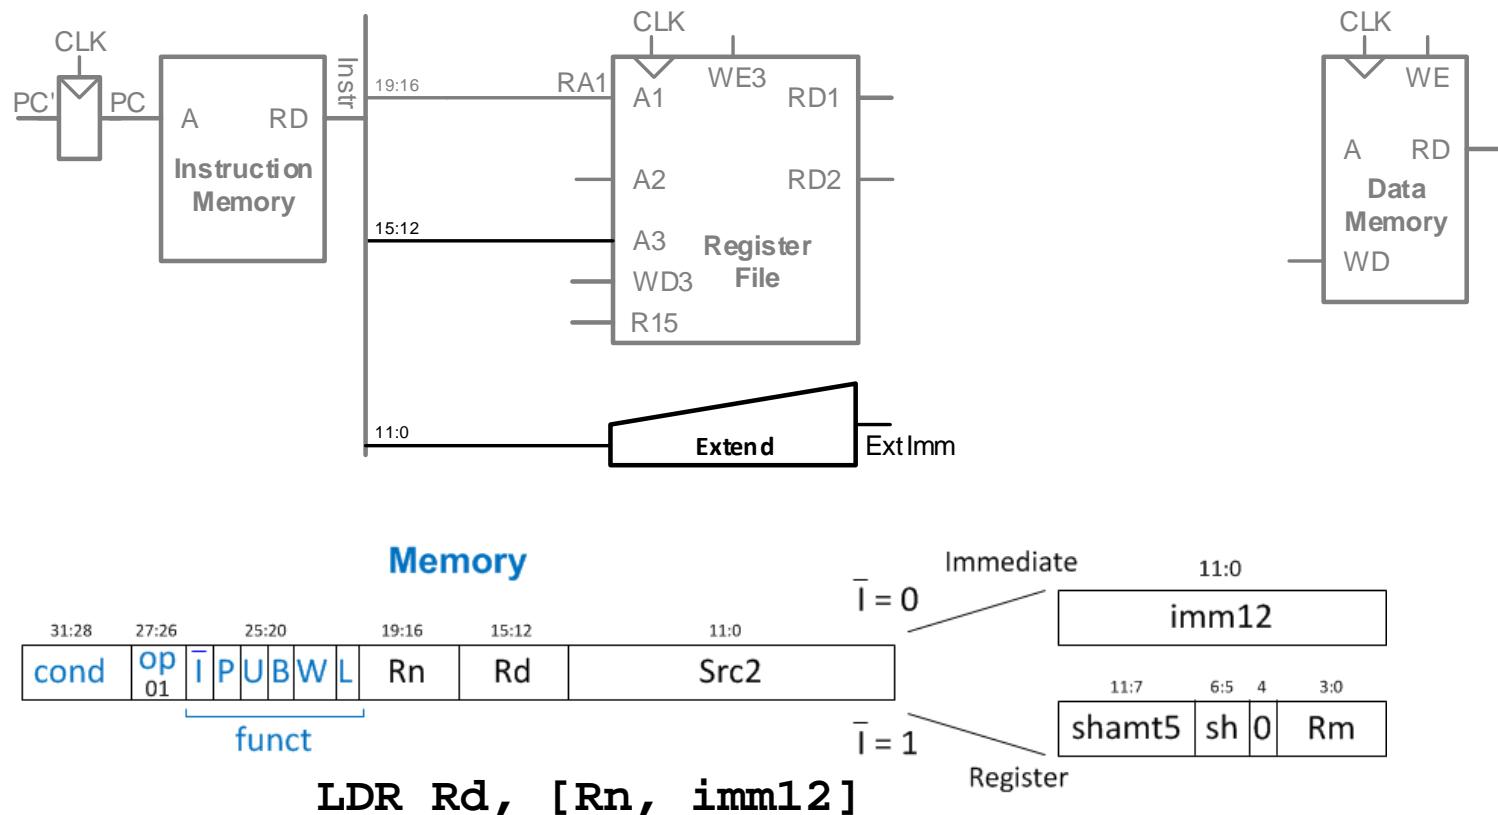
# Single-Cycle Datapath: LDR Reg Read

## STEP 2: Read source operands from RF



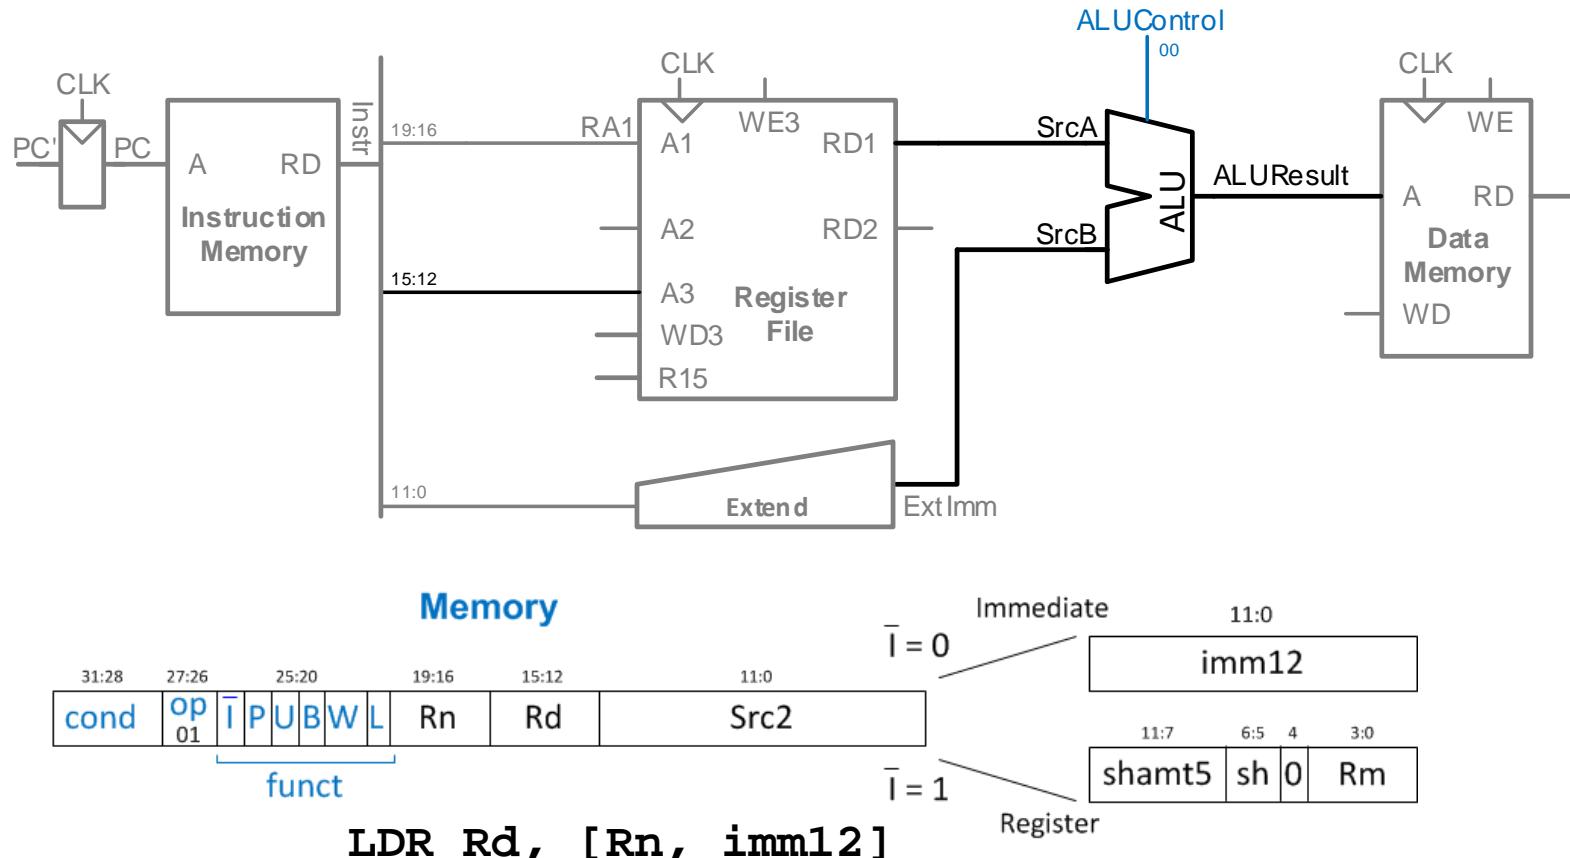
# Single-Cycle Datapath: LDR Immed.

## STEP 3: Extend the immediate



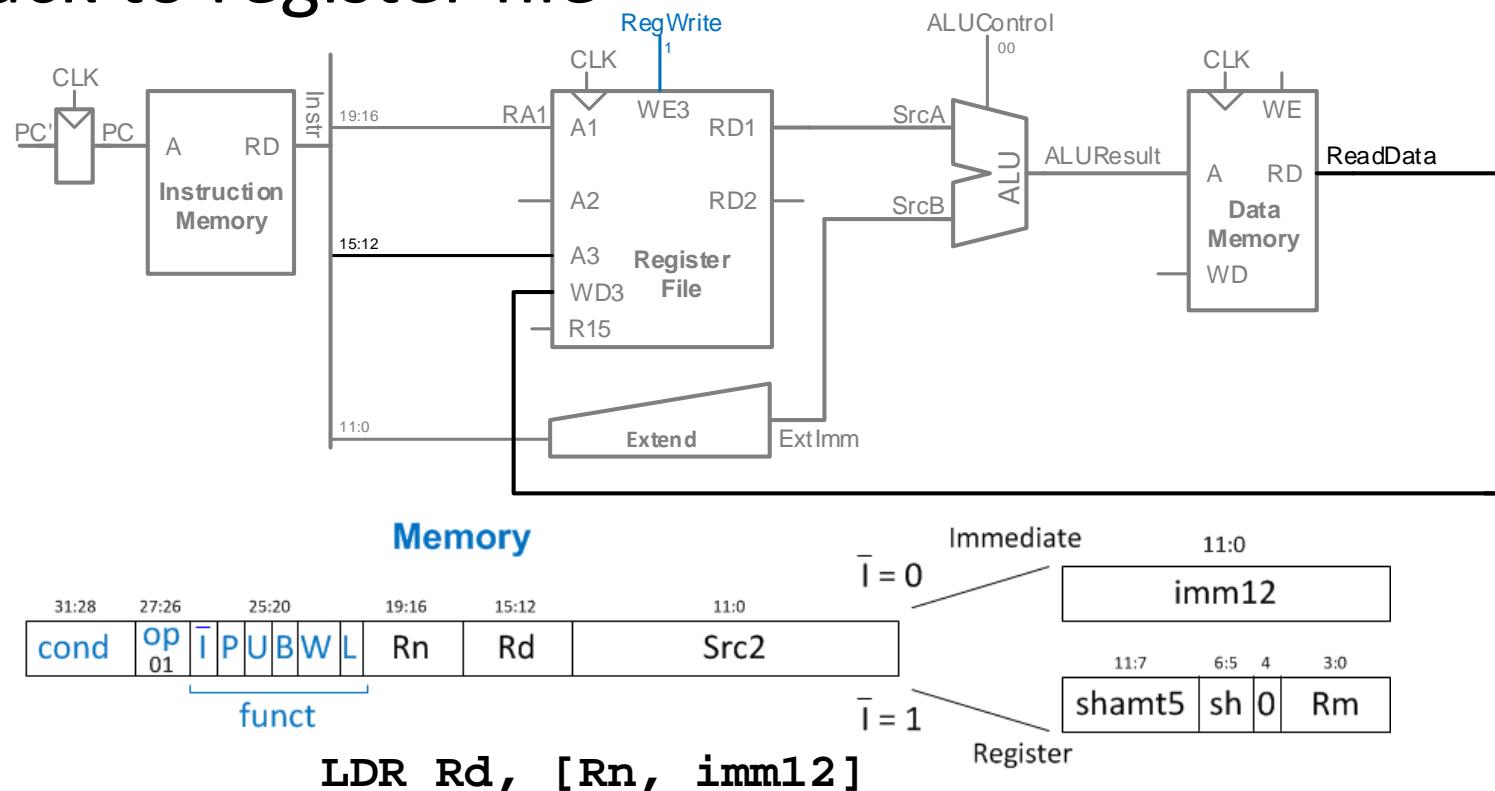
# Single-Cycle Datapath: LDR Address

## STEP 4: Compute the memory address



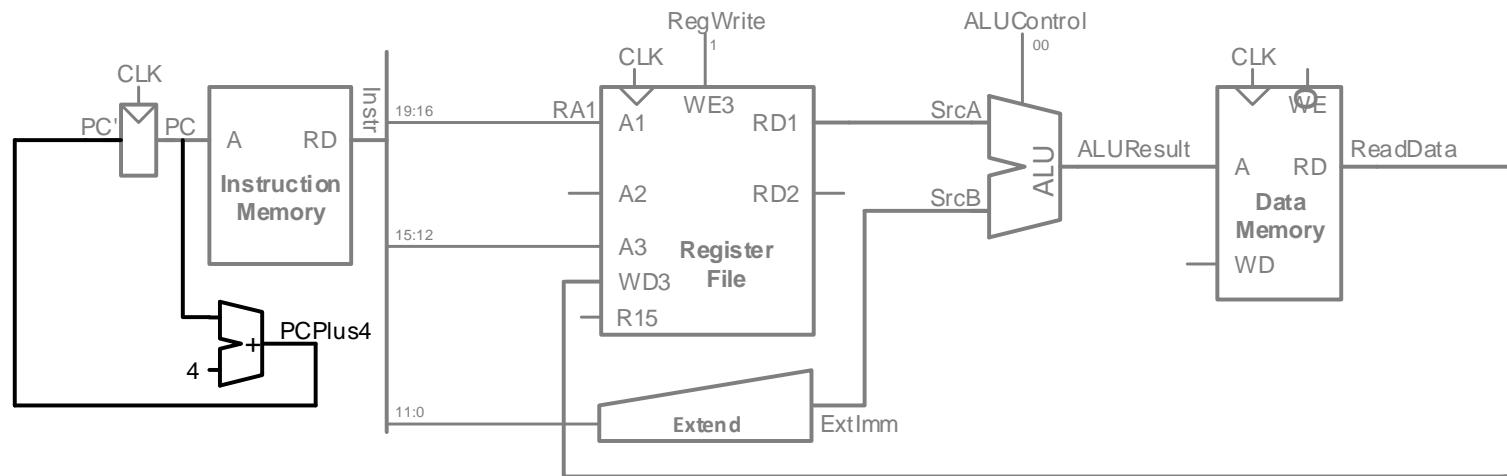
# Single-Cycle Datapath: LDR Mem Read

**STEP 5:** Read data from memory and write it back to register file



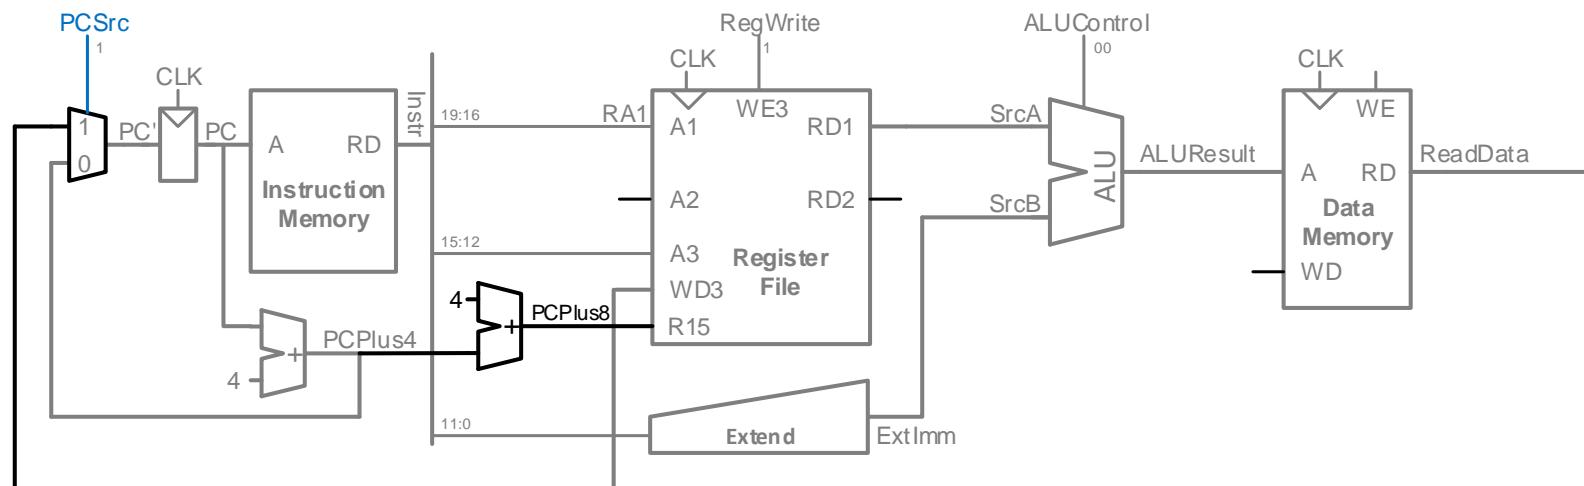
# Single-Cycle Datapath: PC Increment

## STEP 6: Determine address of next instruction



# Single-Cycle Datapath: Access to PC

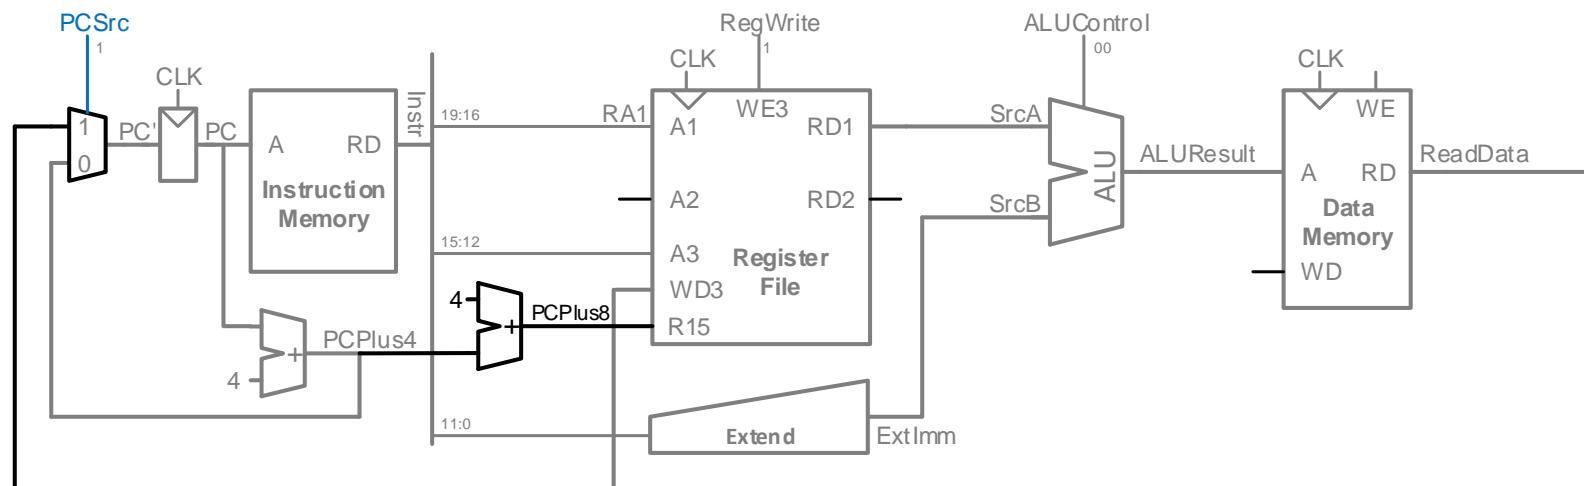
PC can be source/destination of instruction



# Single-Cycle Datapath: Access to PC

PC can be source/destination of instruction

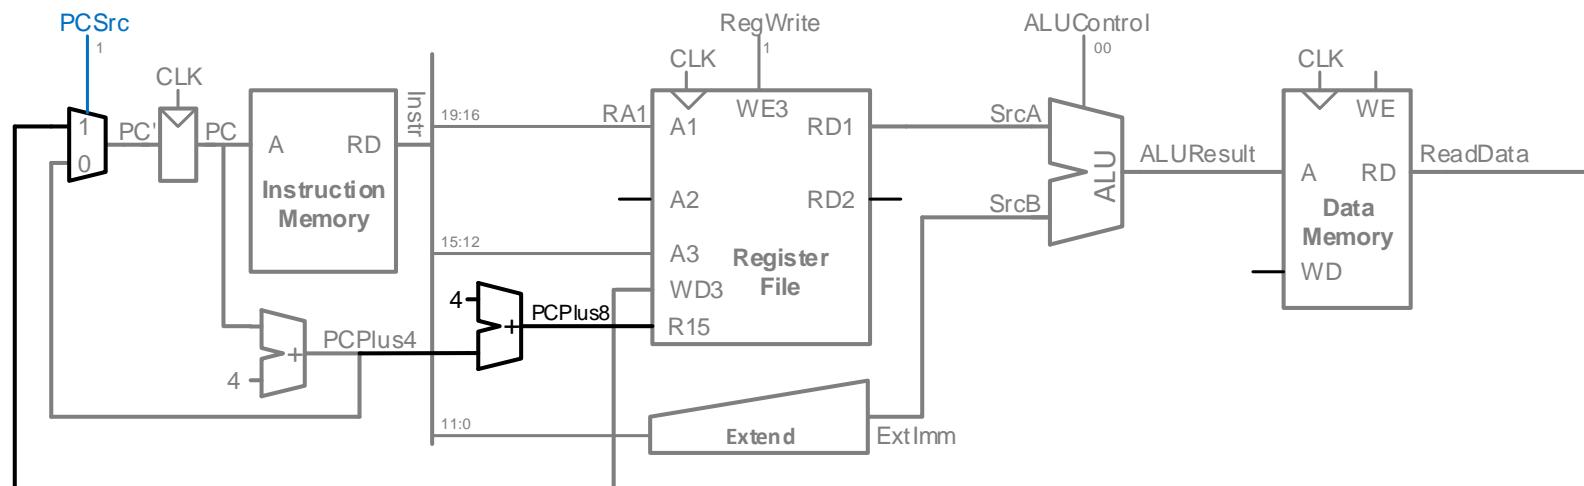
- **Source:** R15 must be available in Register File
  - PC is read as the current PC plus 8



# Single-Cycle Datapath: Access to PC

PC can be source/destination of instruction

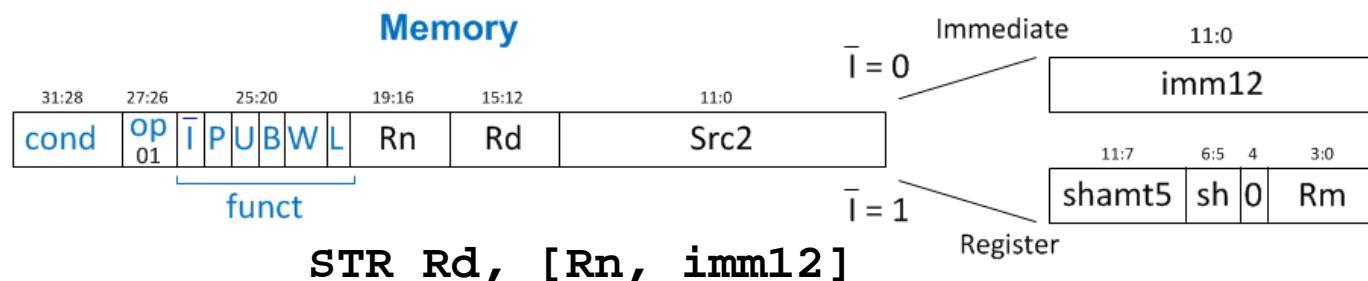
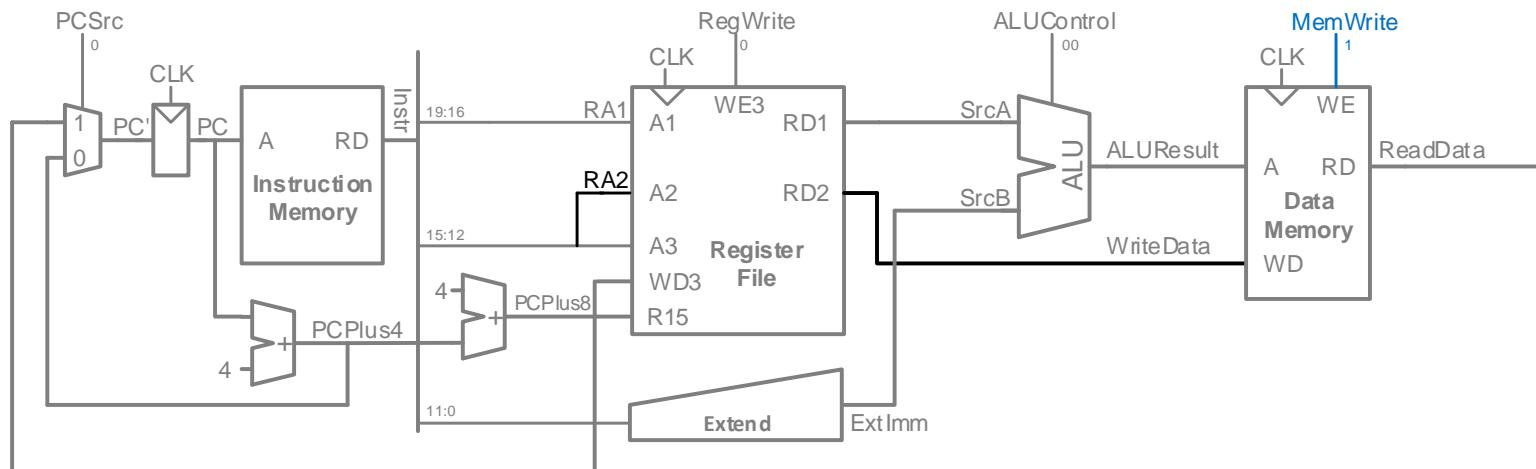
- **Source:** R15 must be available in Register File
  - PC is read as the current PC plus 8
- **Destination:** Be able to write result to PC



# Single-Cycle Datapath: STR

## Expand datapath to handle STR:

- Write data in Rd to memory



# Single-Cycle Datapath: Data-processing

## With immediate Src2:

- Read from Rn and Imm8 (*ImmSrc* chooses the zero-extended Imm8 instead of Imm12)
- Write *ALUResult* to register file
- Write to Rd



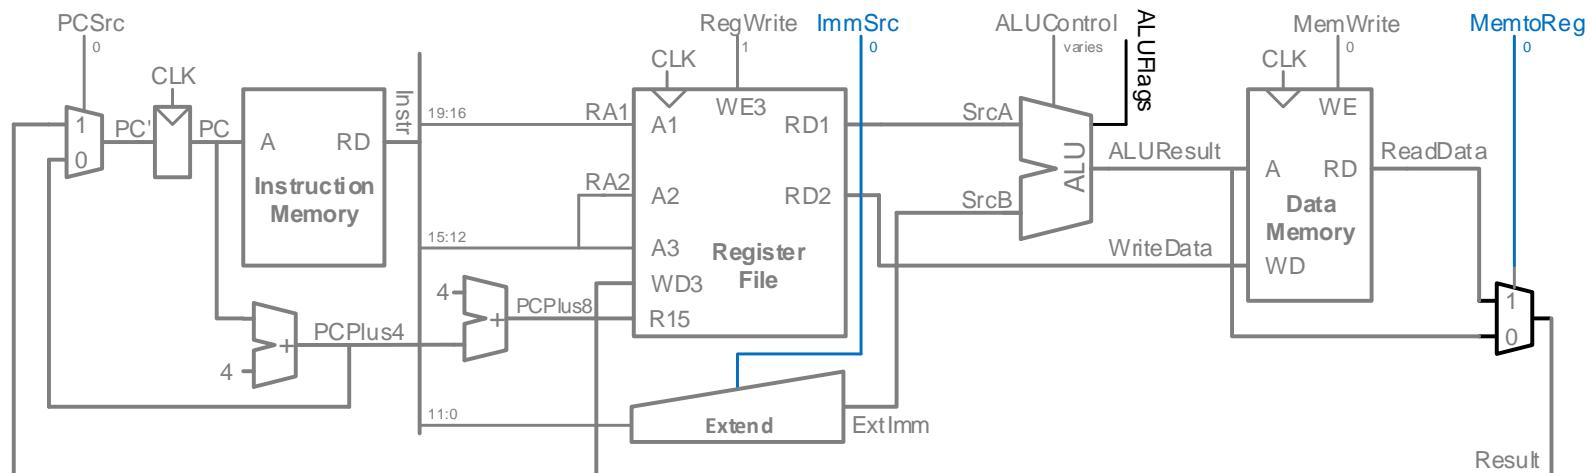
**ADD Rd, Rn, imm8**



# Single-Cycle Datapath: Data-processing

## With immediate Src2:

- Read from Rn and Imm8 (*ImmSrc* chooses the zero-extended Imm8 instead of Imm12)
- Write *ALUResult* to register file
- Write to Rd



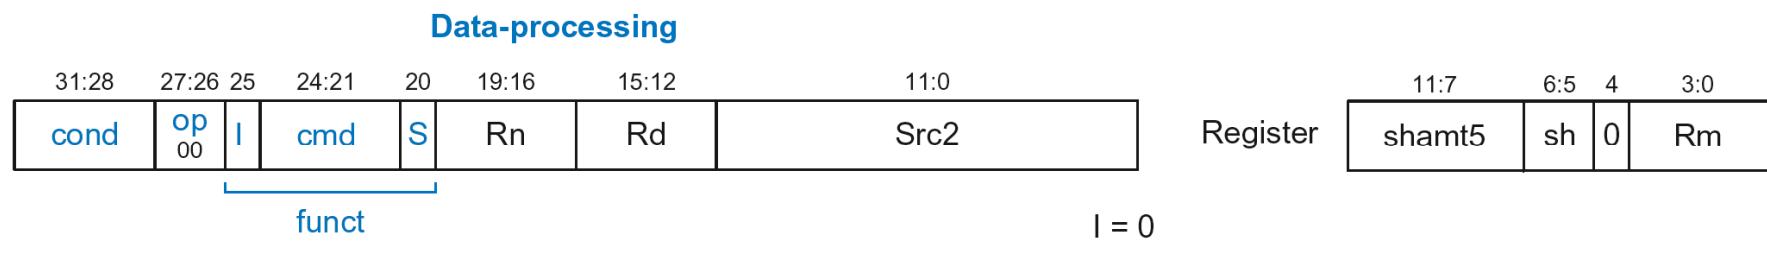
**ADD Rd, Rn, imm8**



# Single-Cycle Datapath: Data-processing

## With register Src2:

- Read from Rn and Rm (instead of Imm8 )
- Write *ALUResult* to register file
- Write to Rd



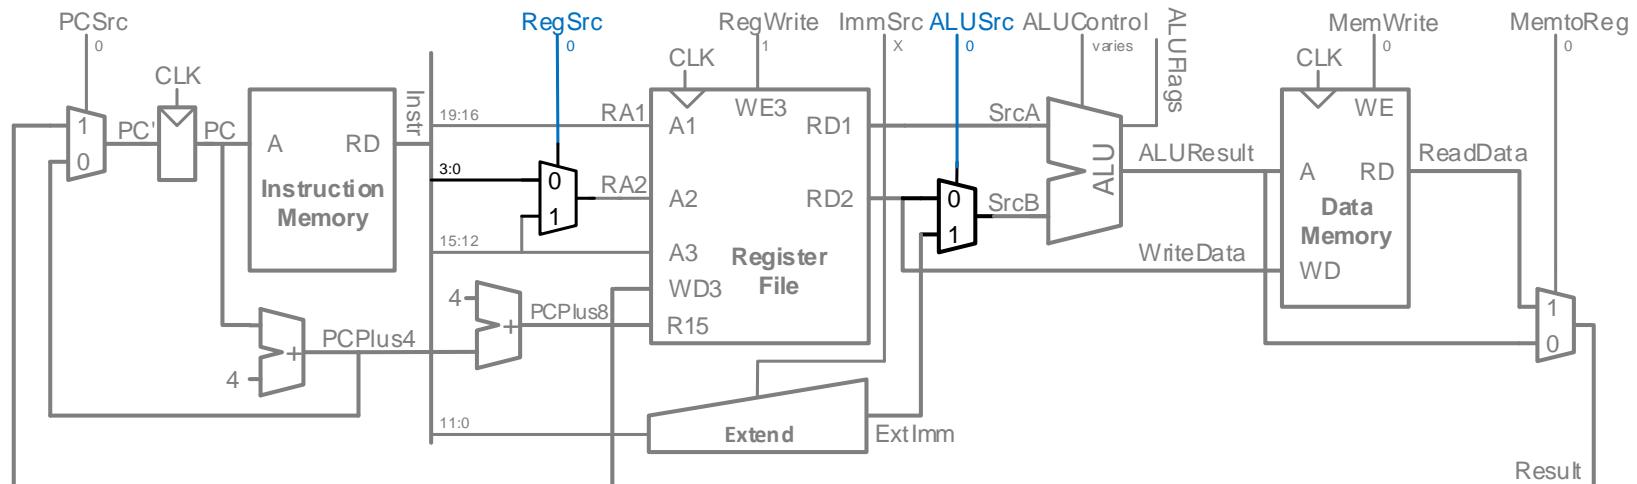
**ADD Rd, Rn, Rm**



# Single-Cycle Datapath: Data-processing

## With register Src2:

- Read from Rn and Rm (instead of Imm8 )
- Write *ALUResult* to register file
- Write to Rd



ADD Rd, Rn, Rm

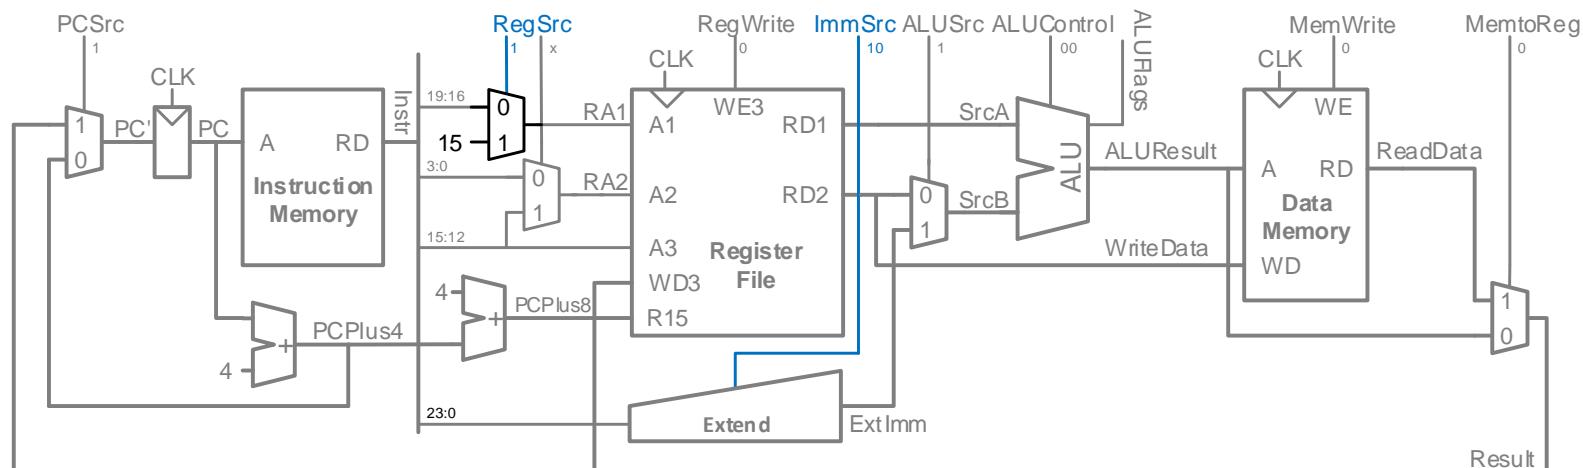


# Single-Cycle Datapath: B

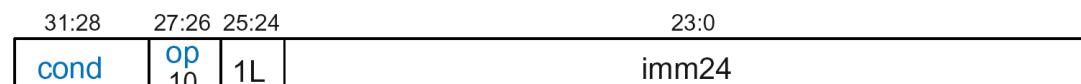
## Calculate branch target address:

$$BTA = (Ext/imm) + (PC + 8)$$

$Ext/imm = Imm24 \ll 2$  and sign-extended



Branch

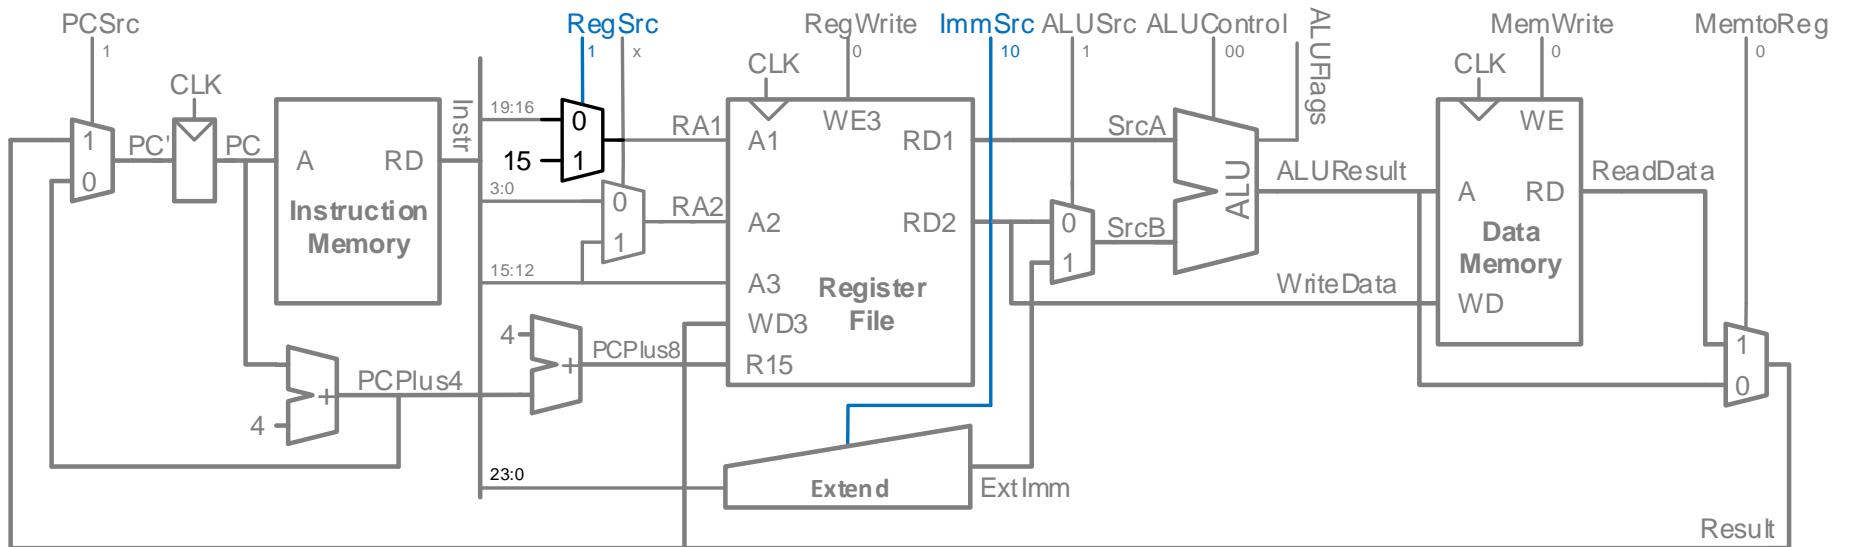


funct

B Label



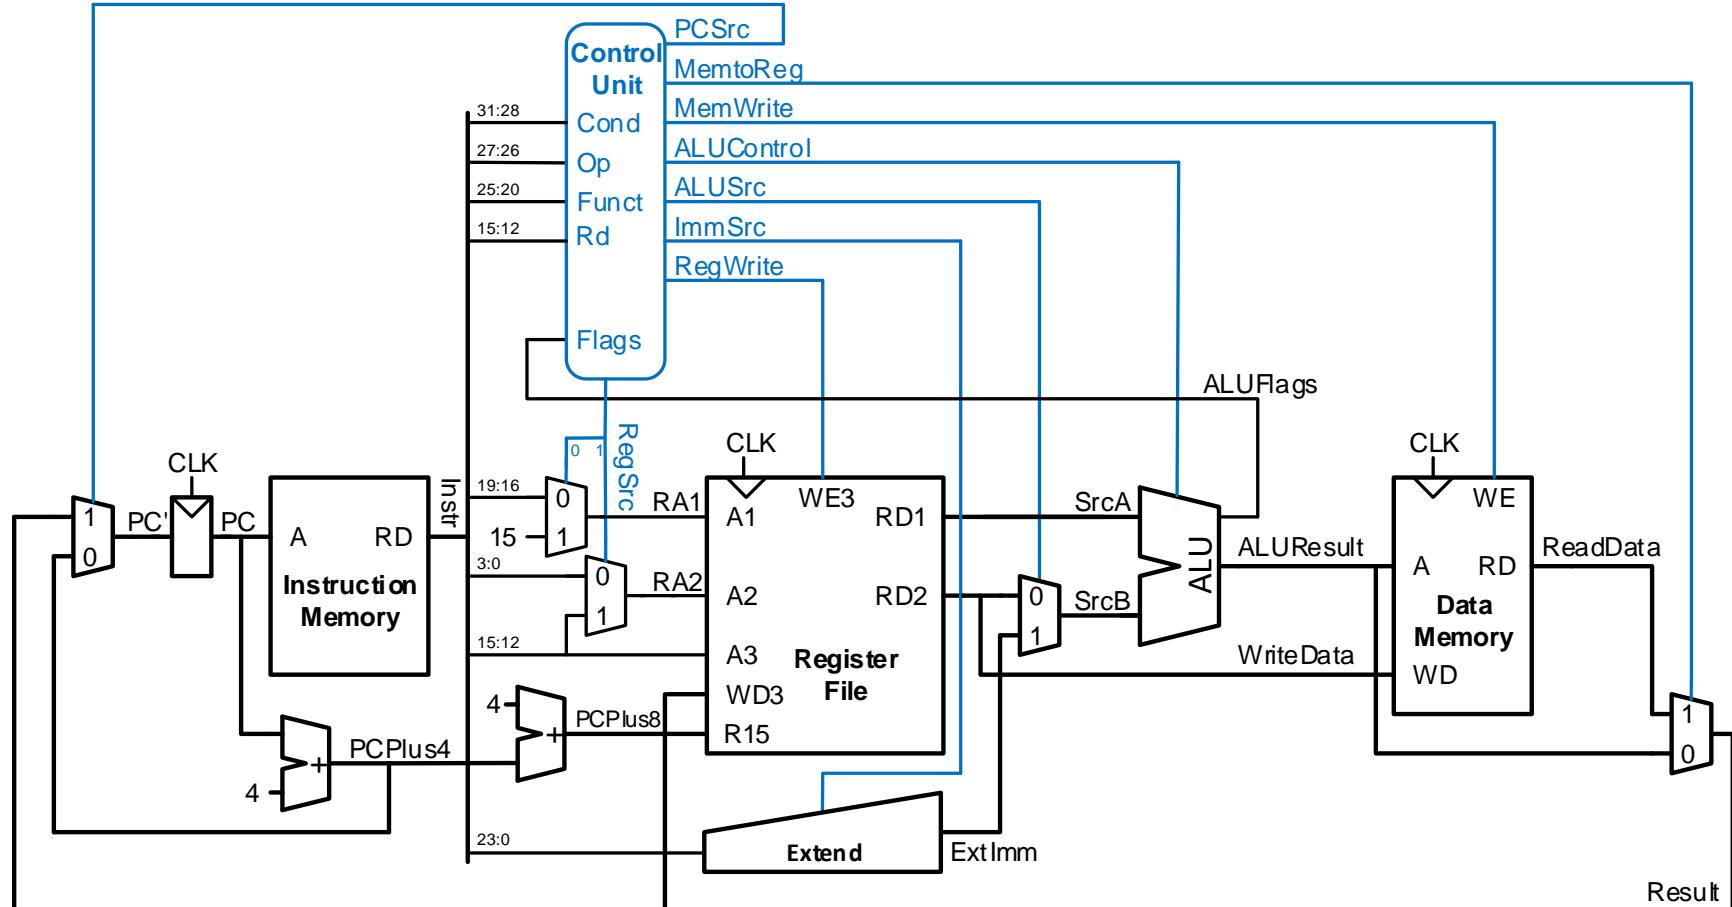
# Single-Cycle Datapath: ExtImm



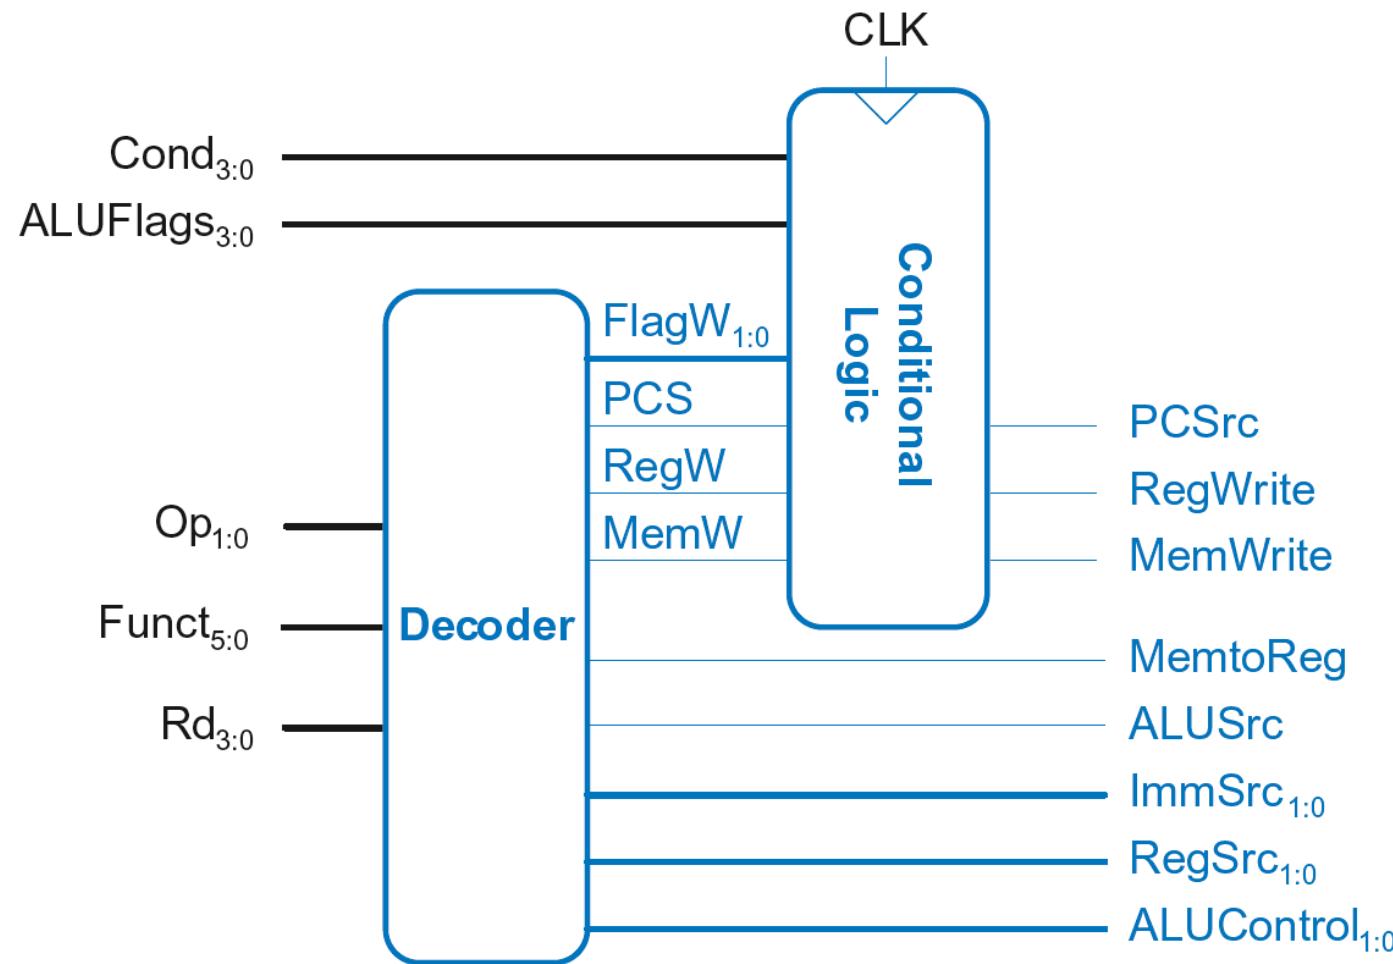
ImmSrc <sub>1:0</sub>	ExtImm	Description
00	{24'b0, Instr <sub>7:0</sub> }	Zero-extended <i>imm8</i>
01	{20'b0, Instr <sub>11:0</sub> }	Zero-extended <i>imm12</i>
10	{6{Instr <sub>23</sub> }, Instr <sub>23:0</sub> }	Sign-extended <i>imm24</i>



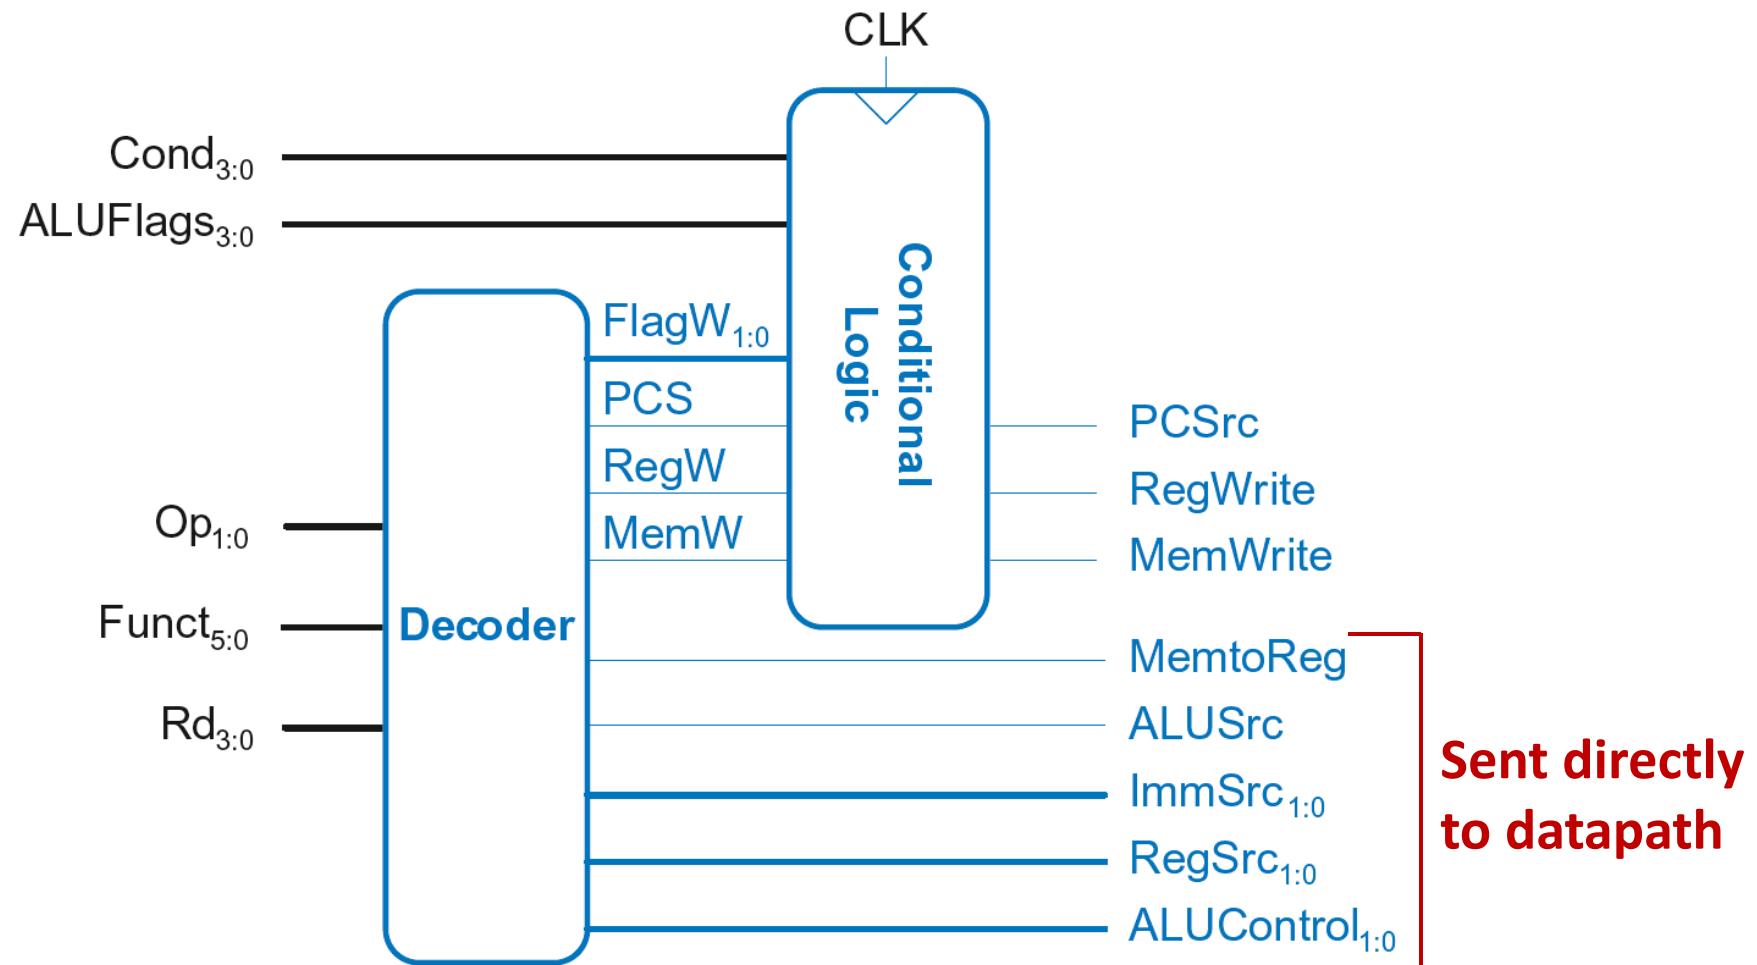
# Single-Cycle ARM Processor



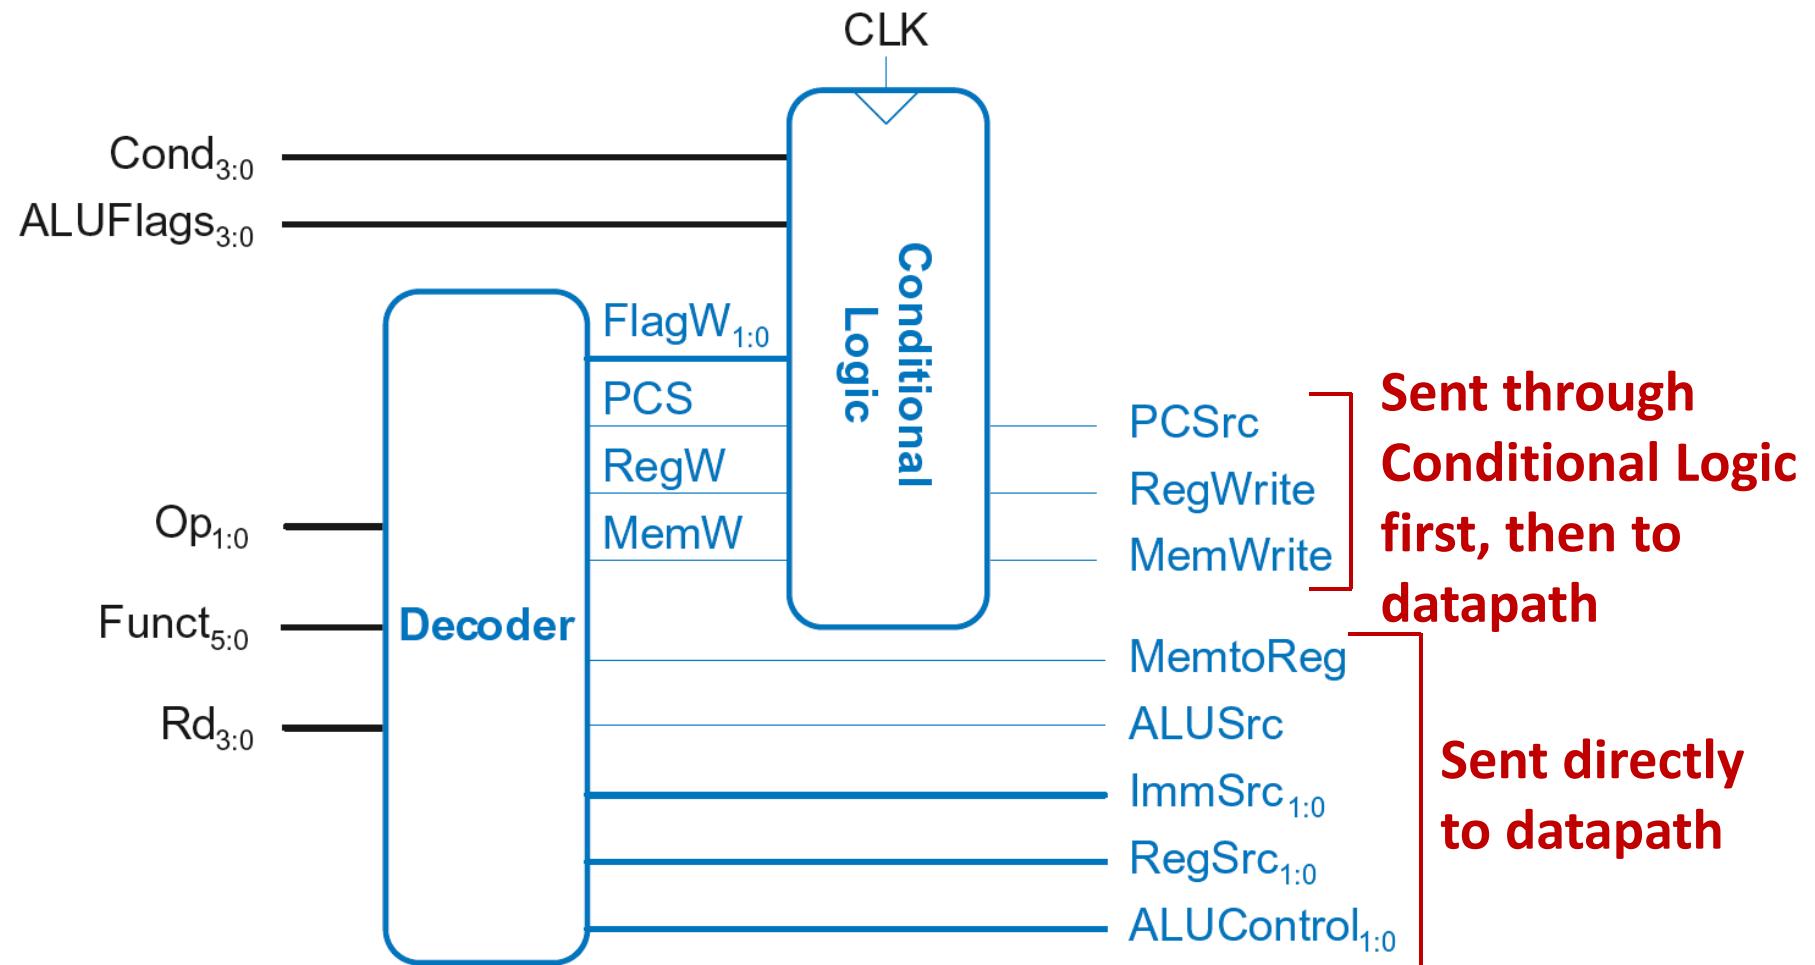
# Single-Cycle Control



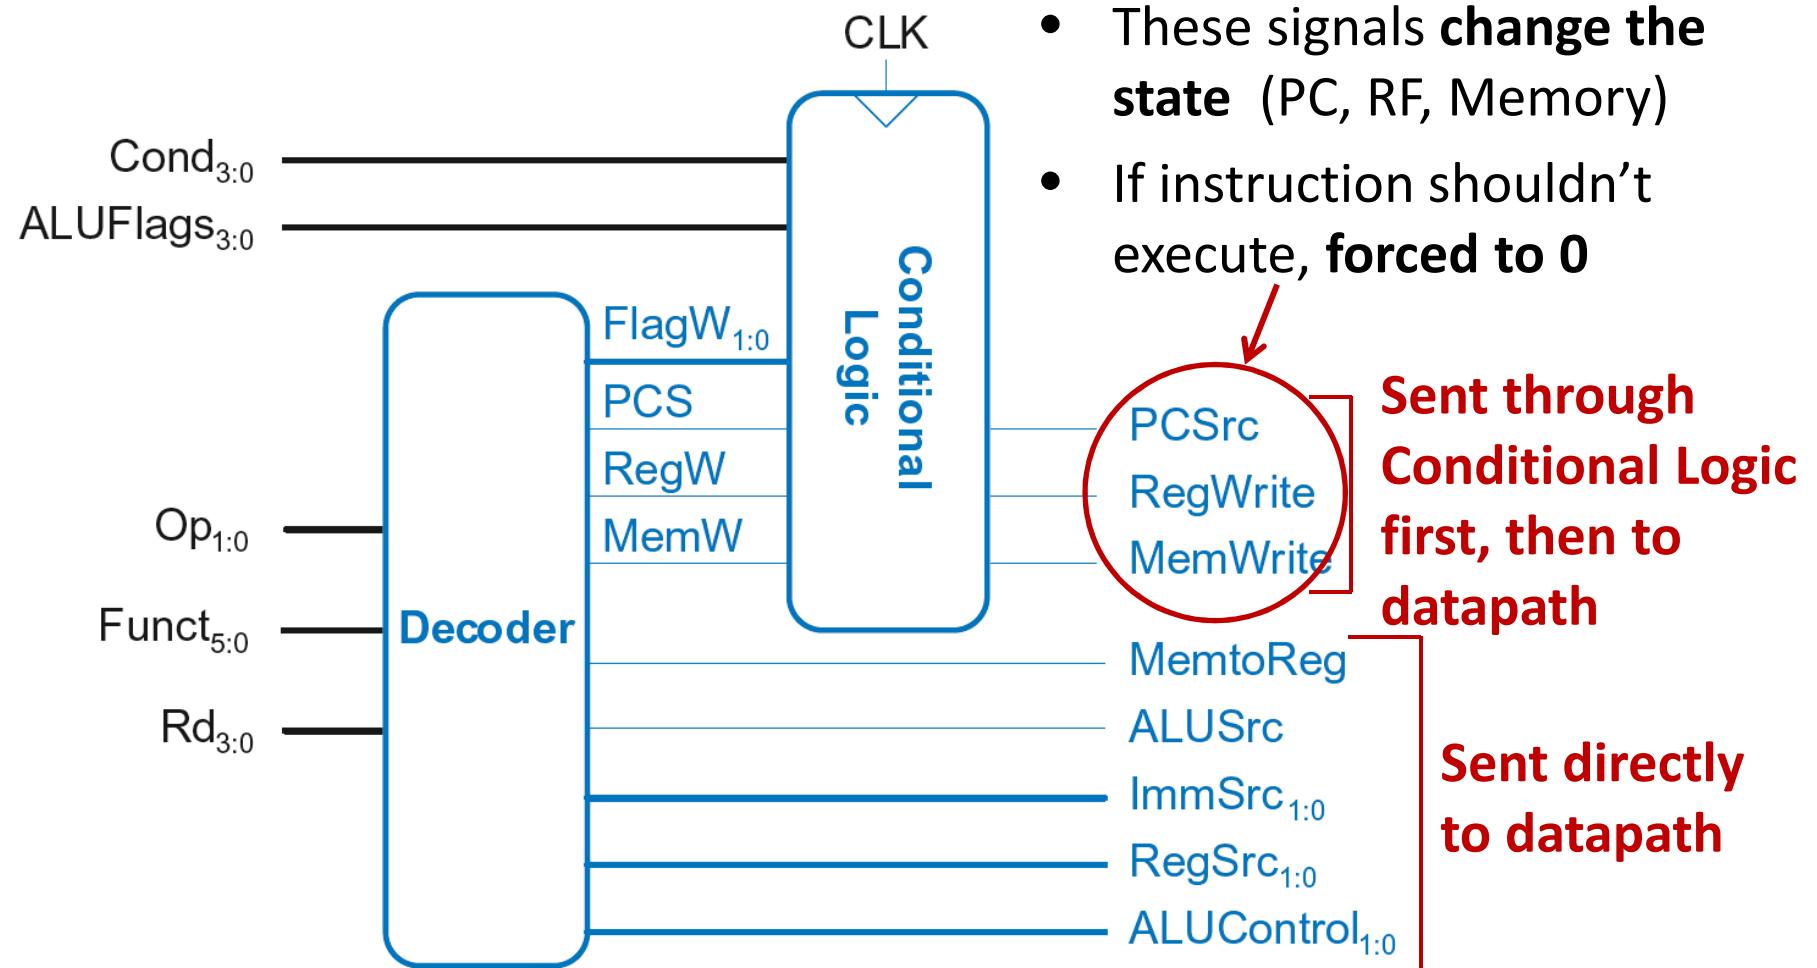
# Single-Cycle Control



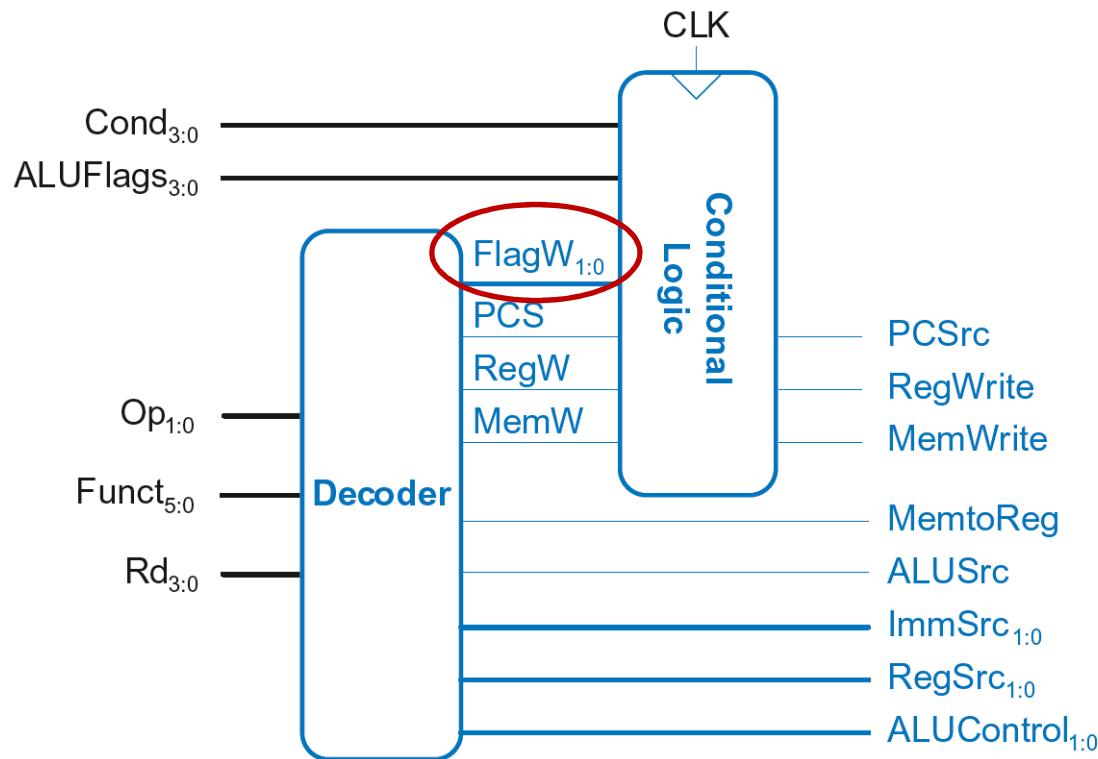
# Single-Cycle Control



# Single-Cycle Control



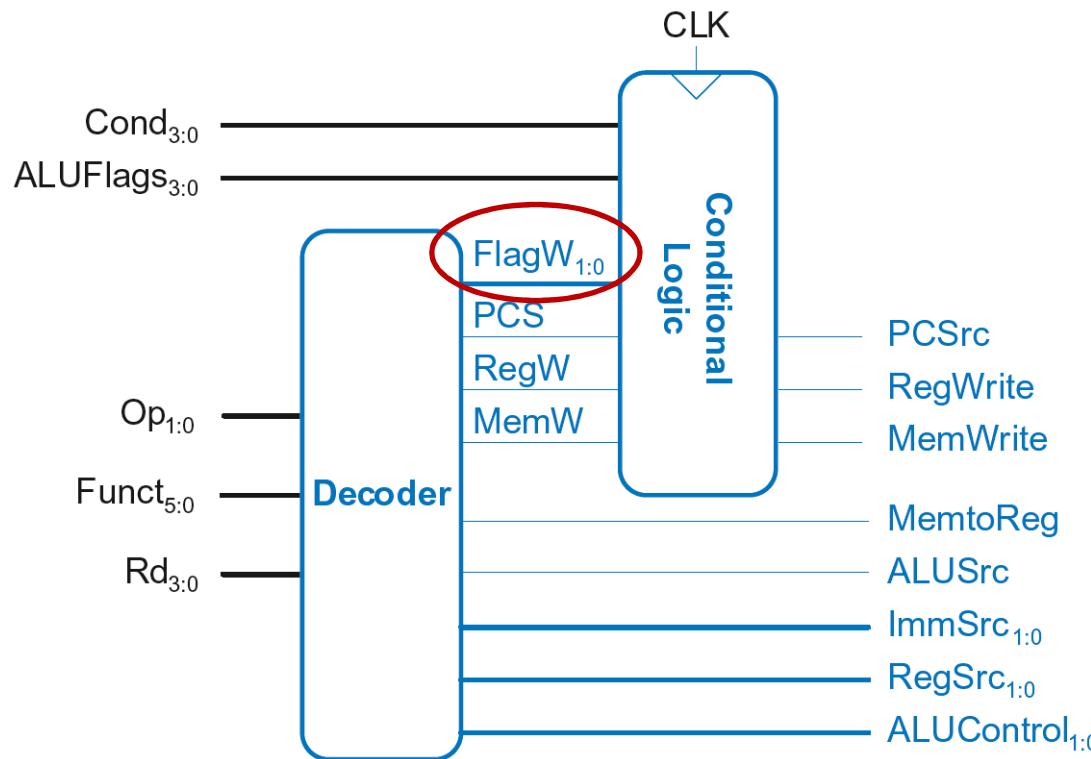
# Single-Cycle Control



- **FlagW<sub>1:0</sub>:** Flag Write signal, asserted when *ALUFlags* should be saved (i.e., on instruction with S=1)



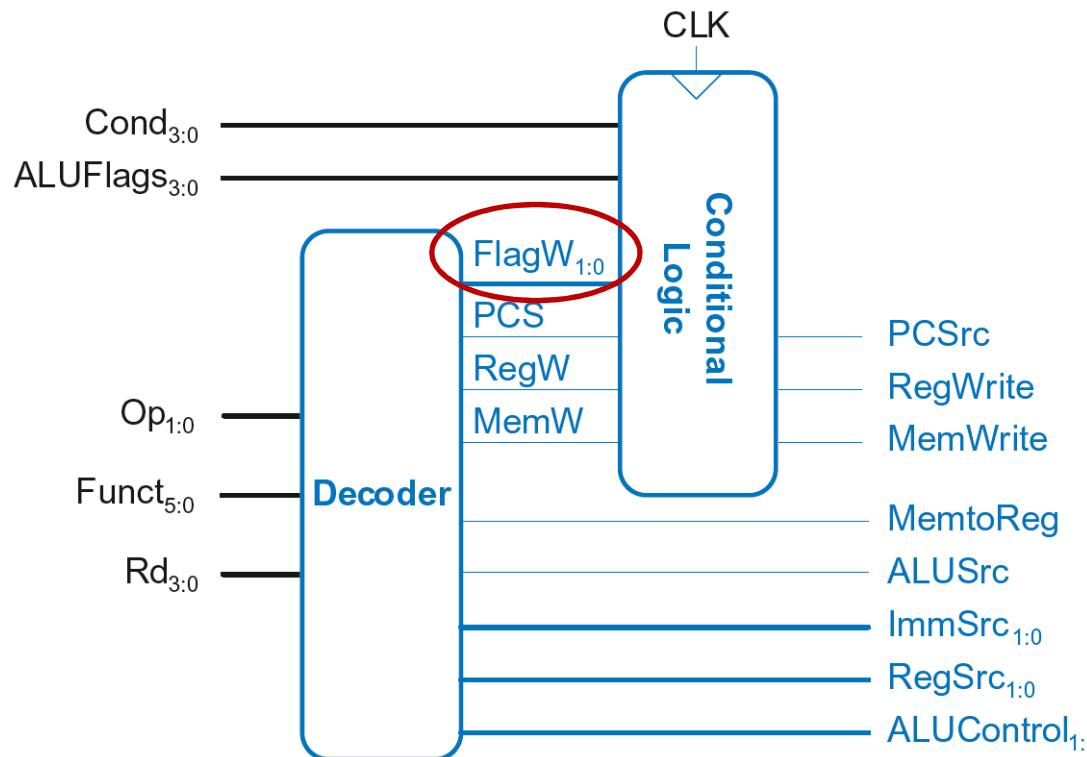
# Single-Cycle Control



- $\text{FlagW}_{1:0}$ : Flag Write signal, asserted when *ALUFlags* should be saved (i.e., on instruction with S=1)
- ADD, SUB update all flags (**NZCV**)
- AND, ORR only update **NZ** flags



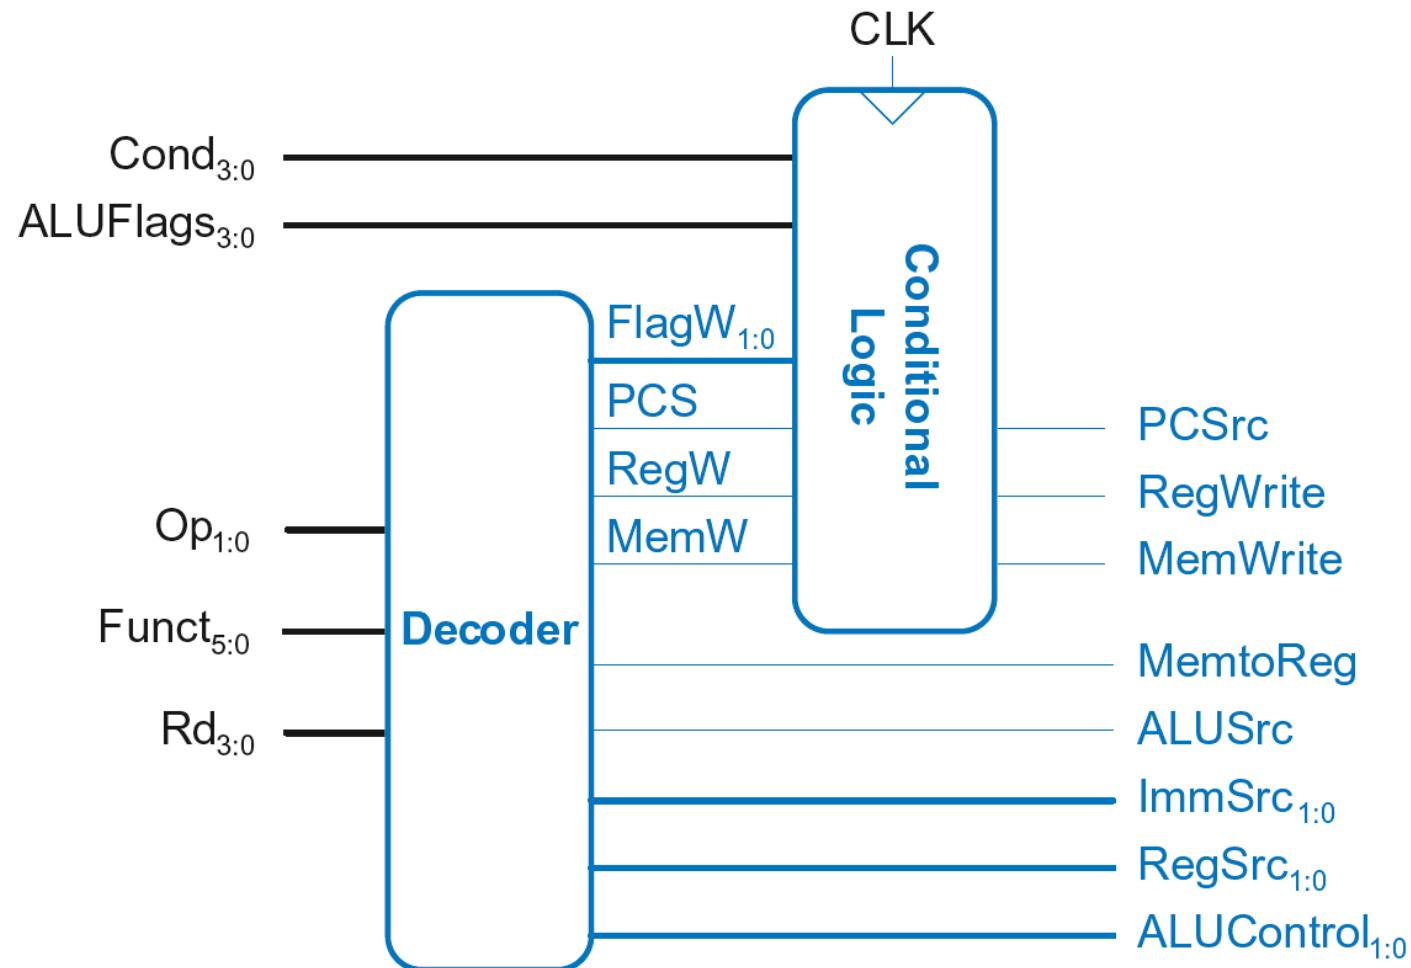
# Single-Cycle Control



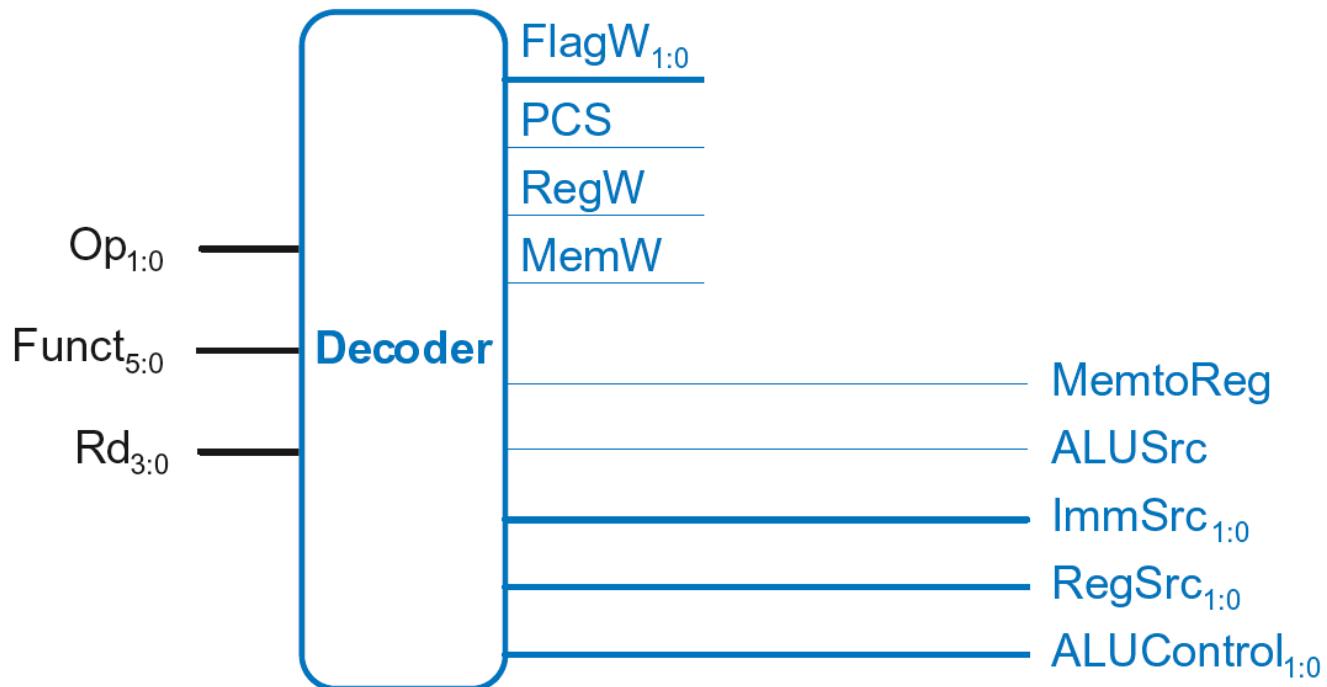
- $\text{FlagW}_{1:0}$ : Flag Write signal, asserted when  $ALUFlags$  should be saved (i.e., on instruction with S=1)
- ADD, SUB update all flags (**NZCV**)
- AND, ORR only update **NZ** flags
- So, two bits needed:
  - $\text{FlagW}_1 = 1$ : NZ saved ( $ALUFlags_{3:2}$  saved)
  - $\text{FlagW}_0 = 1$ : CV saved ( $ALUFlags_{1:0}$  saved)



# Single-Cycle Control



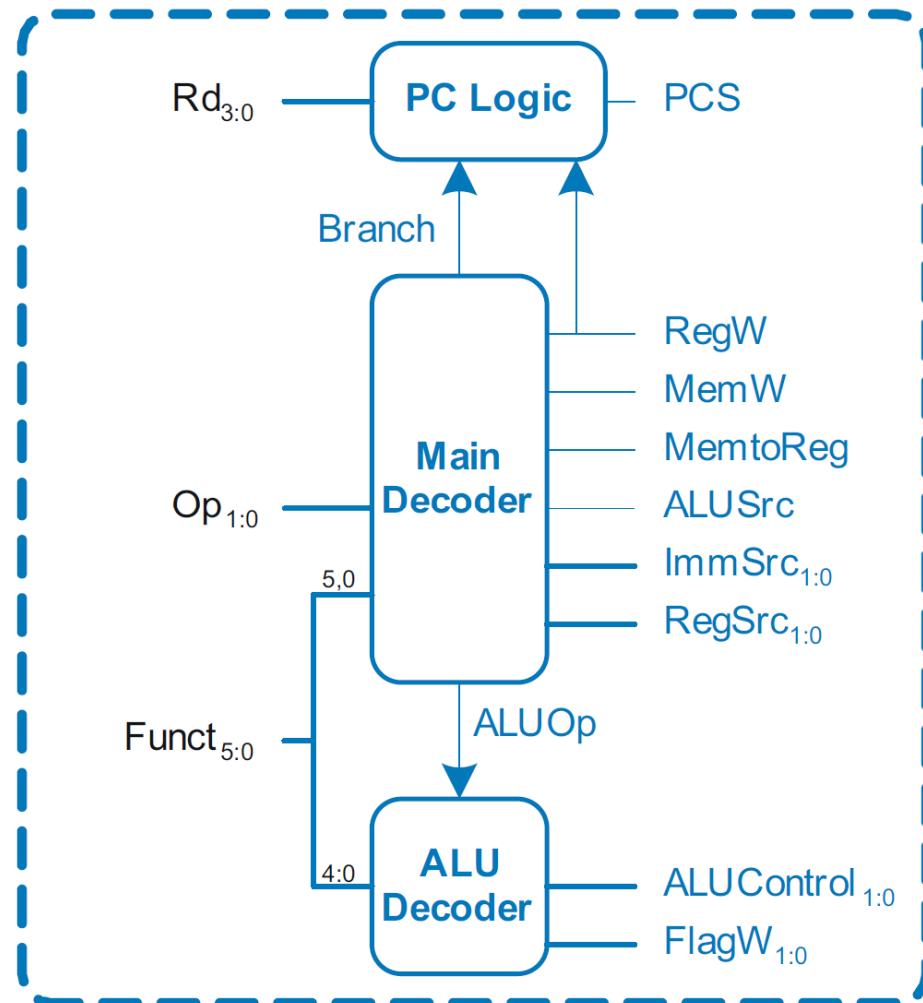
# Single-Cycle Control: Decoder



# Single-Cycle Control: Decoder

## Submodules:

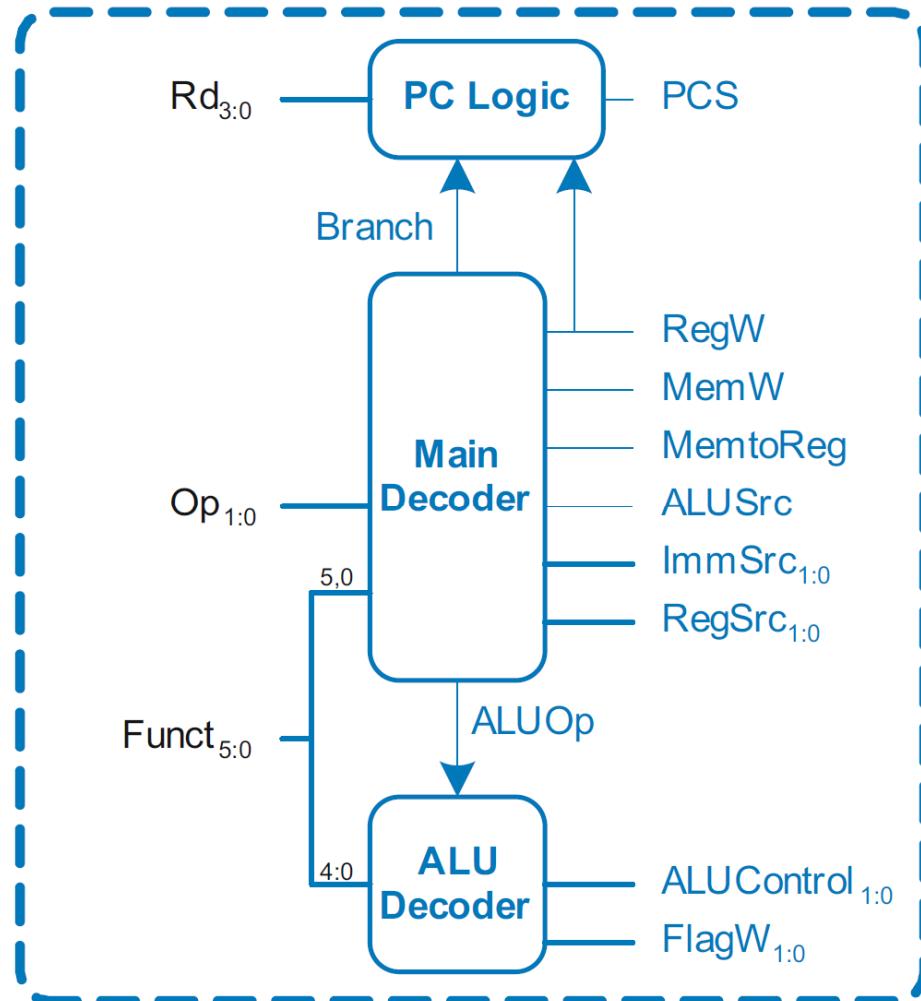
- Main Decoder
- ALU Decoder
- PC Logic



# Single-Cycle Control: Decoder

## Submodules:

- Main Decoder
- ALU Decoder
- PC Logic



# Control Unit: Main Decoder

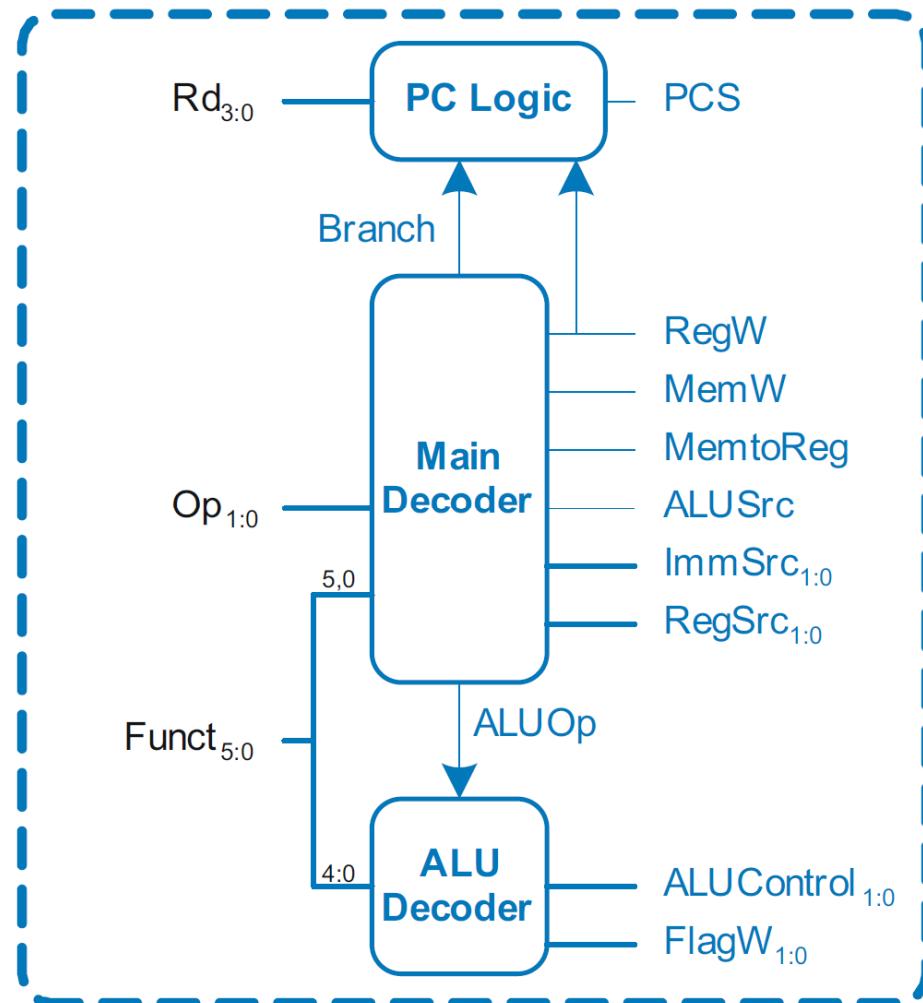
Op	Funct <sub>5</sub>	Funct <sub>0</sub>	Type	Branch	MemtoReg	MemW	ALUSrc	ImmSrc	RegW	RegSrc	ALUOp
00	0	X	DP Reg	0	0	0	0	XX	1	00	1
00	1	X	DP Imm	0	0	0	1	00	1	X0	1
01	X	0	STR	0	X	1	1	01	0	10	0
01	X	1	LDR	0	1	0	1	01	1	X0	0
11	X	X	B	1	0	0	1	10	0	X1	0



# Single-Cycle Control: Decoder

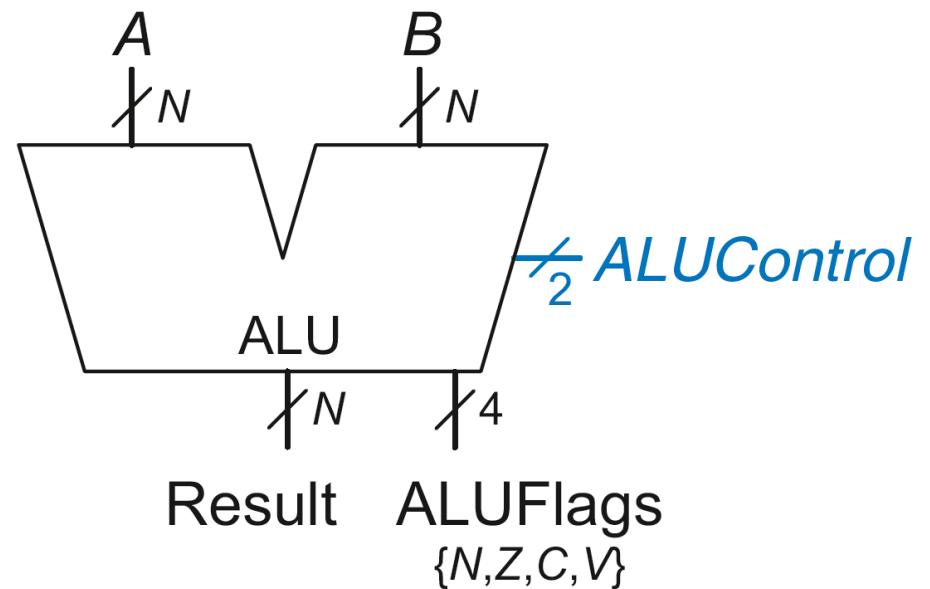
## Submodules:

- Main Decoder
- **ALU Decoder**
- PC Logic

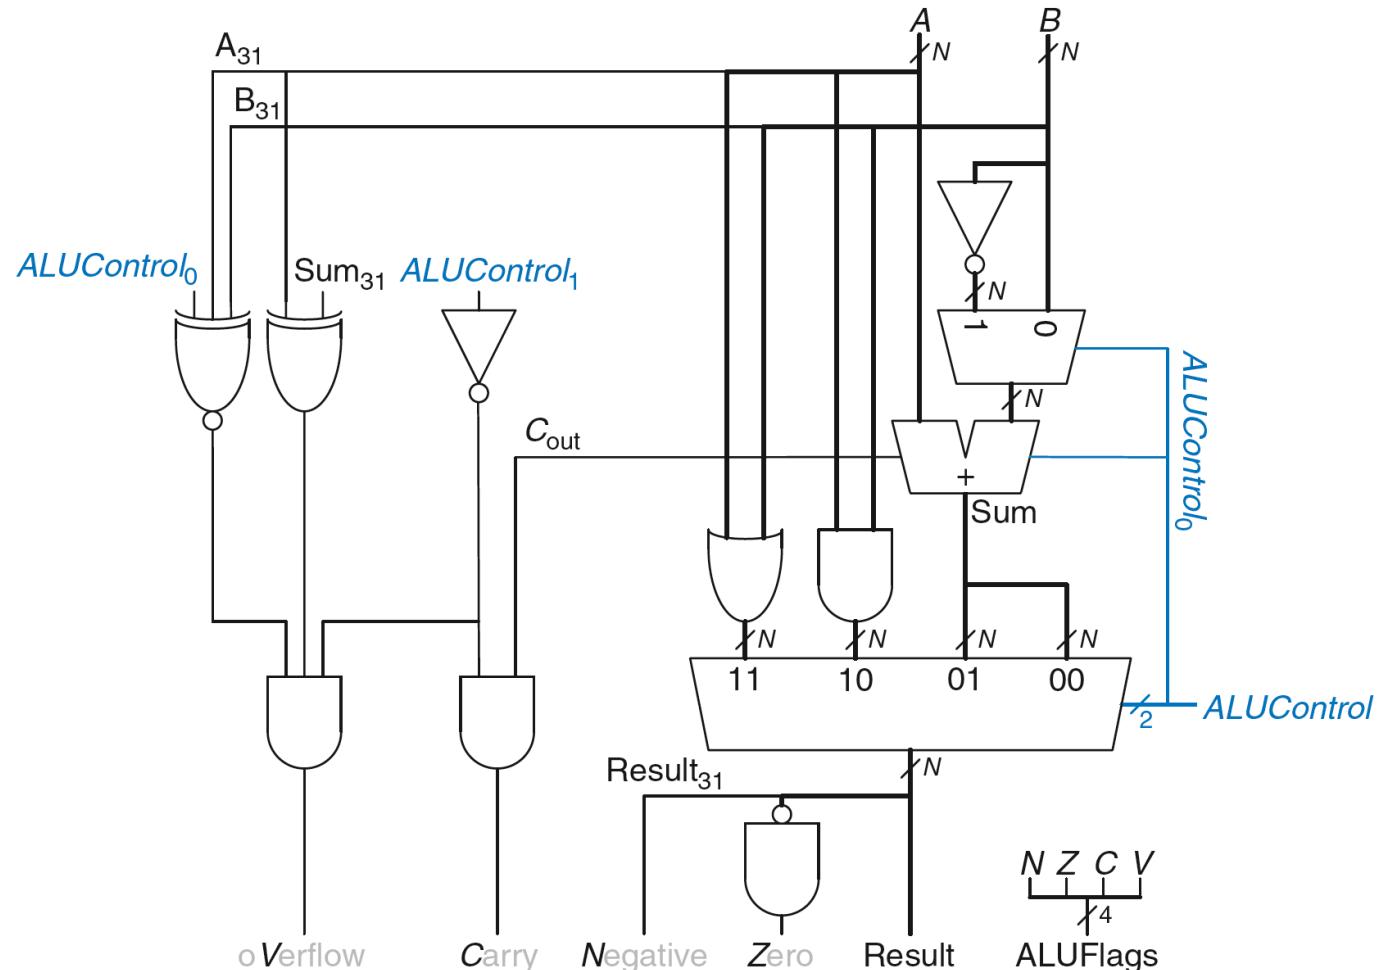


# Review: ALU

ALUControl <sub>1:0</sub>	Function
00	Add
01	Subtract
10	AND
11	OR



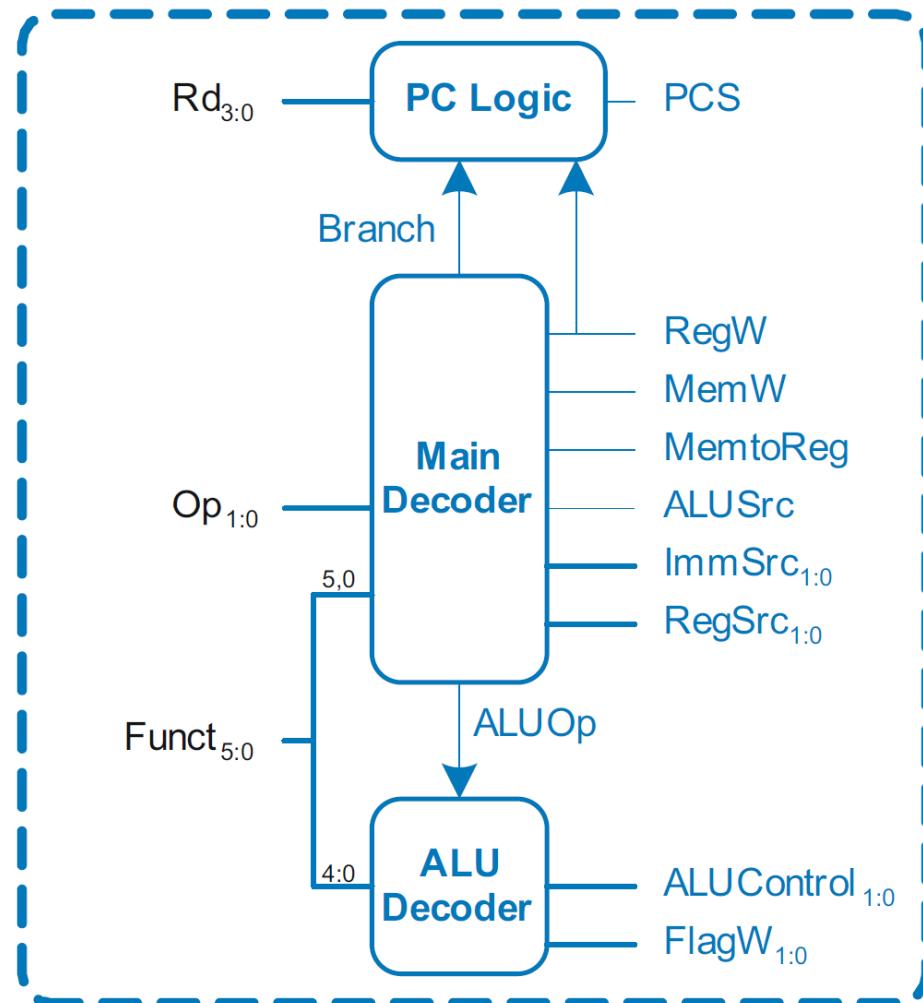
# Review: ALU



# Single-Cycle Control: Decoder

## Submodules:

- Main Decoder
- **ALU Decoder**
- PC Logic



# Control Unit: ALU Decoder

ALUOp	Funct <sub>4:1</sub> (cmd)	Funct <sub>0</sub> (S)	Type	ALUControl <sub>1:0</sub>	FlagW <sub>1:0</sub>
0	X	X	Not DP	00	00
1	0100	0	ADD	00	00
		1			11
	0010	0	SUB	01	00
		1			11
	0000	0	AND	10	00
		1			10
	1100	0	ORR	11	00
		1			10

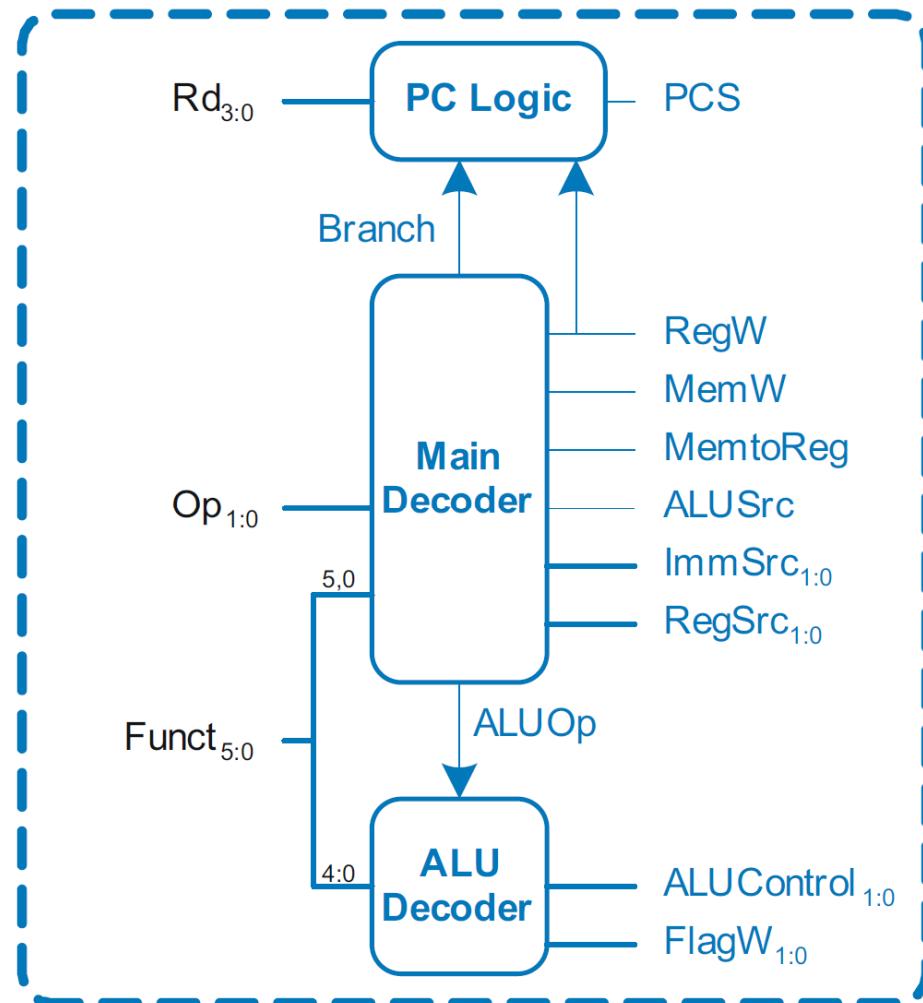
- $\text{FlagW}_1 = 1$ : NZ(Flags<sub>3:2</sub>) should be saved
- $\text{FlagW}_0 = 1$ : CV(Flags<sub>1:0</sub>) should be saved



# Single-Cycle Control: Decoder

## Submodules:

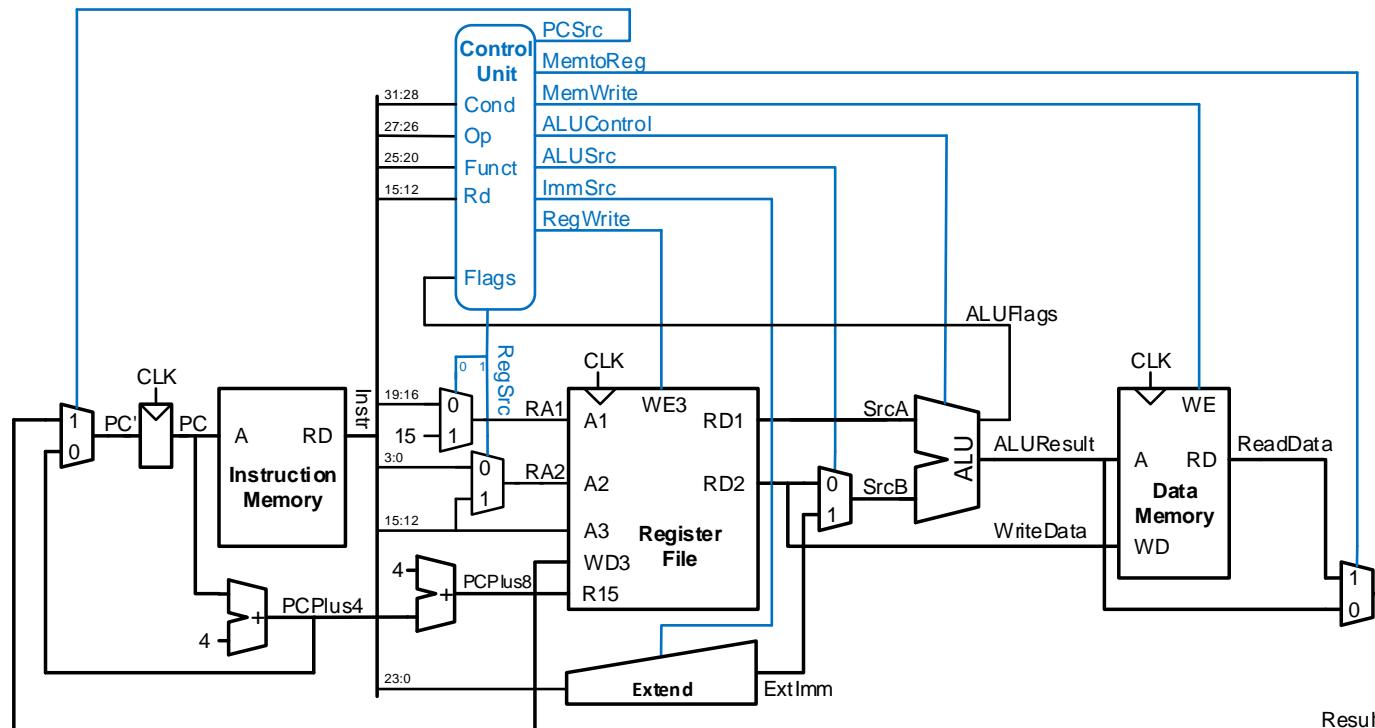
- Main Decoder
- ALU Decoder
- PC Logic



# Single-Cycle Control: PC Logic

$PCS = 1$  if PC is written by an instruction or branch (B):

$$PCS = ((Rd == 15) \& RegW) | Branch$$

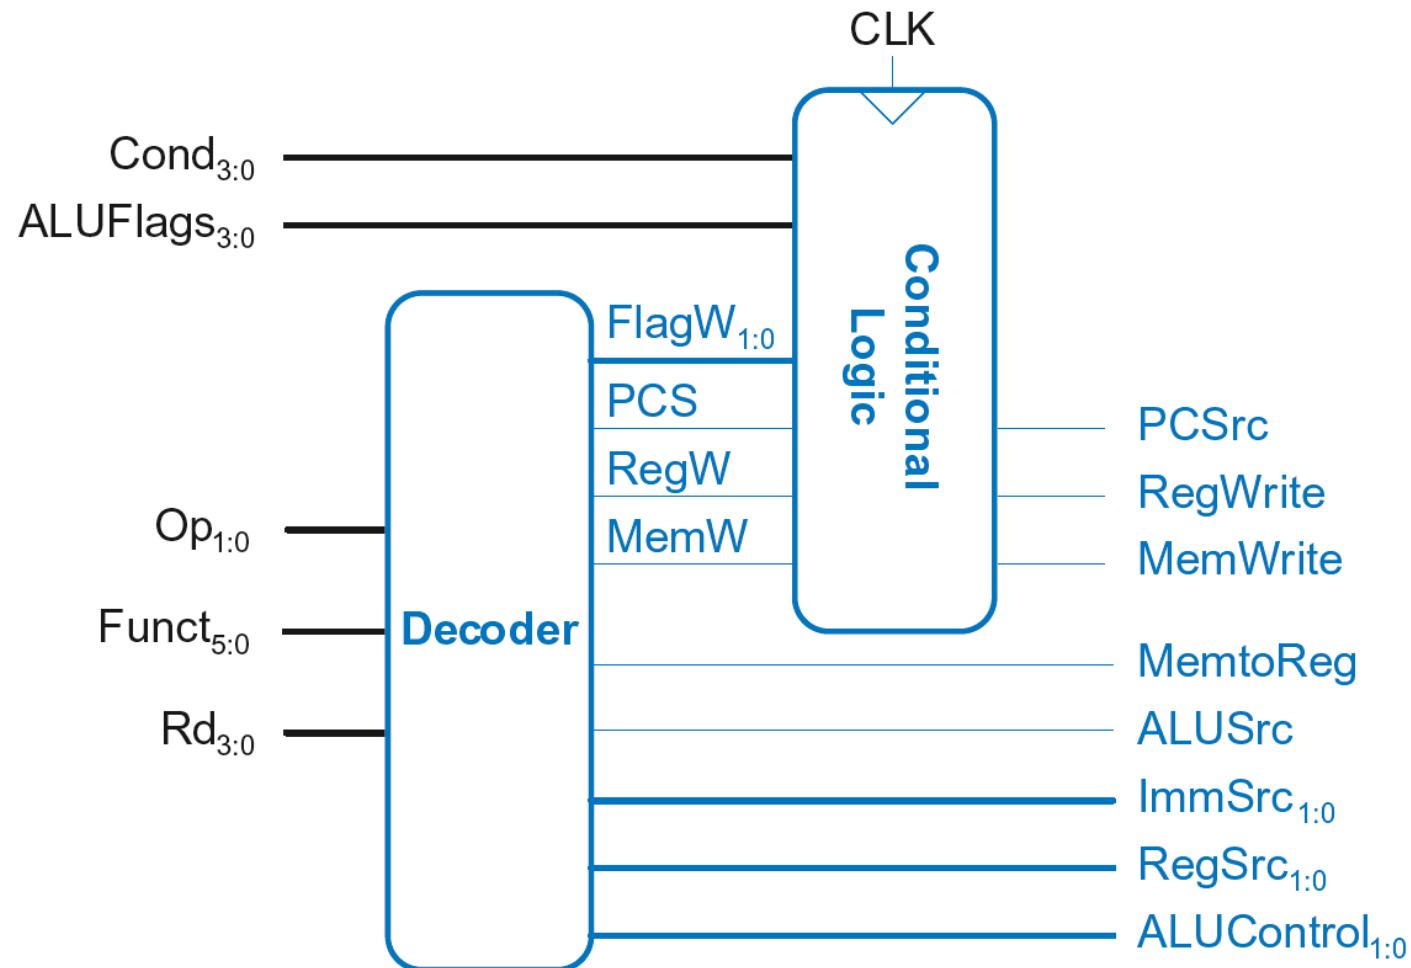


If instruction is executed:  $PCS_{src} = PCS$

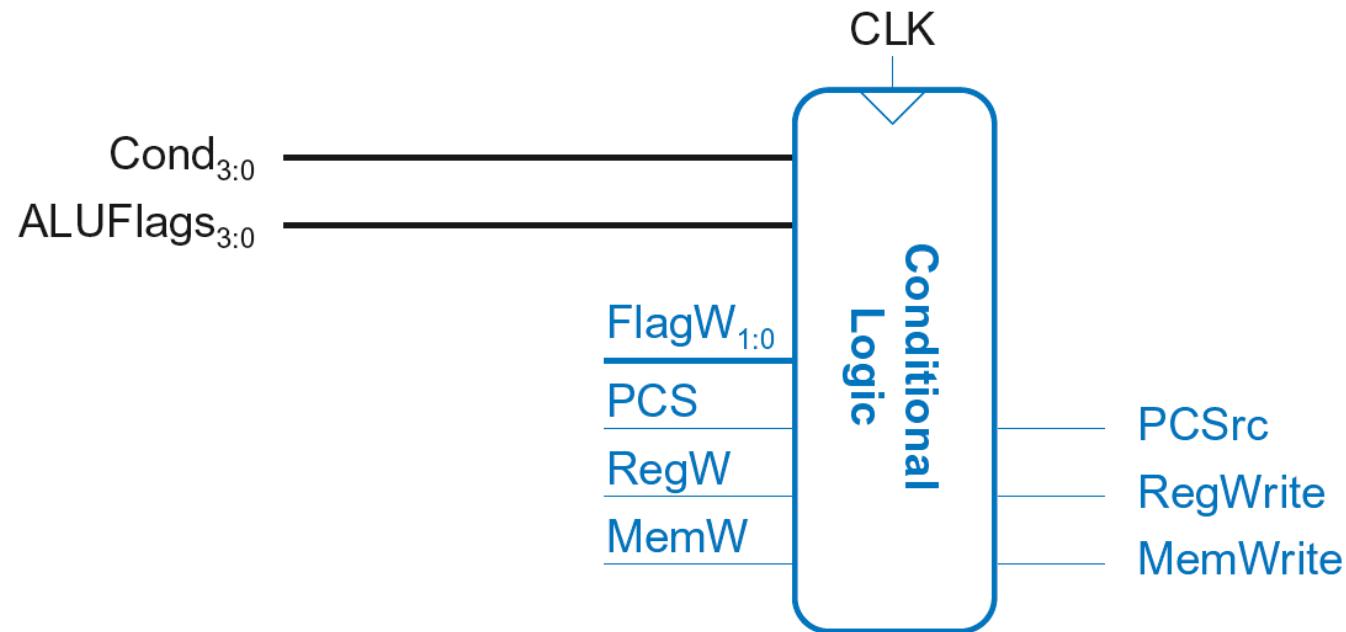
Else  $PCS_{src} = 0$  (i.e.,  $PC = PC + 4$ )



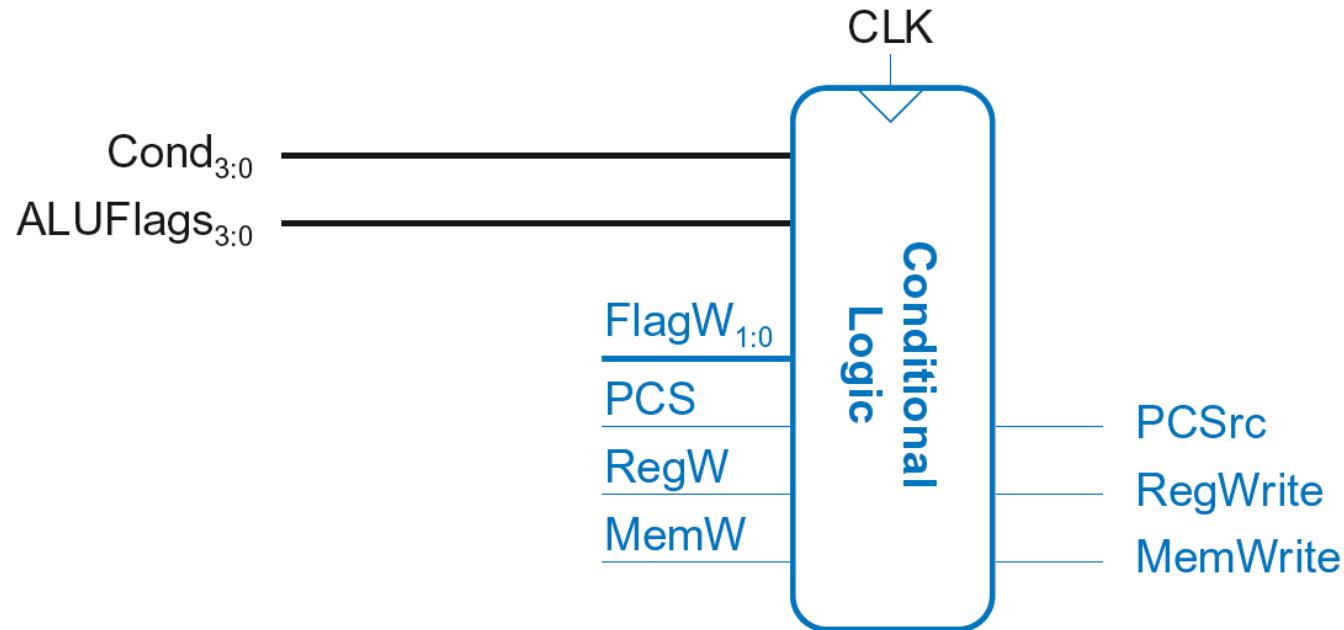
# Single-Cycle Control



# Single-Cycle Control: Cond. Logic



# Conditional Logic

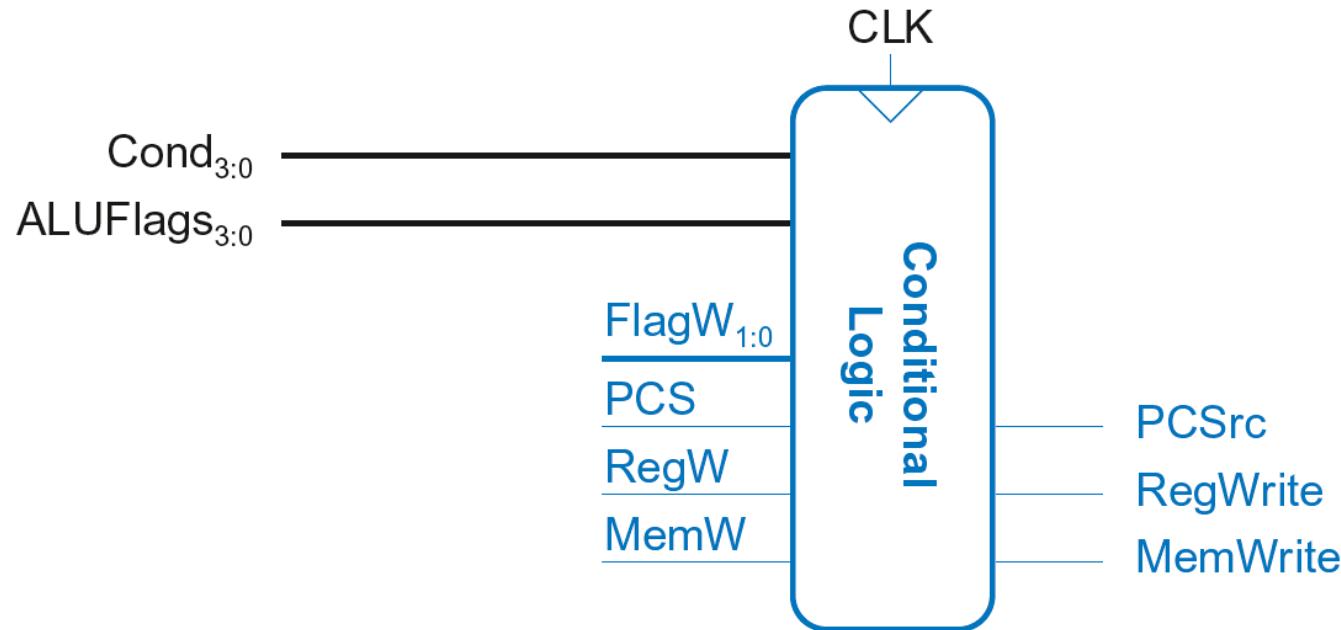


## Function:

1. Check if instruction should execute (if not, force PCSrc, RegWrite, and MemWrite to 0)
2. Possibly update Status Register (Flags<sub>3:0</sub>)



# Conditional Logic

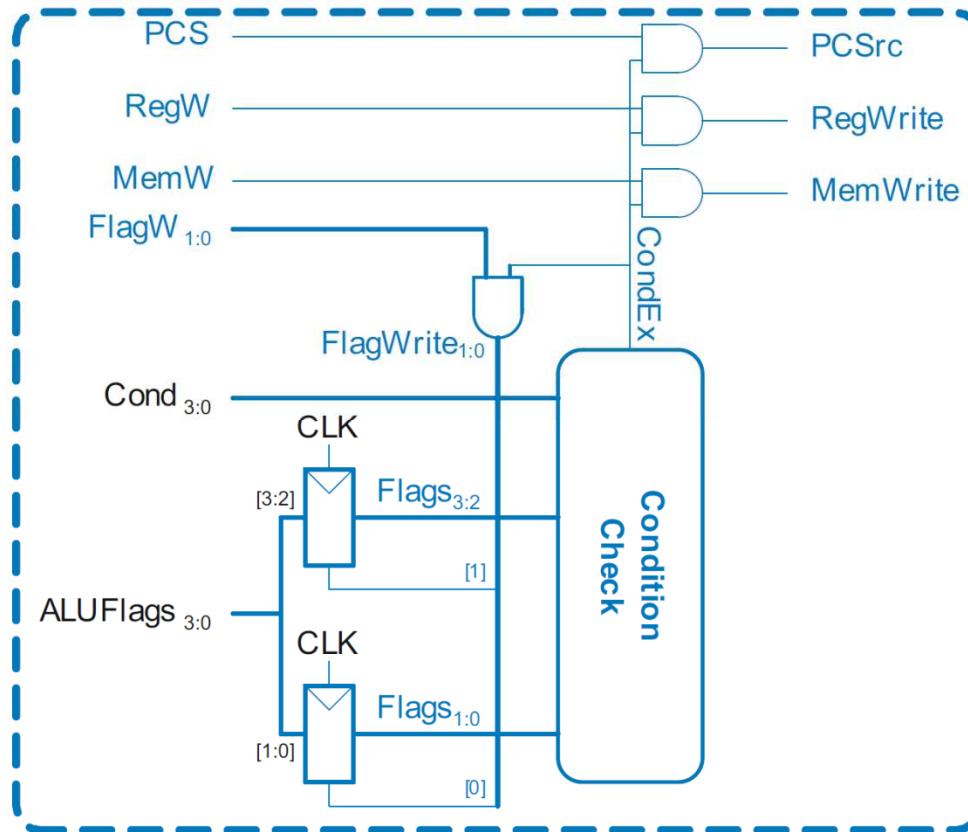


## Function:

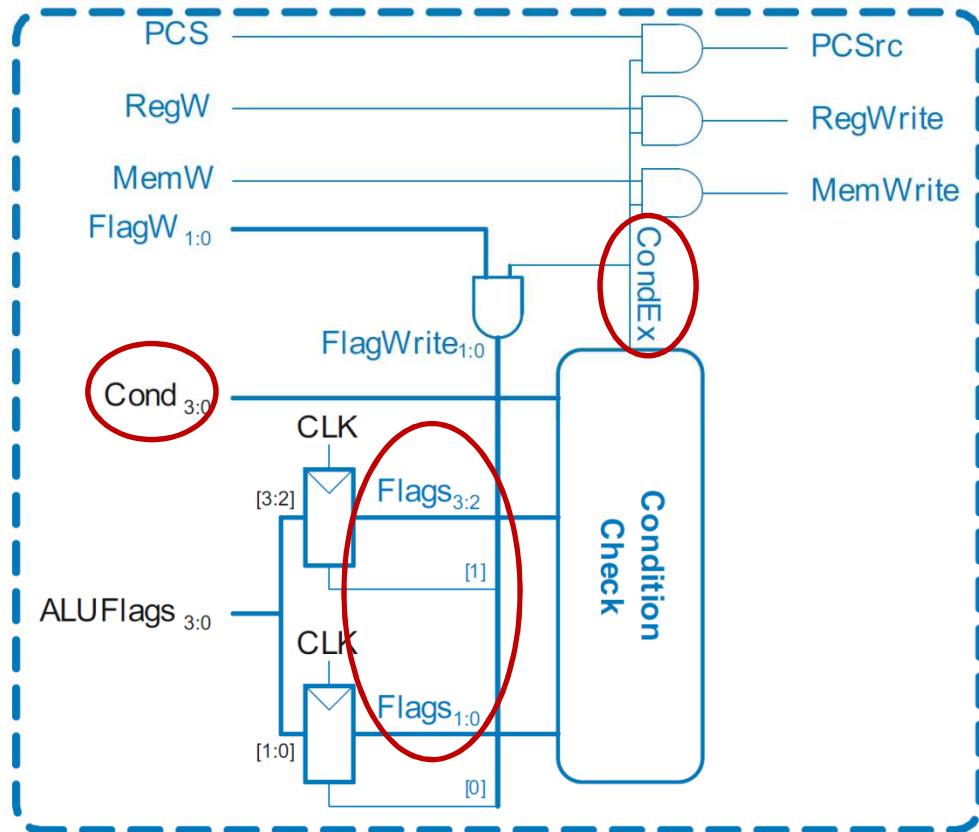
1. Check if instruction should execute (if not, force PCSrc, RegWrite, and MemWrite to 0)
2. Possibly update Status Register (Flags<sub>3:0</sub>)



# Single-Cycle Control: Conditional Logic



# Conditional Logic: Conditional Execution

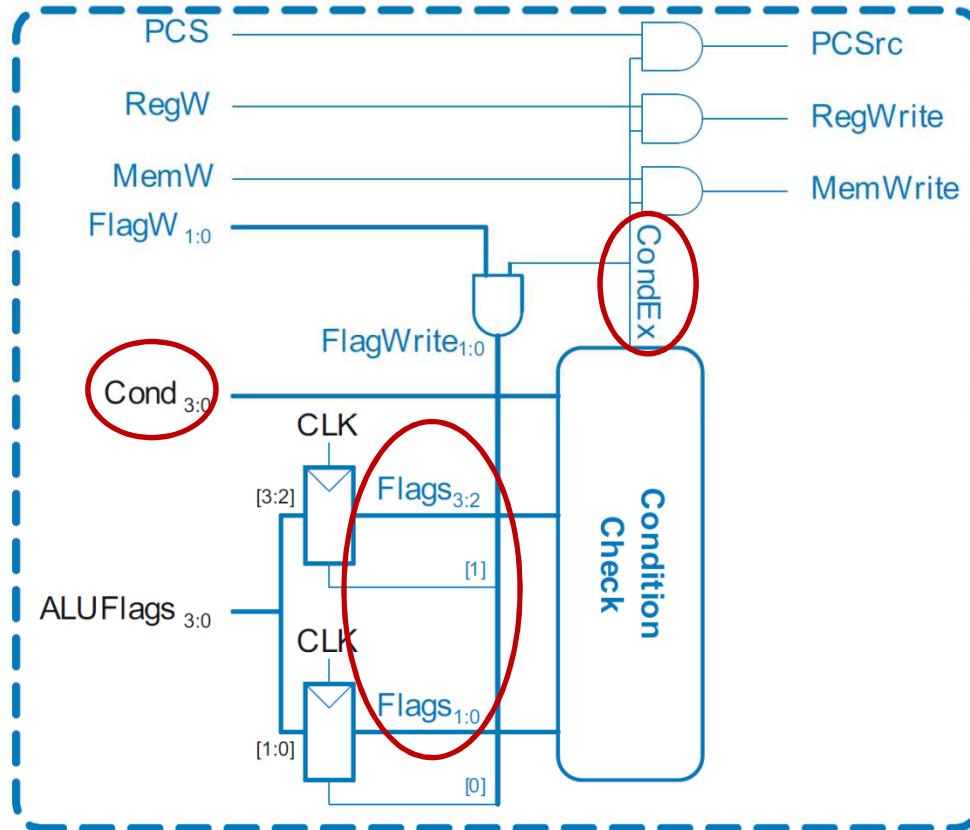


Depending on condition mnemonic ( $Cond_{3:0}$ ) and condition flags ( $Flags_{3:0}$ ) the instruction is executed ( $CondEx = 1$ )



# Conditional Logic: Conditional Execution

**Flags<sub>3:0</sub>** is the status register



Depending on condition mnemonic (**Cond<sub>3:0</sub>**) and condition flags (**Flags<sub>3:0</sub>**) the instruction is executed (**CondEx = 1**)



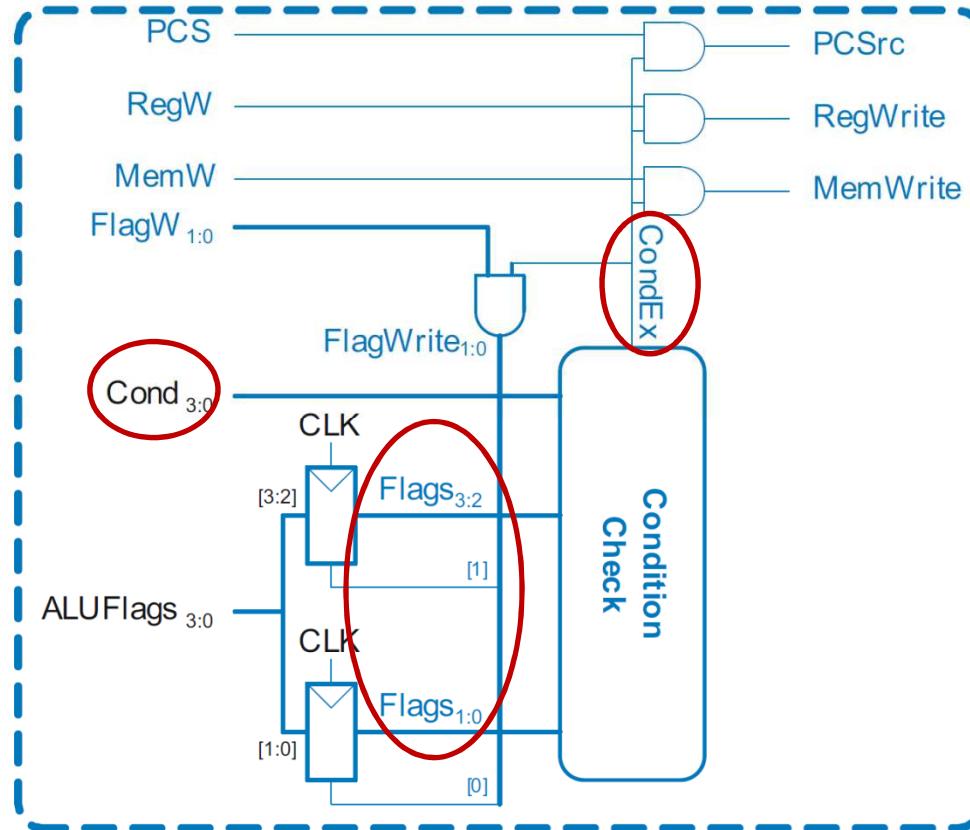
# Review: Condition Mnemonics

Cond <sub>3:0</sub>	Mnemonic	Name	CondEx
0000	EQ	Equal	$Z$
0001	NE	Not equal	$\bar{Z}$
0010	CS / HS	Carry set / Unsigned higher or same	$C$
0011	CC / LO	Carry clear / Unsigned lower	$\bar{C}$
0100	MI	Minus / Negative	$N$
0101	PL	Plus / Positive of zero	$\bar{N}$
0110	VS	Overflow / Overflow set	$V$
0111	VC	No overflow / Overflow clear	$\bar{V}$
1000	HI	Unsigned higher	$\bar{Z}C$
1001	LS	Unsigned lower or same	$Z \text{ OR } \bar{C}$
1010	GE	Signed greater than or equal	$\overline{N \oplus V}$
1011	LT	Signed less than	$N \oplus V$
1100	GT	Signed greater than	$\bar{Z}(\overline{N \oplus V})$
1101	LE	Signed less than or equal	$Z \text{ OR } (N \oplus V)$
1110	AL (or none)	Always / unconditional	ignored



# Conditional Logic: Conditional Execution

$\text{Flags}_{3:0} = \text{NZCV}$



Example:

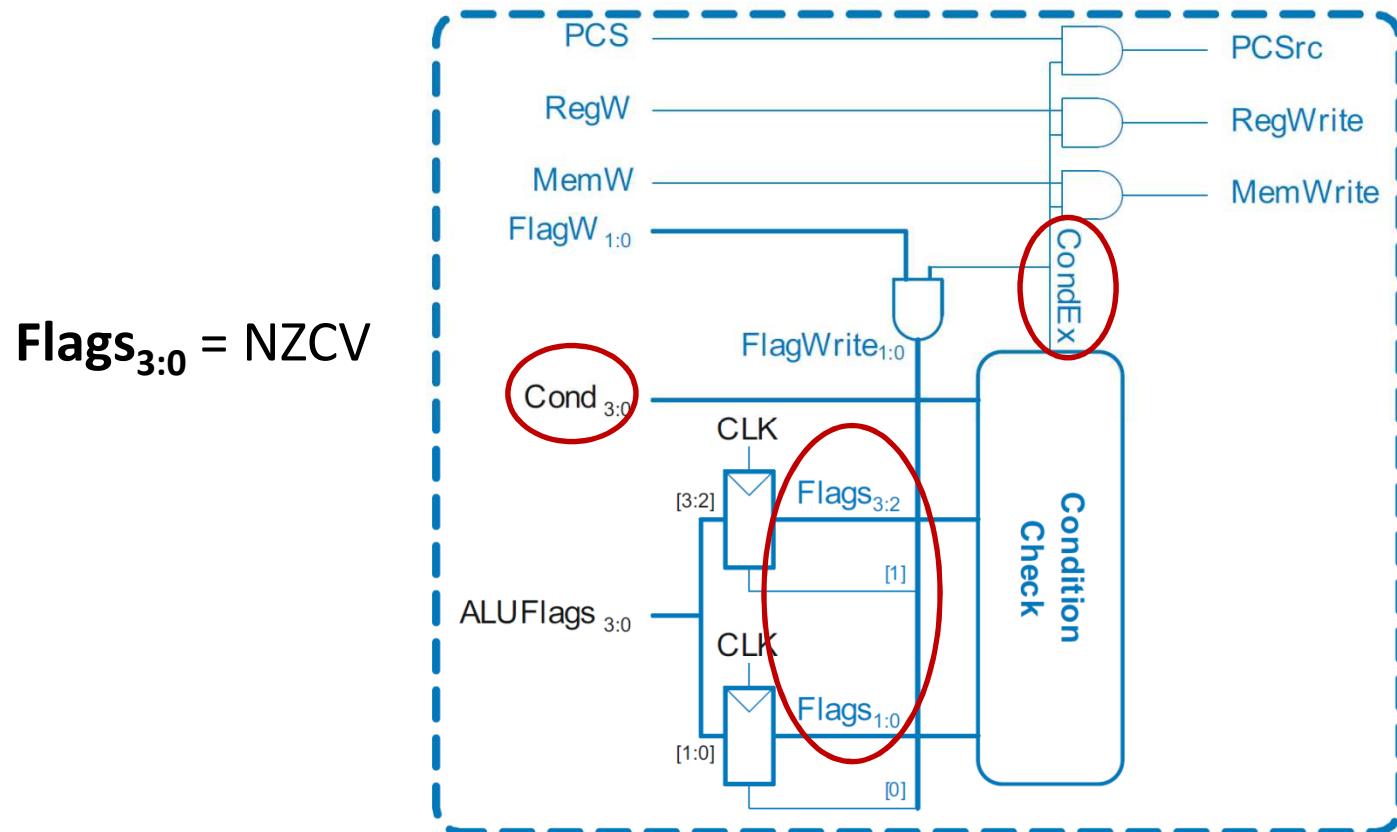
AND R1, R2, R3

$\text{Cond}_{3:0} = 1110$  (unconditional)

$\Rightarrow \text{CondEx} = 1$



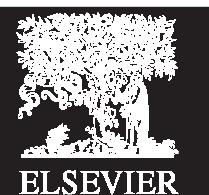
# Conditional Logic: Conditional Execution



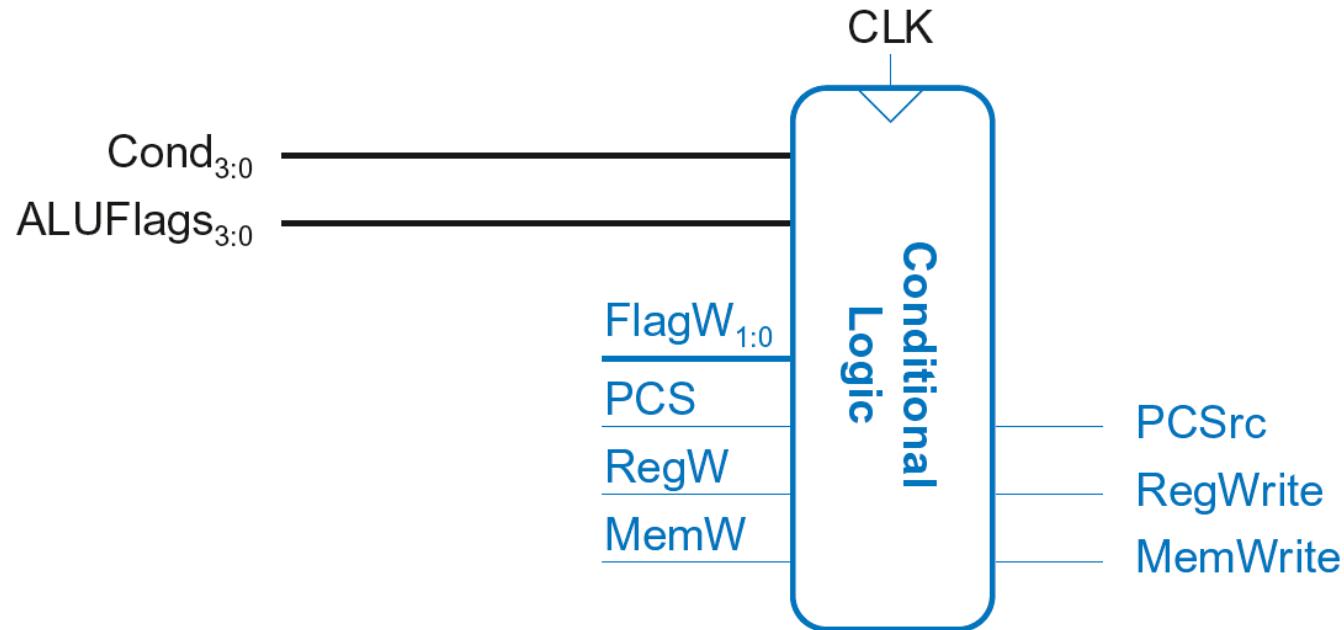
## Example:

EOREQ R5, R6, R7

**Cond<sub>3:0</sub>**=0000 (EQ):      if **Flags<sub>3:2</sub>**=0100 => **CondEx = 1**



# Conditional Logic

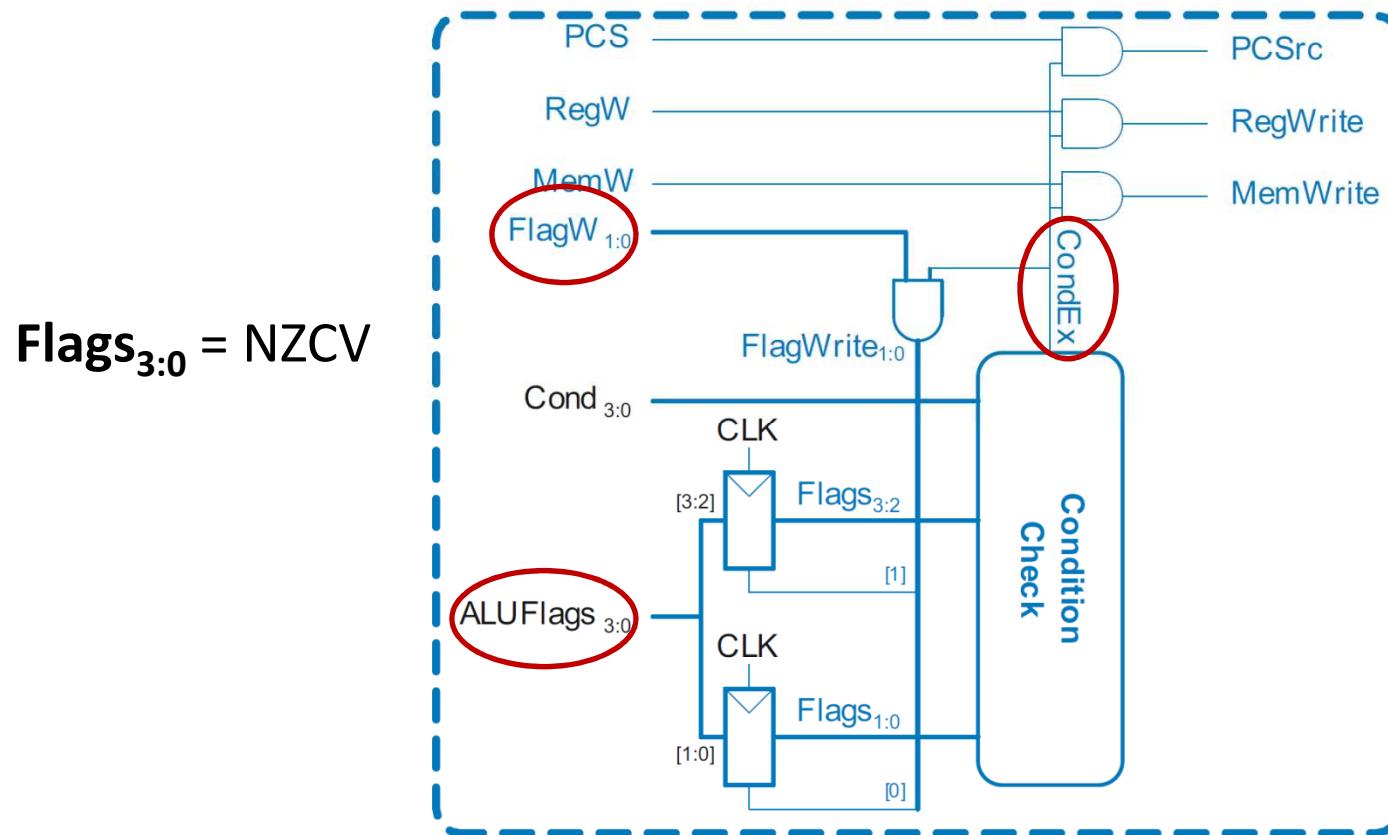


## Function:

1. Check if instruction should execute (if not, force PCSrc, RegWrite, and MemWrite to 0)
2. **Possibly update Status Register (Flags<sub>3:0</sub>)**



# Conditional Logic: Update (Set) Flags

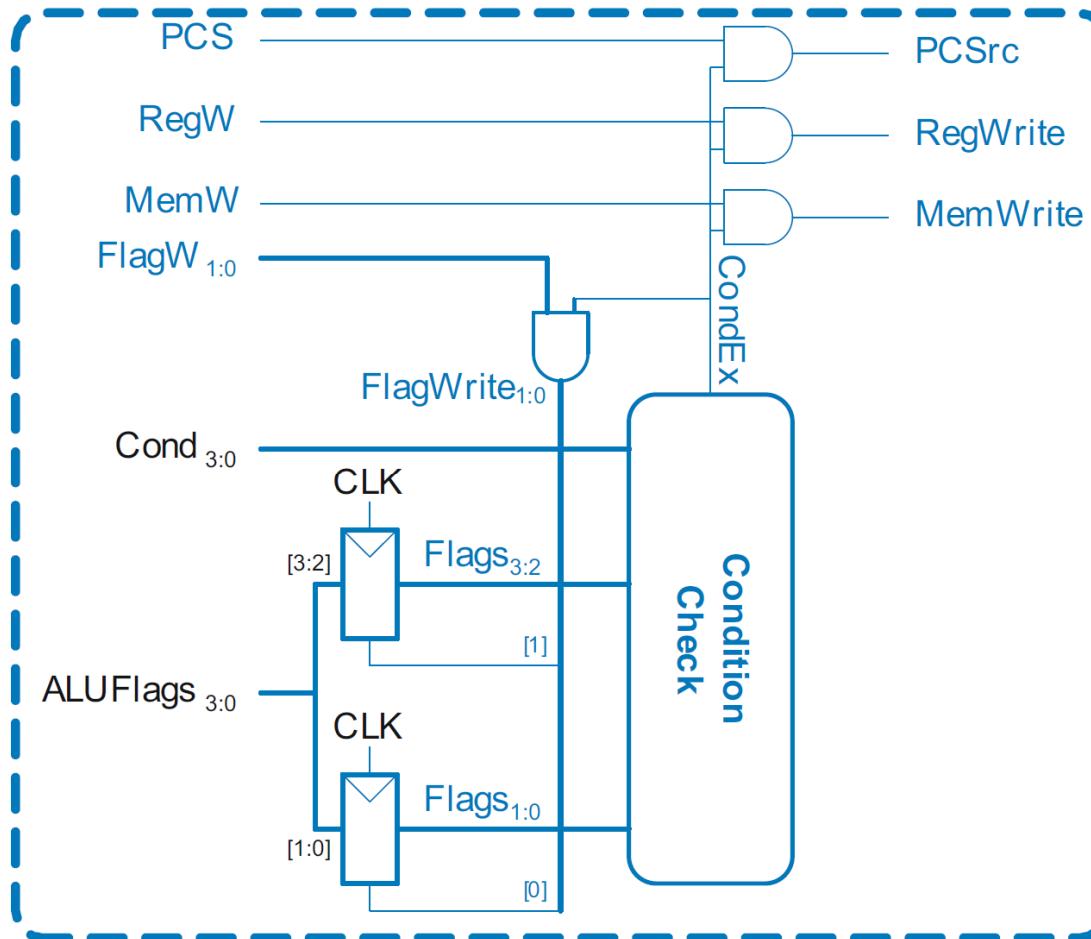


Flags<sub>3:0</sub> updated (with ALUFlags<sub>3:0</sub>) if:

- FlagW is 1 (i.e., the instruction's S-bit is 1) AND
- CondEx is 1 (the instruction should be executed)



# Conditional Logic: Update (Set) Flags



## Recall:

- ADD, SUB update **all Flags**
- AND, OR update **NZ only**
- So Flags status register has two write enables:  
**FlagW<sub>1:0</sub>**



# Review: ALU Decoder

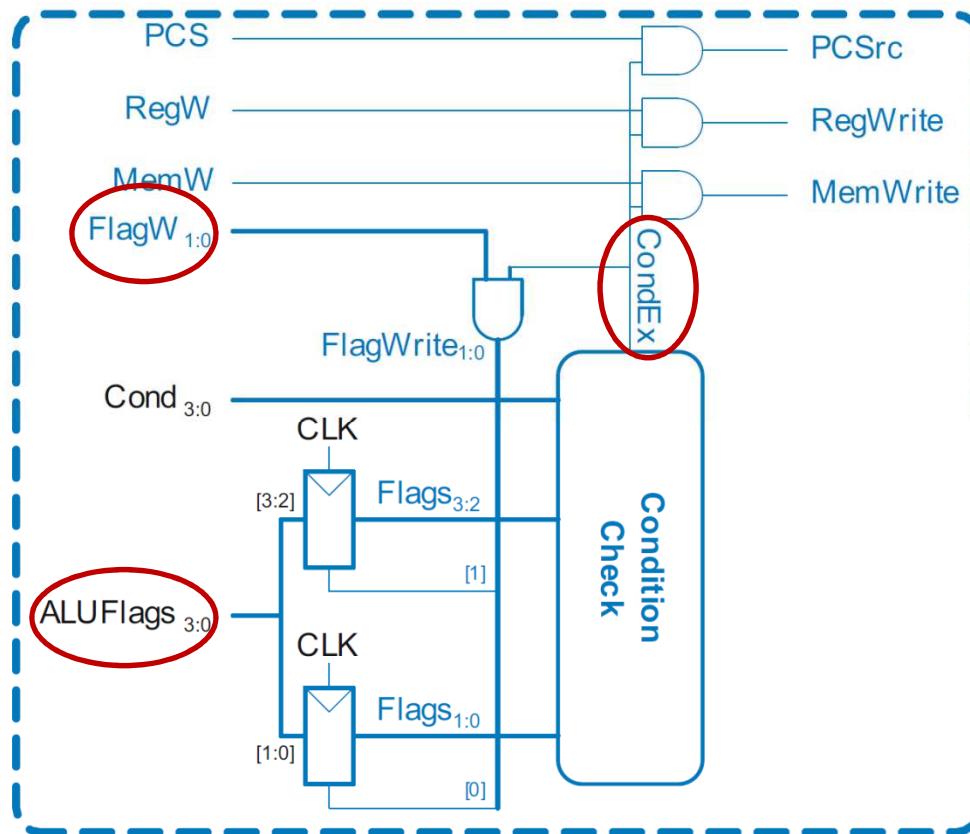
ALUOp	Funct <sub>4:1</sub> (cmd)	Funct <sub>0</sub> (S)	Type	ALUControl <sub>1:0</sub>	FlagW <sub>1:0</sub>
0	X	X	Not DP	00	00
1	0100	0	ADD	00	00
		1			11
	0010	0	SUB	01	00
		1			11
	0000	0	AND	10	00
		1			10
	1100	0	ORR	11	00
		1			10

- $\text{FlagW}_1 = 1$ : NZ(Flags<sub>3:2</sub>) should be saved
- $\text{FlagW}_0 = 1$ : CV(Flags<sub>1:0</sub>) should be saved



# Conditional Logic: Update (Set) Flags

All Flags updated



**Example:** SUBS R5, R6, R7

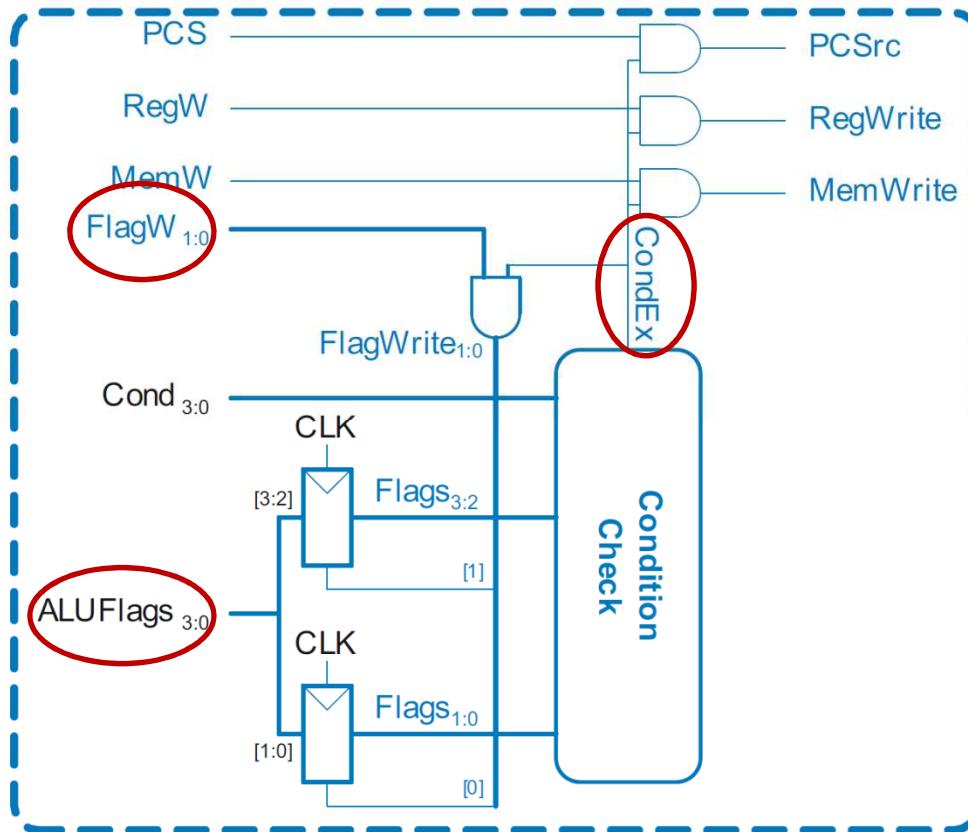
$\text{FlagW}_{1:0} = 11$  AND  $\text{CondEx} = 1$  (unconditional)  $\Rightarrow \text{FlagWrite}_{1:0} = 11$



# Conditional Logic: Update (Set) Flags

**Flags<sub>3:0</sub> = NZCV**

- Only **Flags<sub>3:2</sub>** updated
- i.e., only **NZ** Flags updated



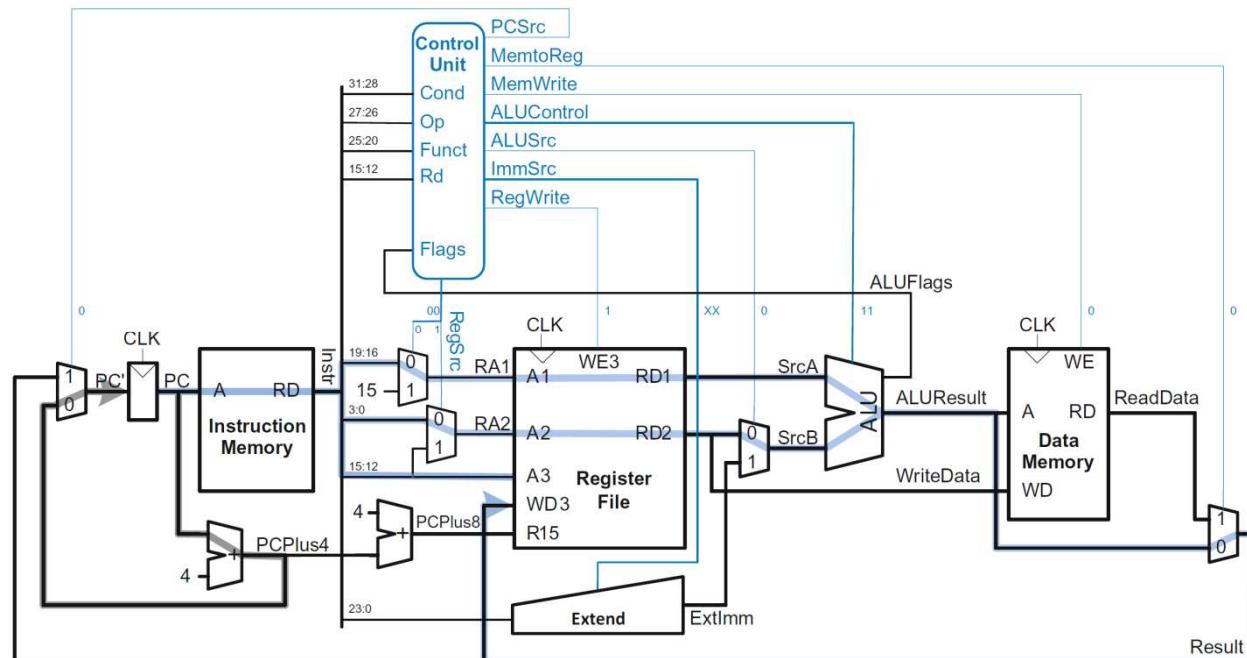
**Example:** ANDS R7, R1, R3

**FlagW<sub>1:0</sub> = 10 AND CondEx = 1 (unconditional) => FlagWrite<sub>1:0</sub> = 10**

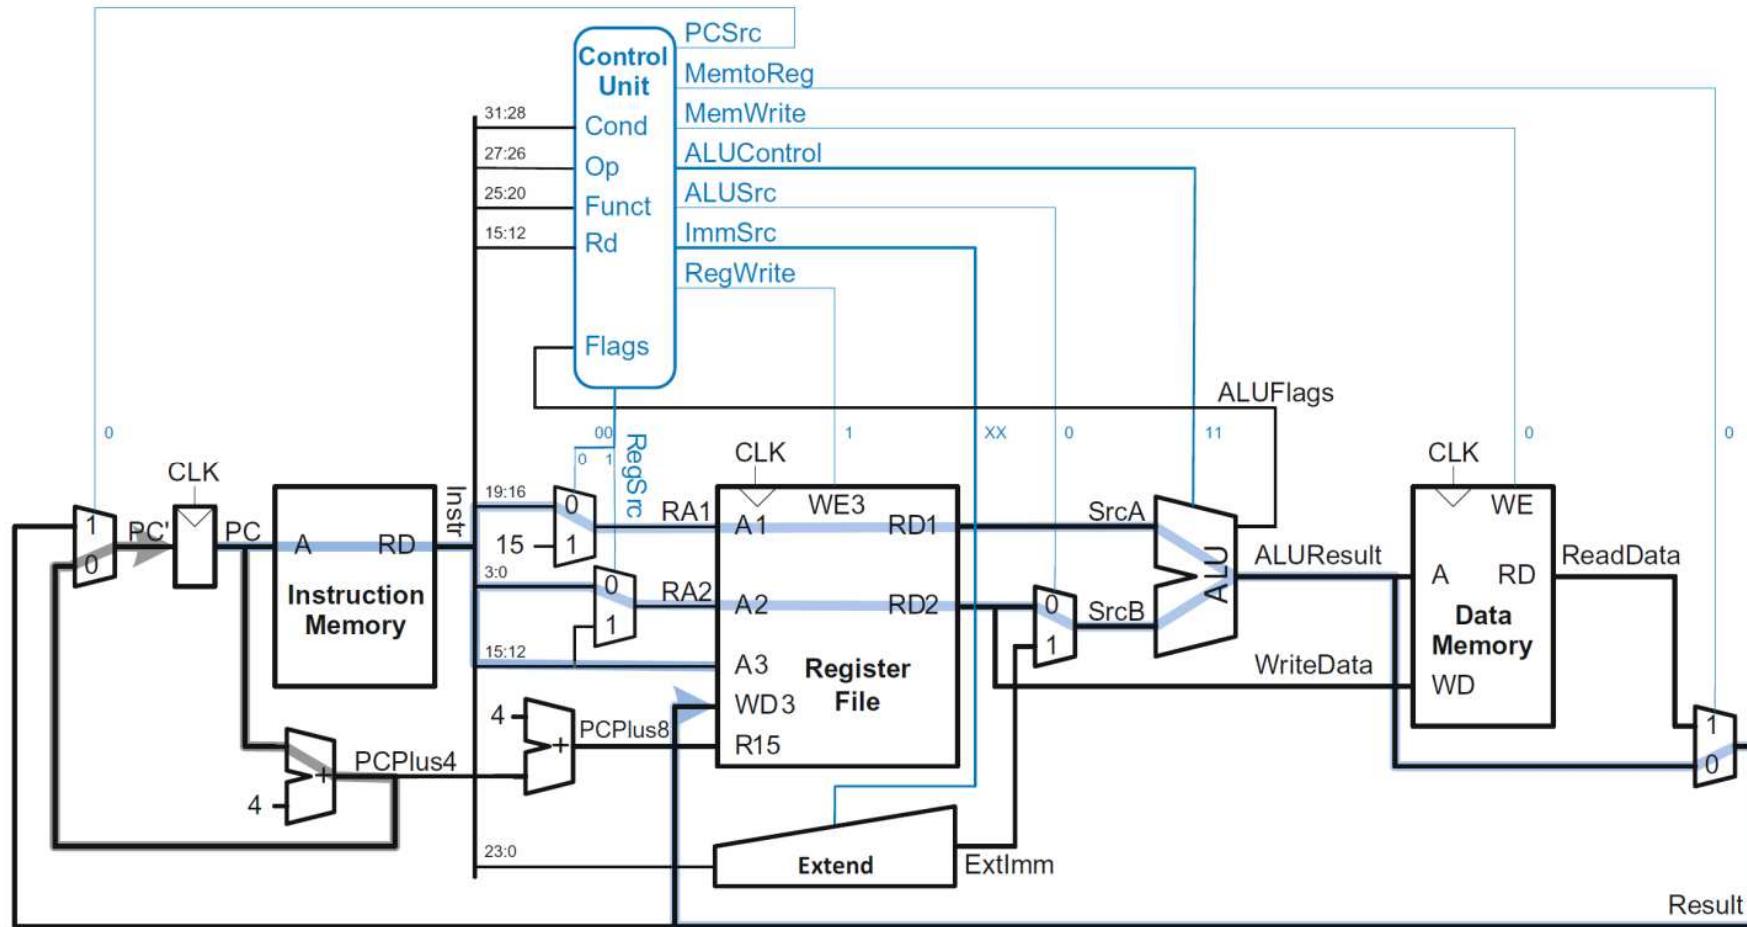


# Example: ORR

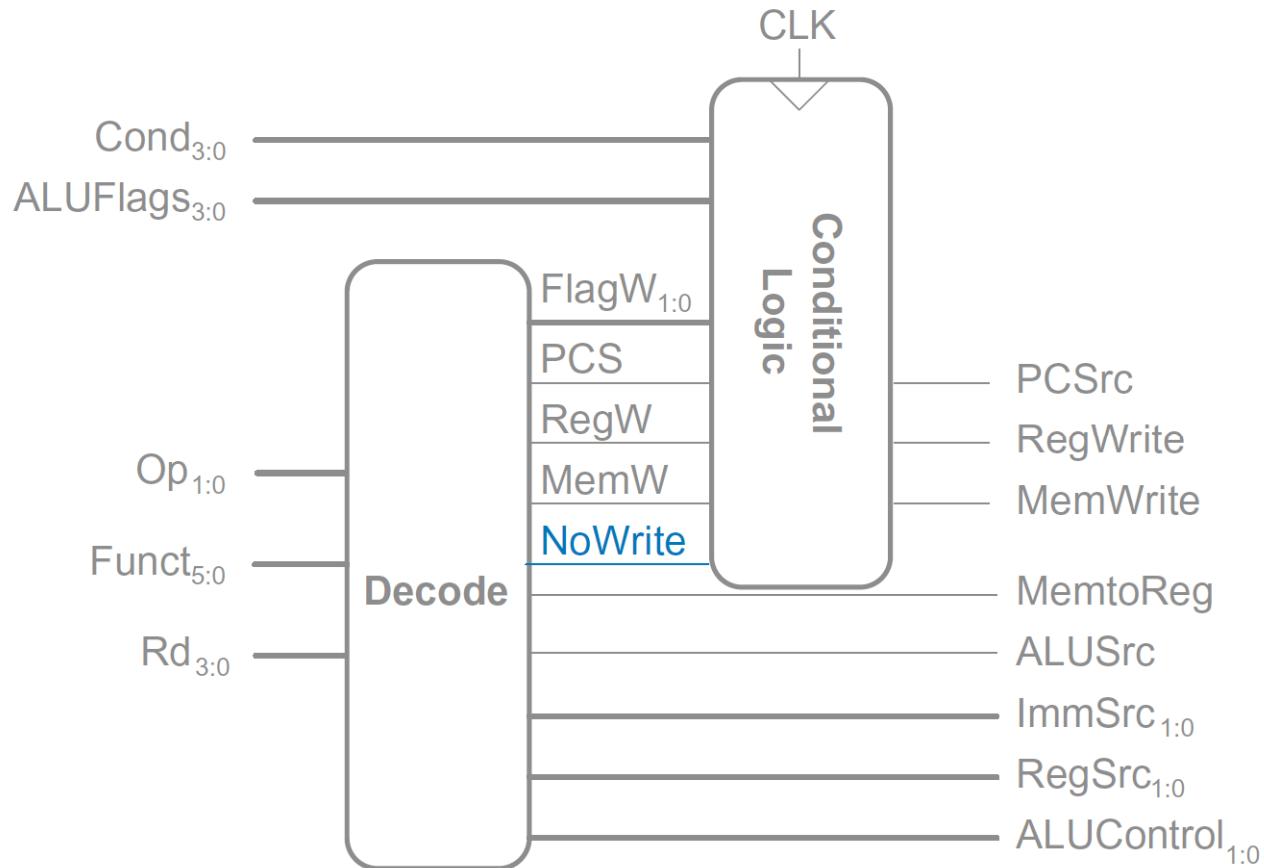
Op	Funct <sub>5</sub>	Funct <sub>0</sub>	Type	Branch	MemtoReg	MemW	ALUSrc	ImmSrc	RegW	RegSrc	ALUOp
00	0	X	DP Reg	0	0	0	0	XX	1	00	1



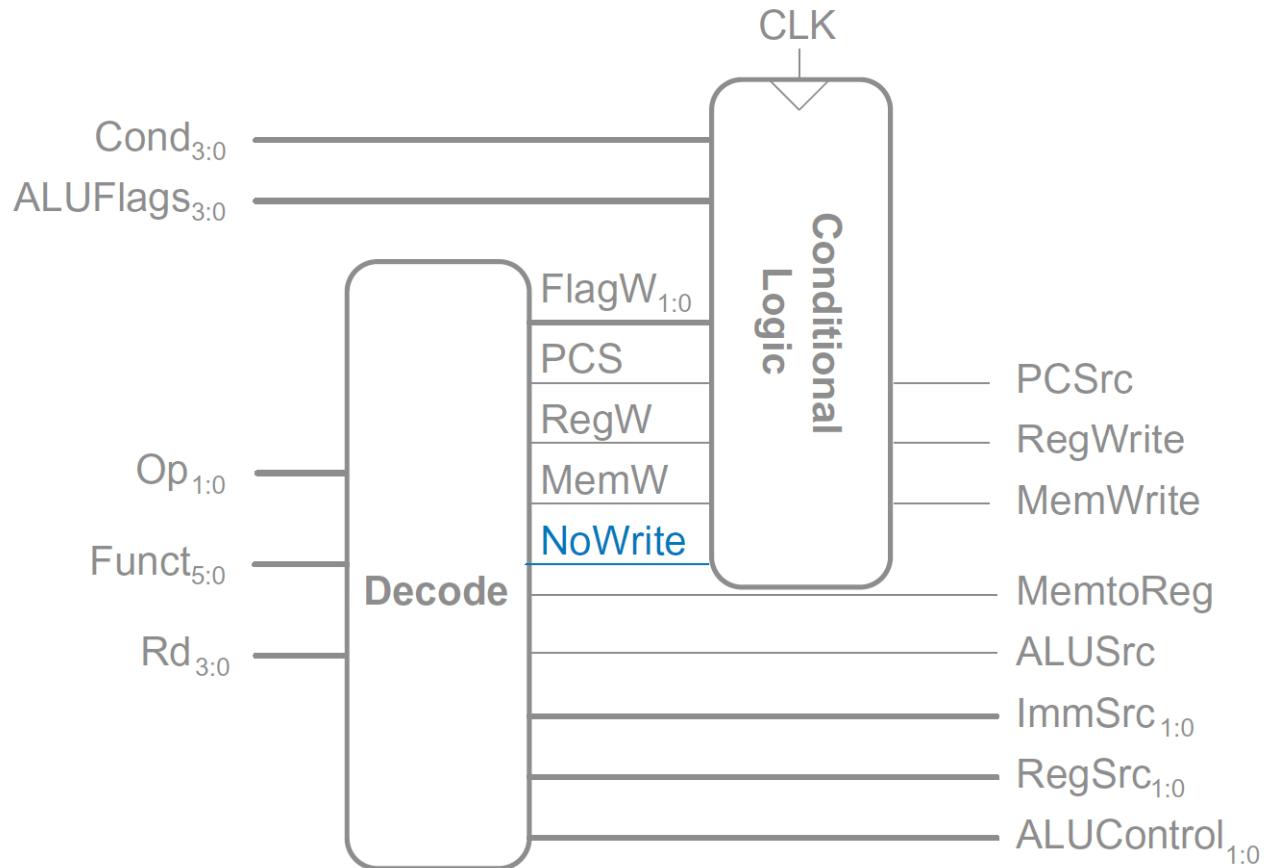
# Example: ORR



# Extended Functionality: CMP



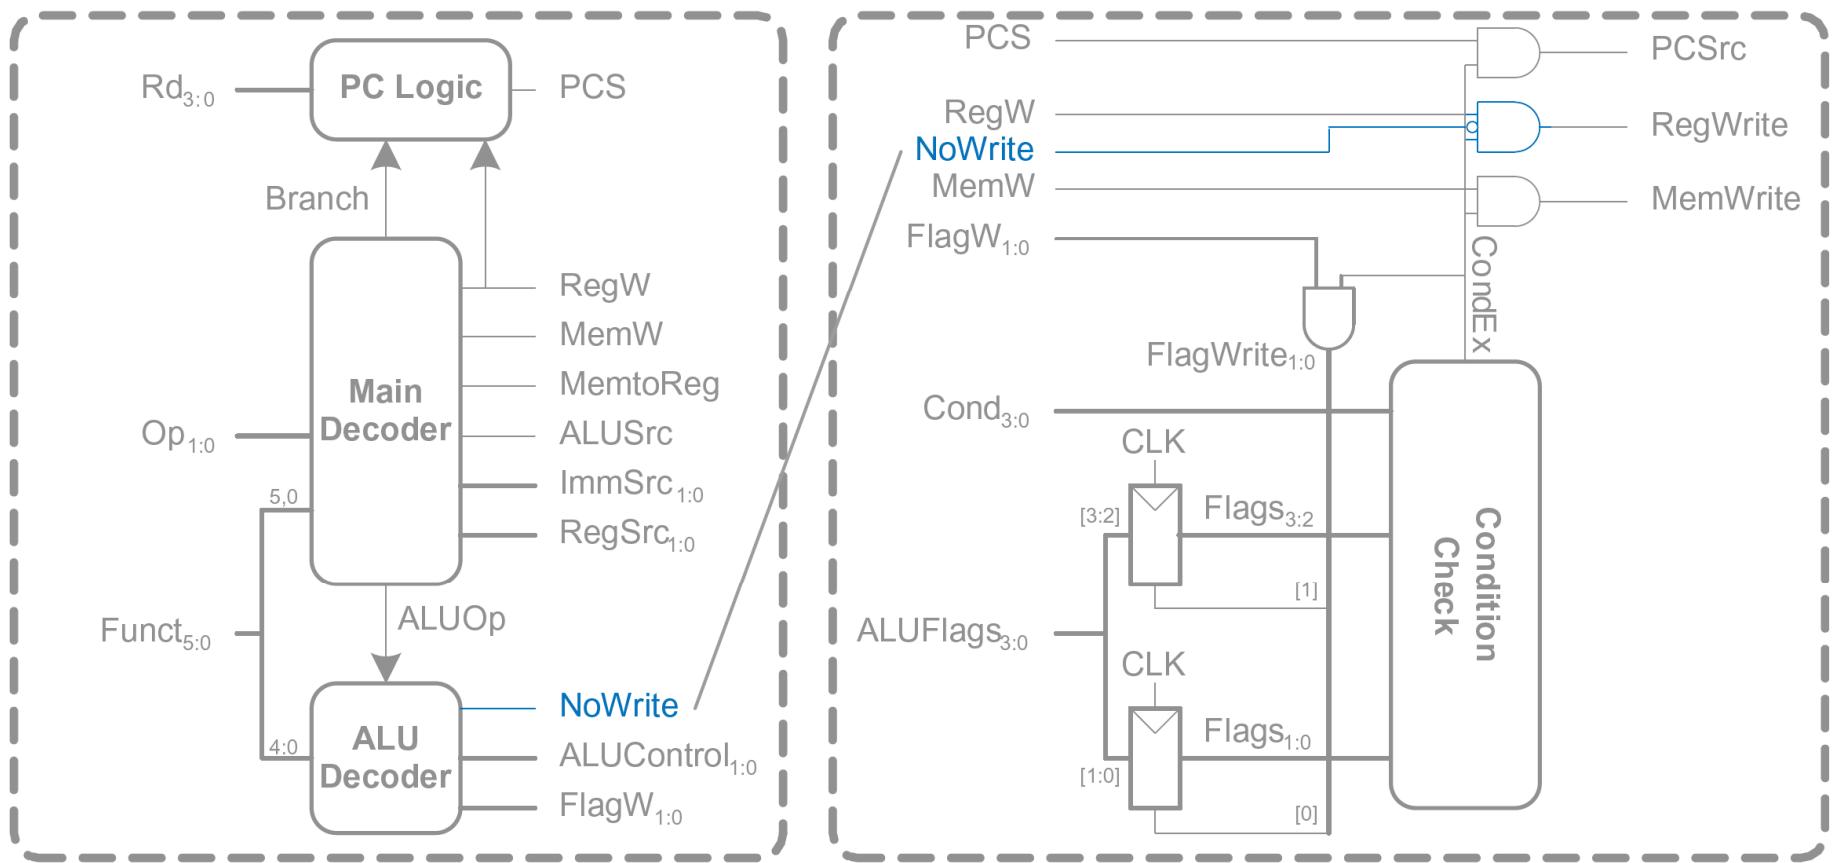
# Extended Functionality: CMP



**No change to datapath**



# Extended Functionality: CMP

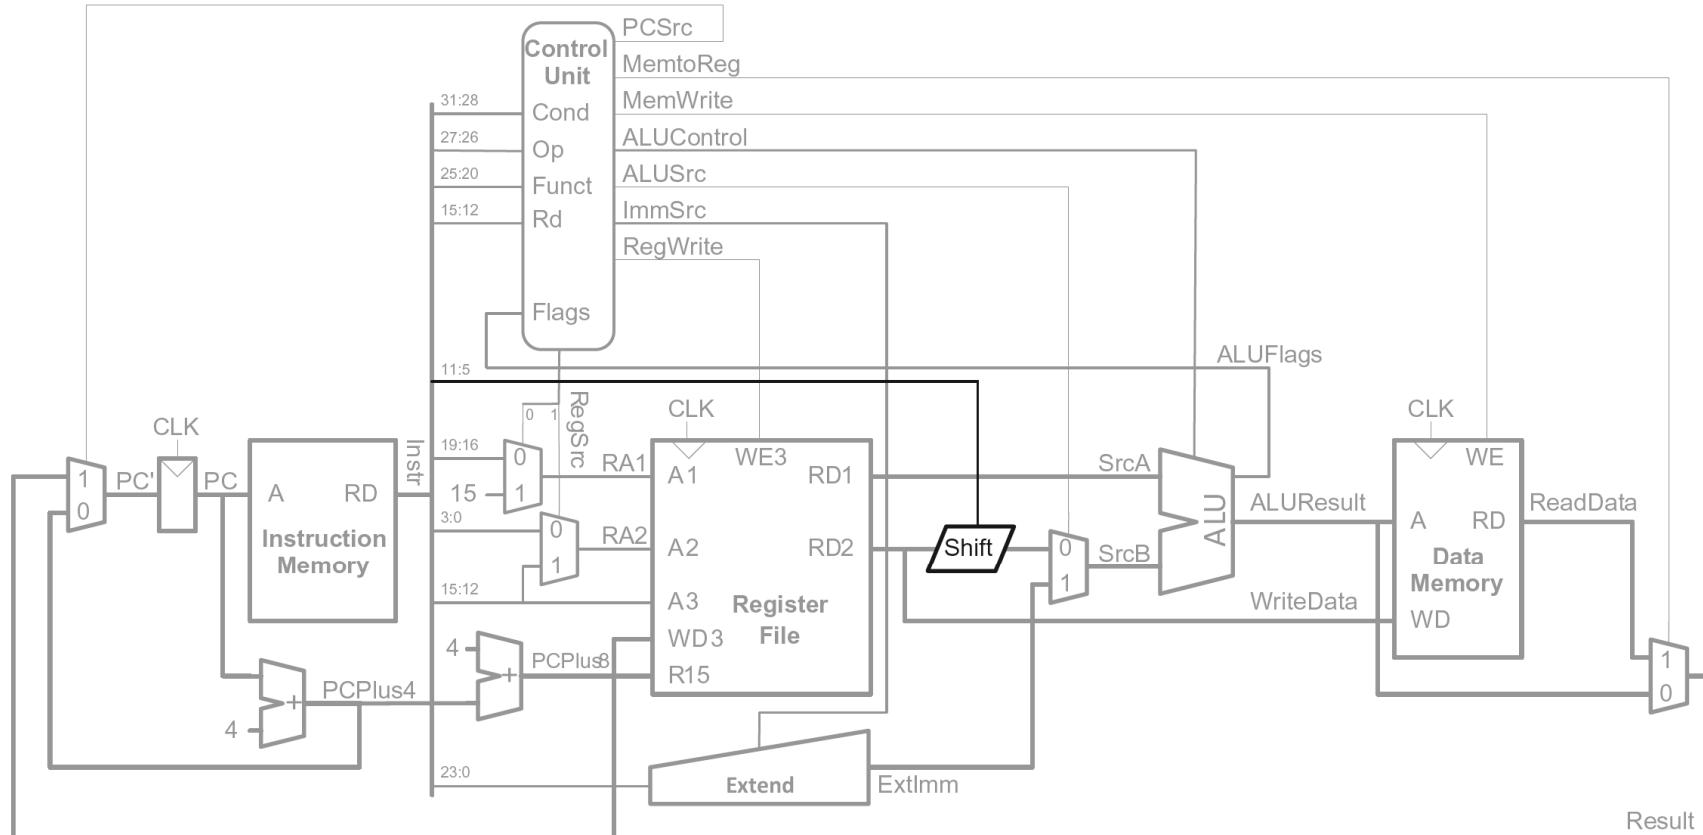


# Extended Functionality: CMP

<b>ALUOp</b>	<b>Funct<sub>4:1</sub> (cmd)</b>	<b>Funct<sub>0</sub> (S)</b>	<b>Type</b>	<b>ALUControl<sub>1:0</sub></b>	<b>FlagW<sub>1:0</sub></b>	<b>NoWrite</b>
0	X	X	Not DP	00	00	0
1	0100	0	ADD	00	00	0
		1			11	0
	0010	0	SUB	01	00	0
		1			11	0
	0000	0	AND	10	00	0
		1			10	0
	1100	0	ORR	11	00	0
		1			10	0
	<b>1010</b>	<b>1</b>	<b>CMP</b>	<b>01</b>	<b>11</b>	<b>1</b>



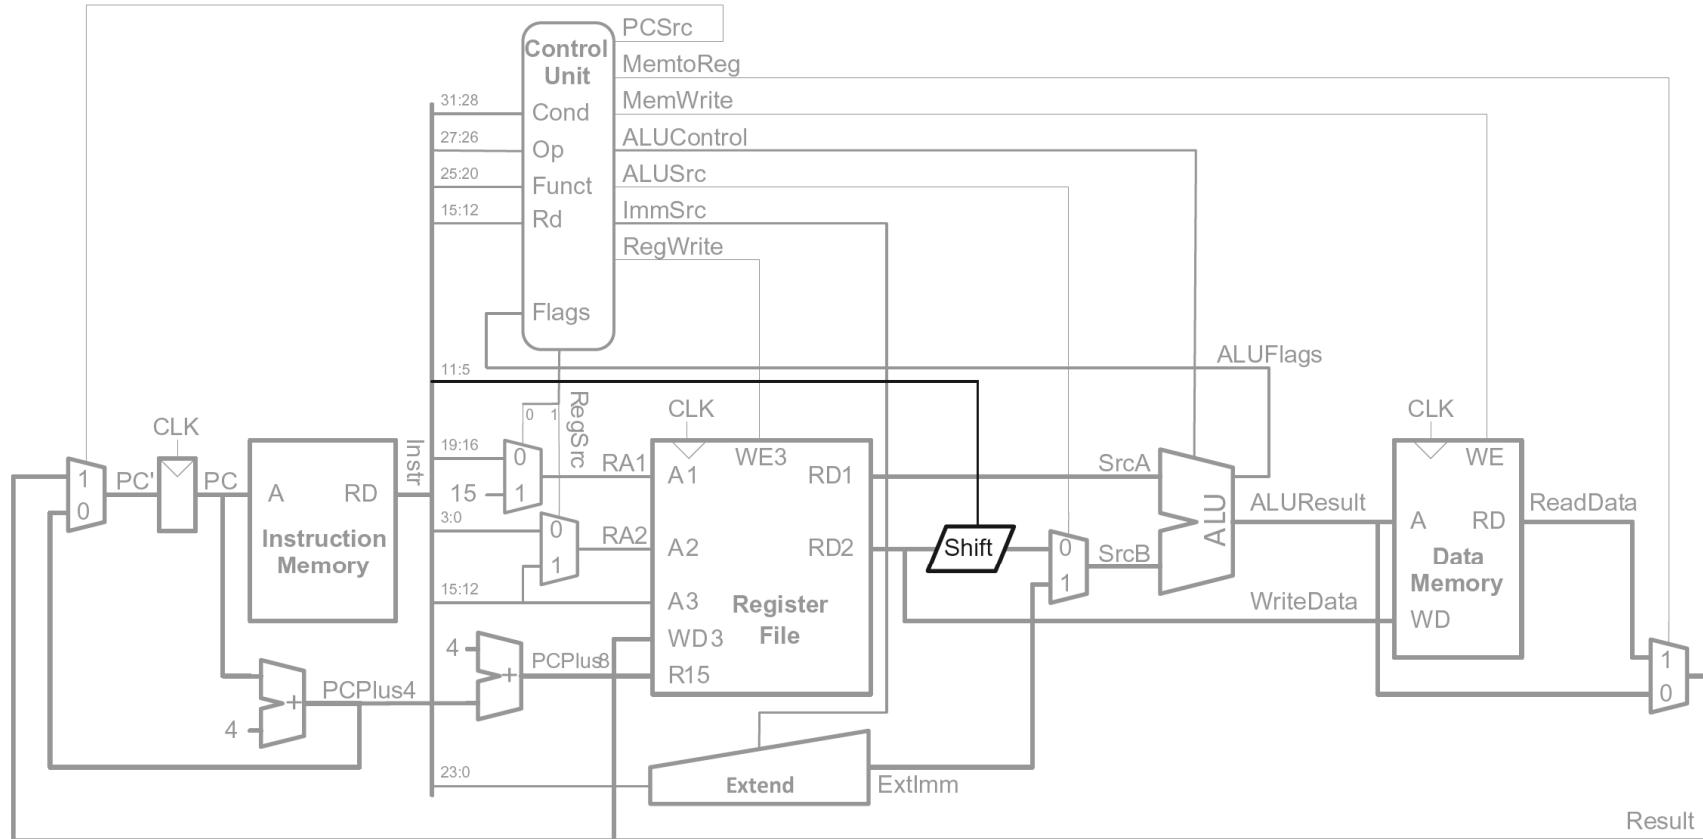
# Extended Functionality: Shifted Register



ADD R7, R2, R12, LSR #5	31:28	27:26	25	24:21	20	19:16	15:12	11:7	6:5	4	3:0
	cond	op	I	cmd	S	rn	rd	shamt5	sh	rm	
	14	0	0	4	0	2	7	5	01 <sub>2</sub>	0	12



# Extended Functionality: Shifted Register



No change to controller

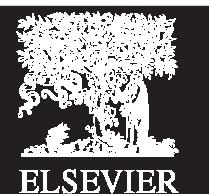


# Review: Processor Performance

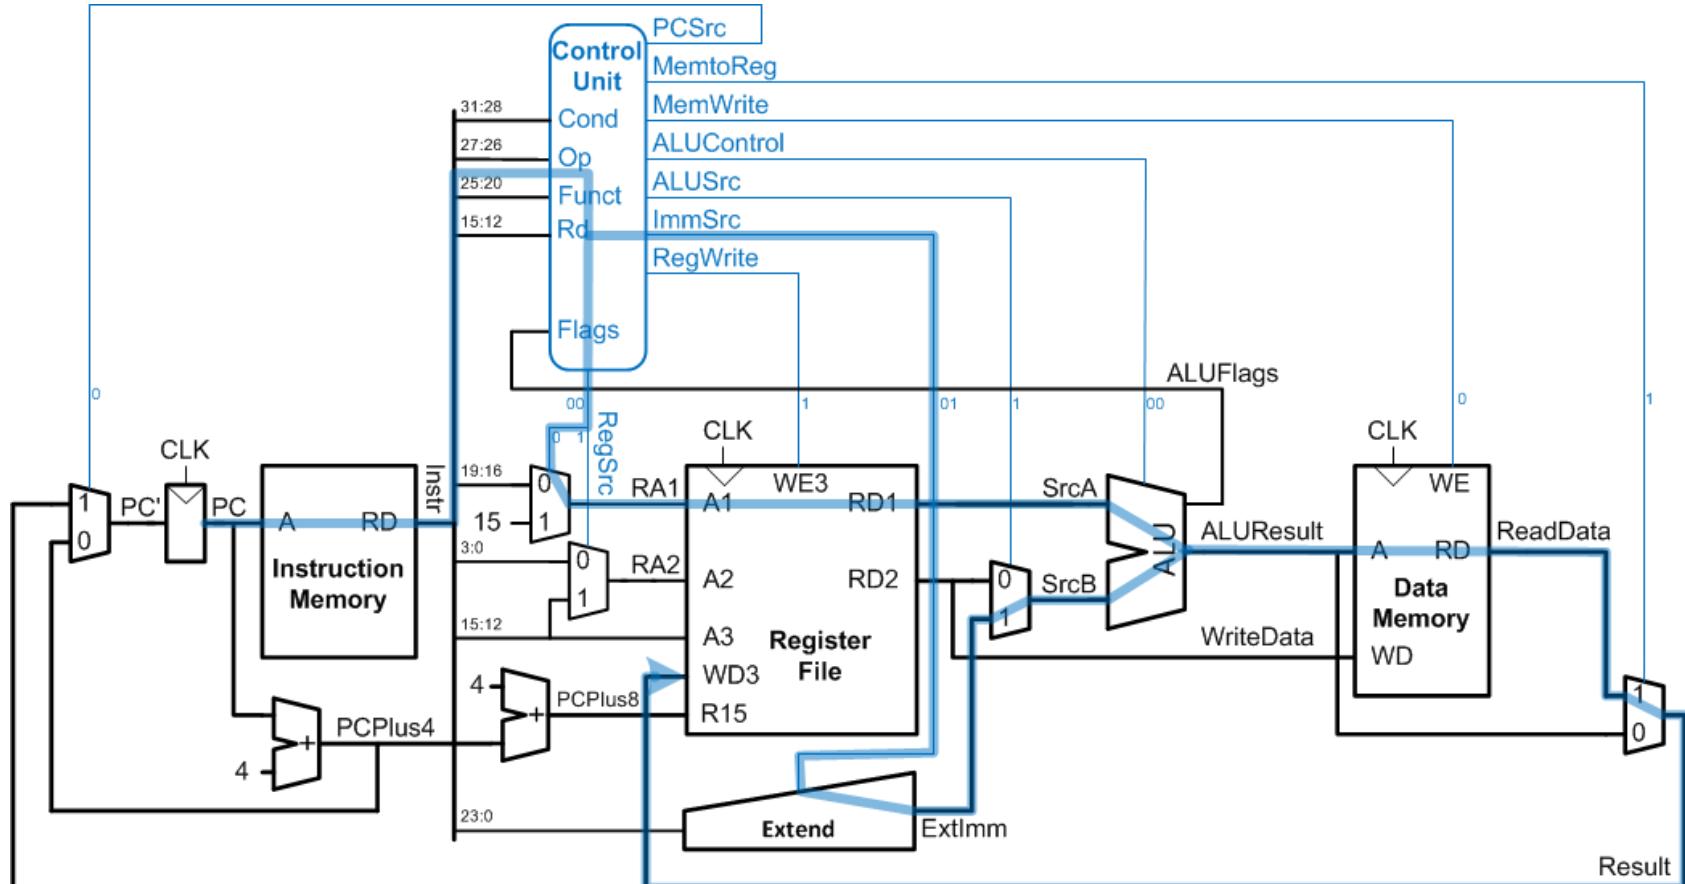
## Program Execution Time

$$= (\# \text{instructions})(\text{cycles/instruction})(\text{seconds/cycle})$$

$$= \# \text{ instructions} \times \text{CPI} \times T_C$$



# Single-Cycle Performance



$T_C$  limited by critical path (LDR)



# Single-Cycle Performance

- **Single-cycle critical path:**

$$T_{c1} = t_{pcq\_PC} + t_{mem} + t_{dec} + \max[t_{mux} + t_{RFread}, t_{sext} + t_{mux}] + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

- **Typically, limiting paths are:**
  - memory, ALU, register file
  - $T_{c1} = t_{pcq\_PC} + 2t_{mem} + t_{dec} + t_{RFread} + t_{ALU} + 2t_{mux} + t_{RFsetup}$



# Single-Cycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	40
Register setup	$t_{\text{setup}}$	50
Multiplexer	$t_{\text{mux}}$	25
ALU	$t_{\text{ALU}}$	120
Decoder	$t_{\text{dec}}$	70
Memory read	$t_{\text{mem}}$	200
Register file read	$t_{RF\text{read}}$	100
Register file setup	$t_{RF\text{setup}}$	60

$$T_{c1} = ?$$



# Single-Cycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	40
Register setup	$t_{\text{setup}}$	50
Multiplexer	$t_{\text{mux}}$	25
ALU	$t_{\text{ALU}}$	120
Decoder	$t_{\text{dec}}$	70
Memory read	$t_{\text{mem}}$	200
Register file read	$t_{RF\text{read}}$	100
Register file setup	$t_{RF\text{setup}}$	60

$$\begin{aligned}T_{c1} &= t_{pcq\_PC} + 2t_{\text{mem}} + t_{\text{dec}} + t_{RF\text{read}} + t_{\text{ALU}} + 2t_{\text{mux}} + t_{RF\text{setup}} \\&= [50 + 2(200) + 70 + 100 + 120 + 2(25) + 60] \text{ ps} \\&= \mathbf{840 \text{ ps}}\end{aligned}$$



# Single-Cycle Performance Example

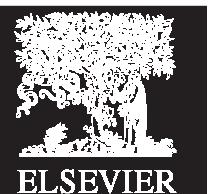
Program with 100 billion instructions:

$$\begin{aligned}\textbf{Execution Time} &= \# \text{ instructions} \times \text{CPI} \times T_C \\ &= (100 \times 10^9)(1)(840 \times 10^{-12} \text{ s}) \\ &= \mathbf{84 \text{ seconds}}\end{aligned}$$



# Multicycle ARM Processor

- **Single-cycle:**
  - + simple
  - cycle time limited by longest instruction (LDR)
  - separate memories for instruction and data
  - 3 adders/ALUs
- **Multicycle processor addresses these issues by breaking instruction into shorter steps**
  - shorter instructions take fewer steps
  - can re-use hardware
  - cycle time is faster



# Multicycle ARM Processor

- **Single-cycle:**

- + simple
- cycle time limited by longest instruction (LDR)
- separate memories for instruction and data
- 3 adders/ALUs

- **Multicycle:**

- + higher clock speed
- + simpler instructions run faster
- + reuse expensive hardware on multiple cycles
- sequencing overhead paid many times



# Multicycle ARM Processor

- **Single-cycle:**

- + simple
- cycle time limited by longest instruction (LDR)
- separate memories for instruction and data
- 3 adders/ALUs

- **Multicycle:**

- + higher clock speed
- + simpler instructions run faster
- + reuse expensive hardware on multiple cycles
- sequencing overhead paid many times

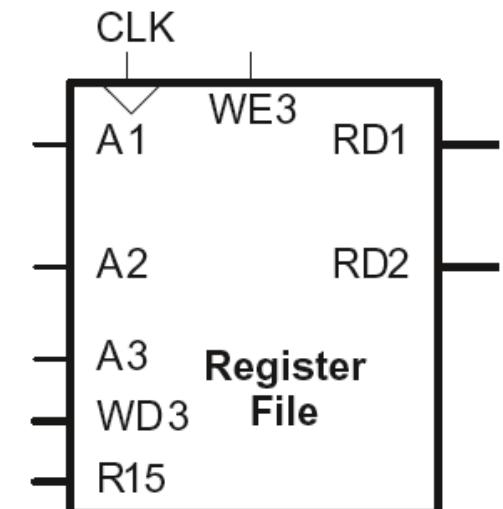
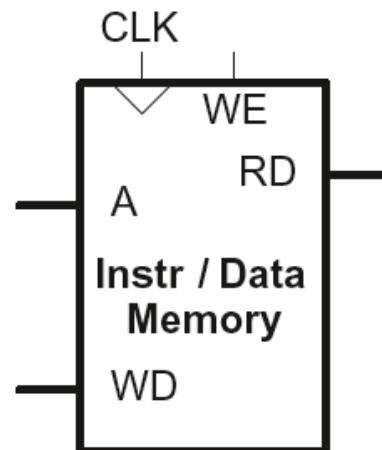
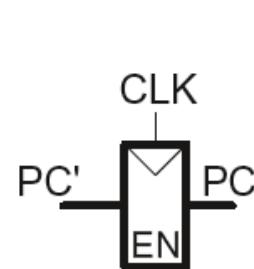
Same design steps  
as single-cycle:

- first datapath
- then control



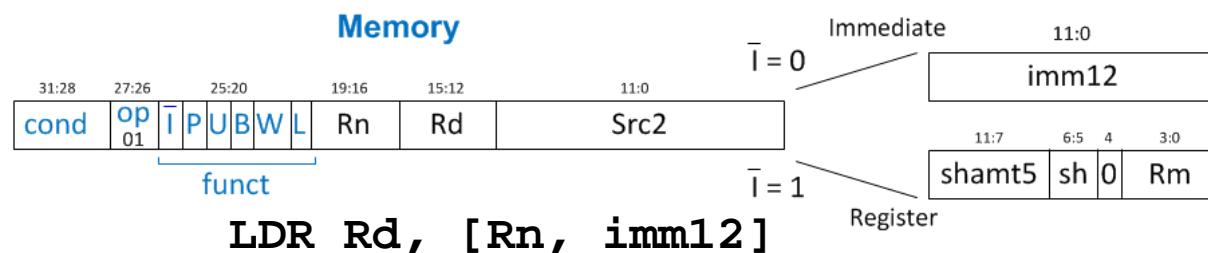
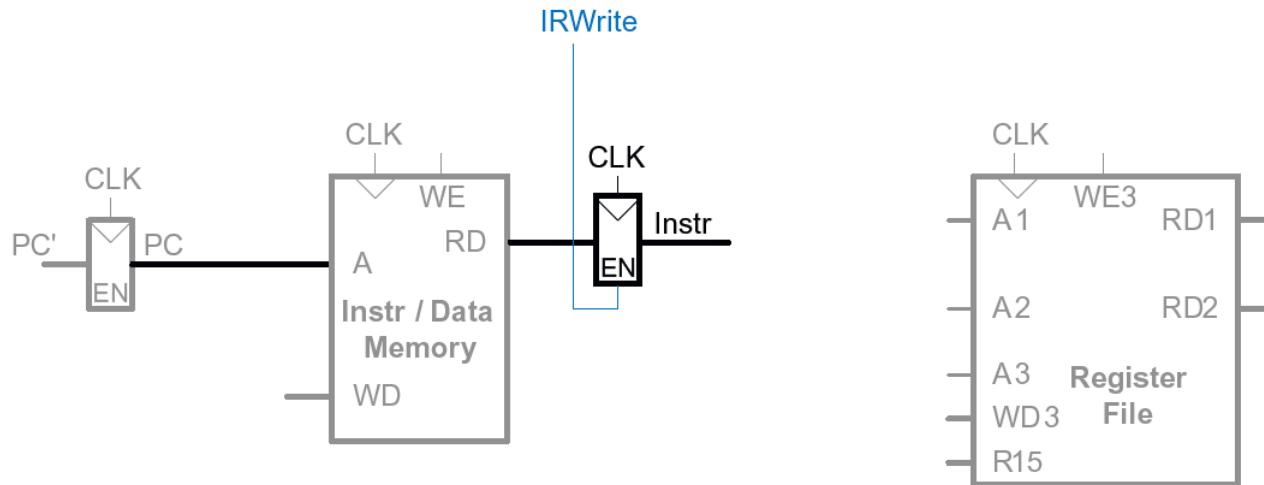
# Multicycle State Elements

Replace Instruction and Data memories with a single unified memory – more realistic



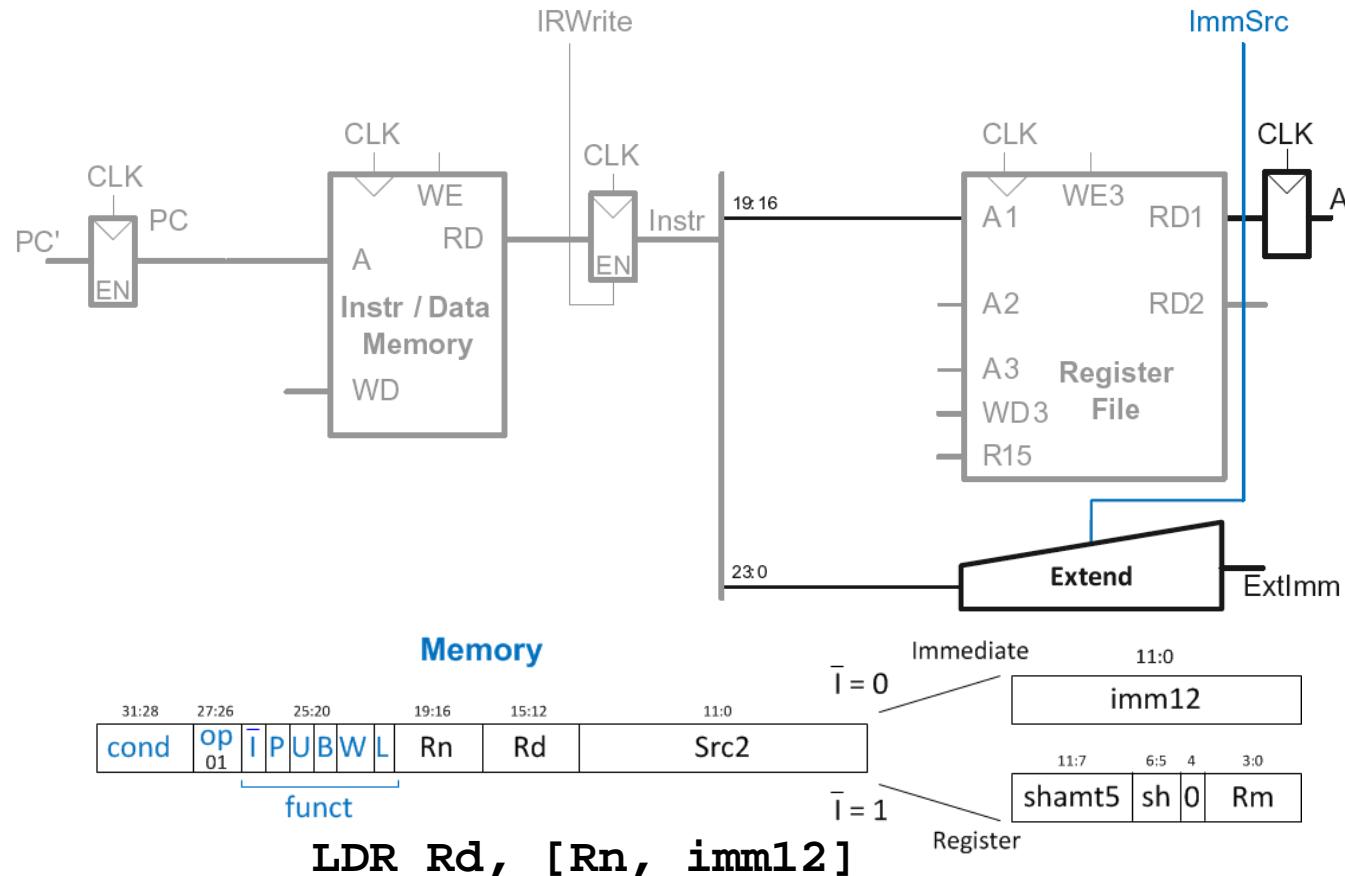
# Multicycle Datapath: Instruction Fetch

## STEP 1: Fetch instruction



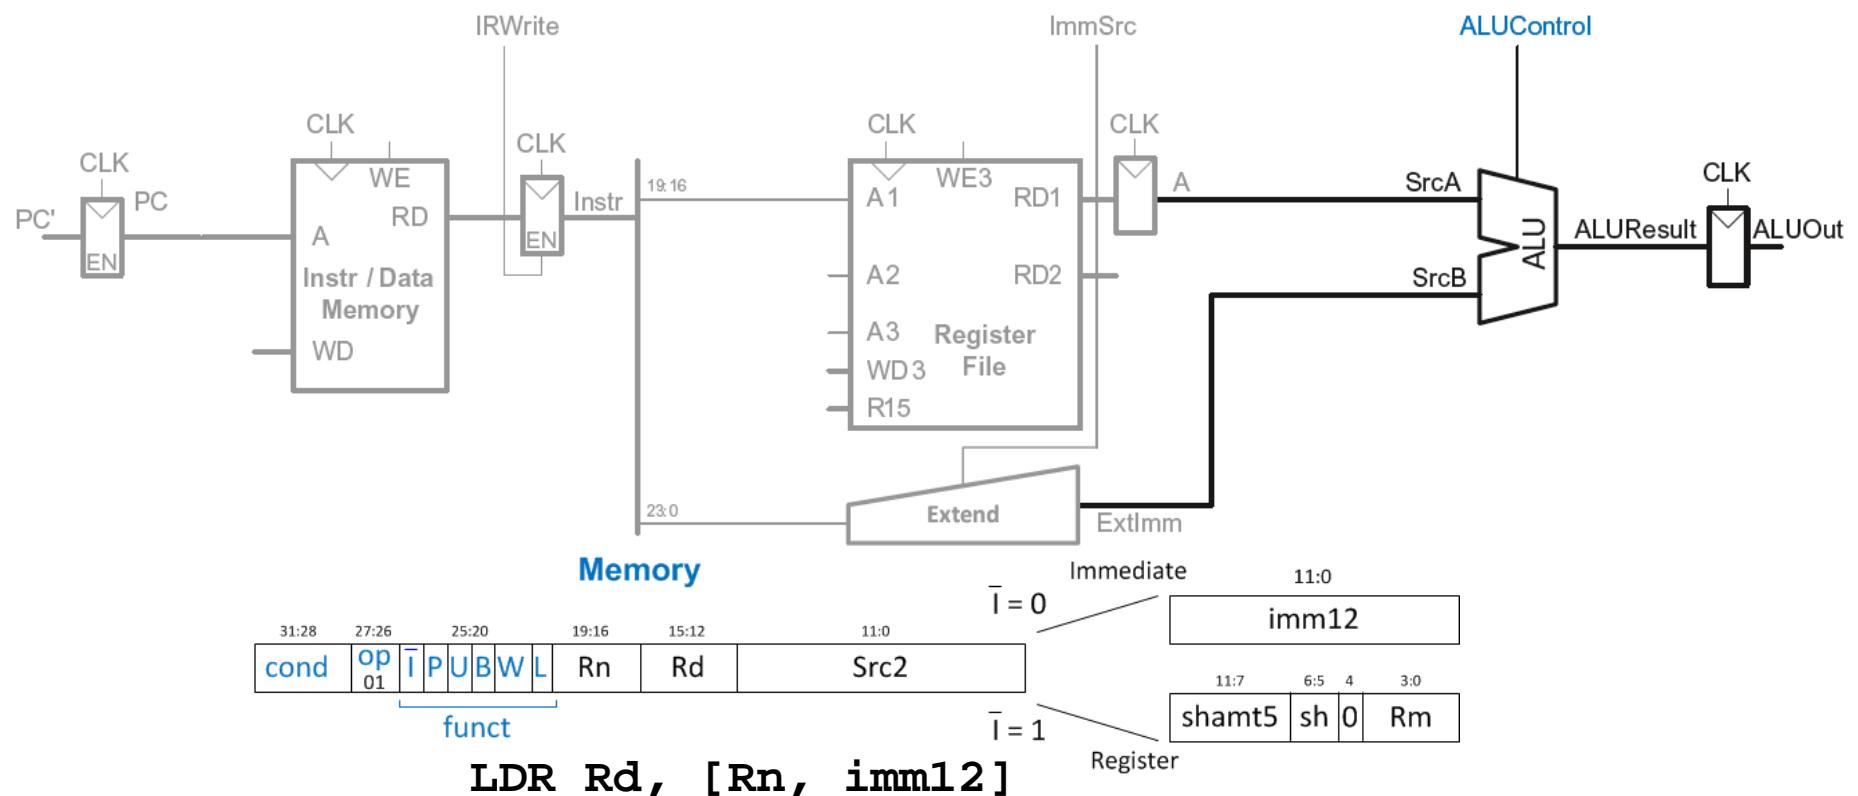
# Multicycle Datapath: LDR Register Read

## STEP 2: Read source operands from RF



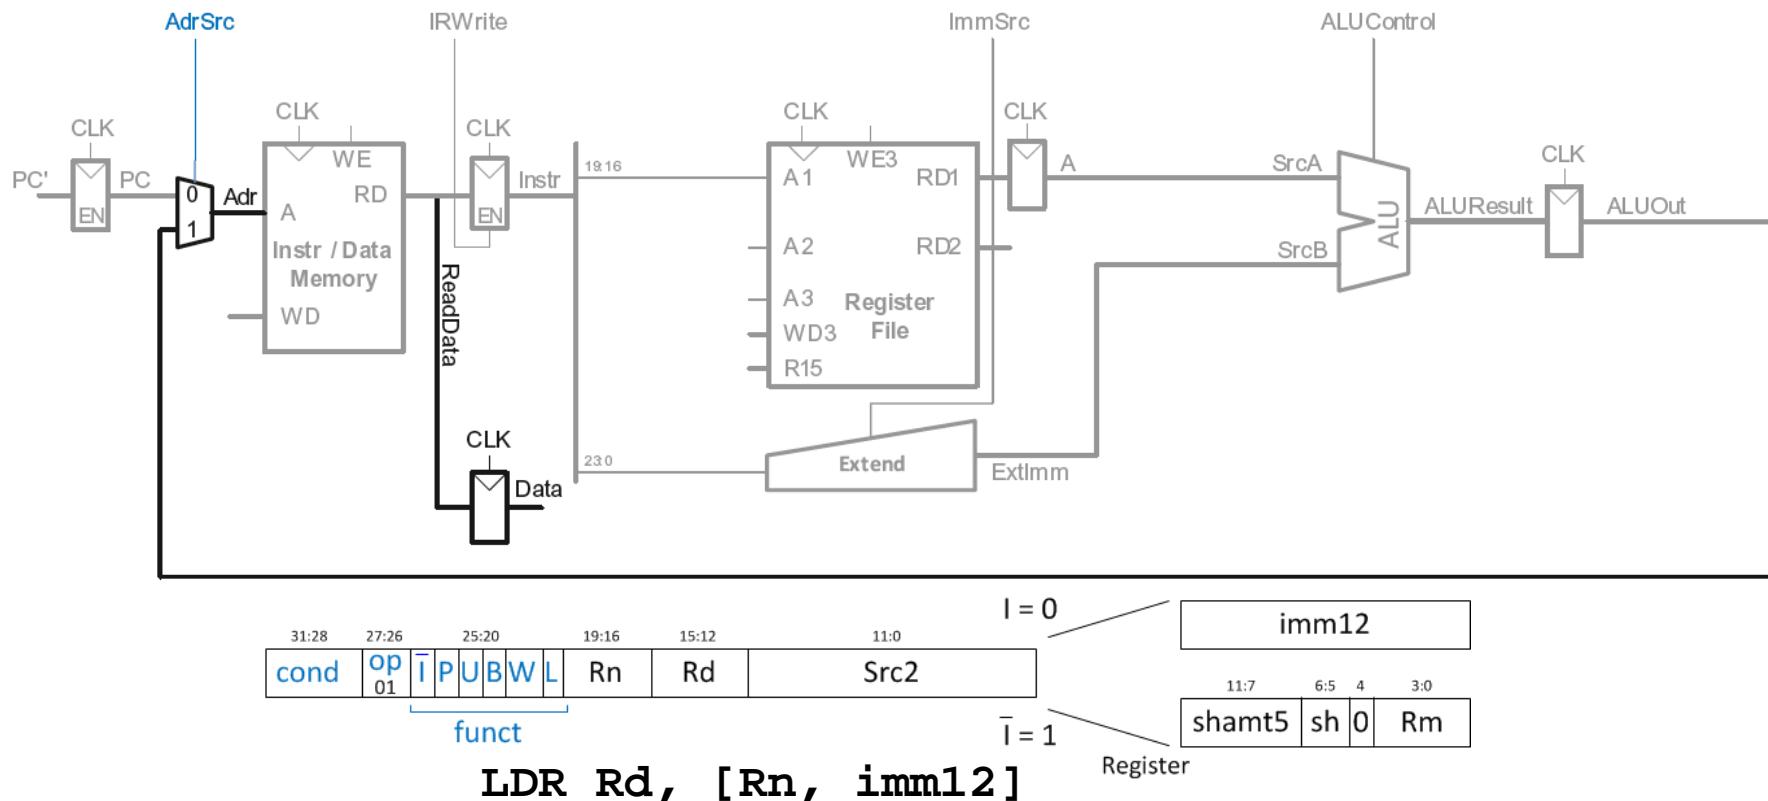
# Multicycle Datapath: LDR Address

## STEP 3: Compute the memory address



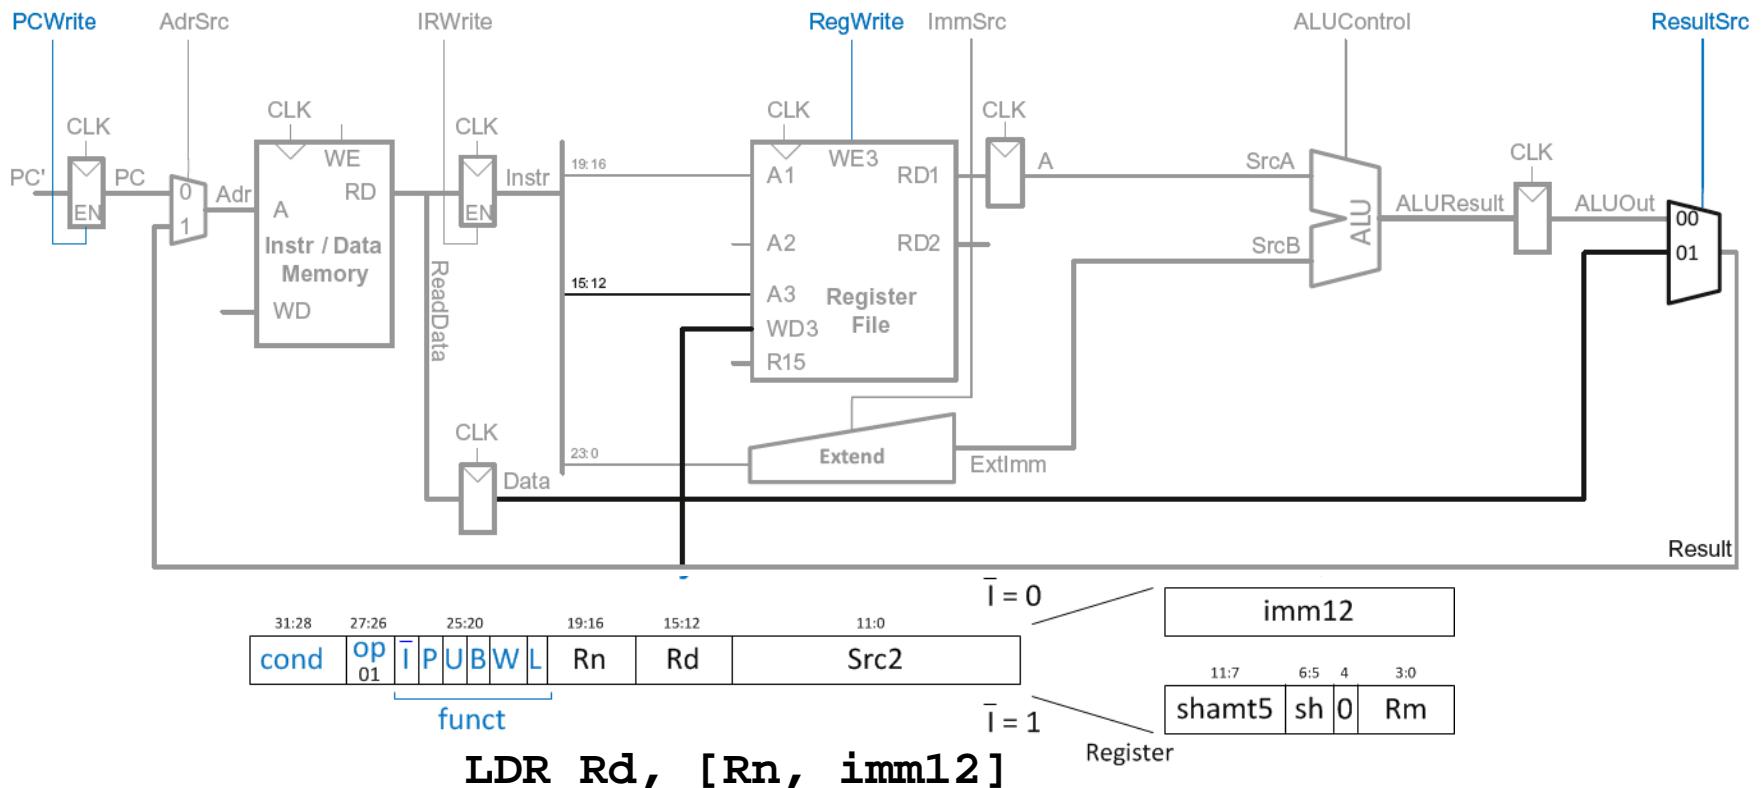
# Multicycle Datapath: LDR Memory Read

## STEP 4: Read data from memory



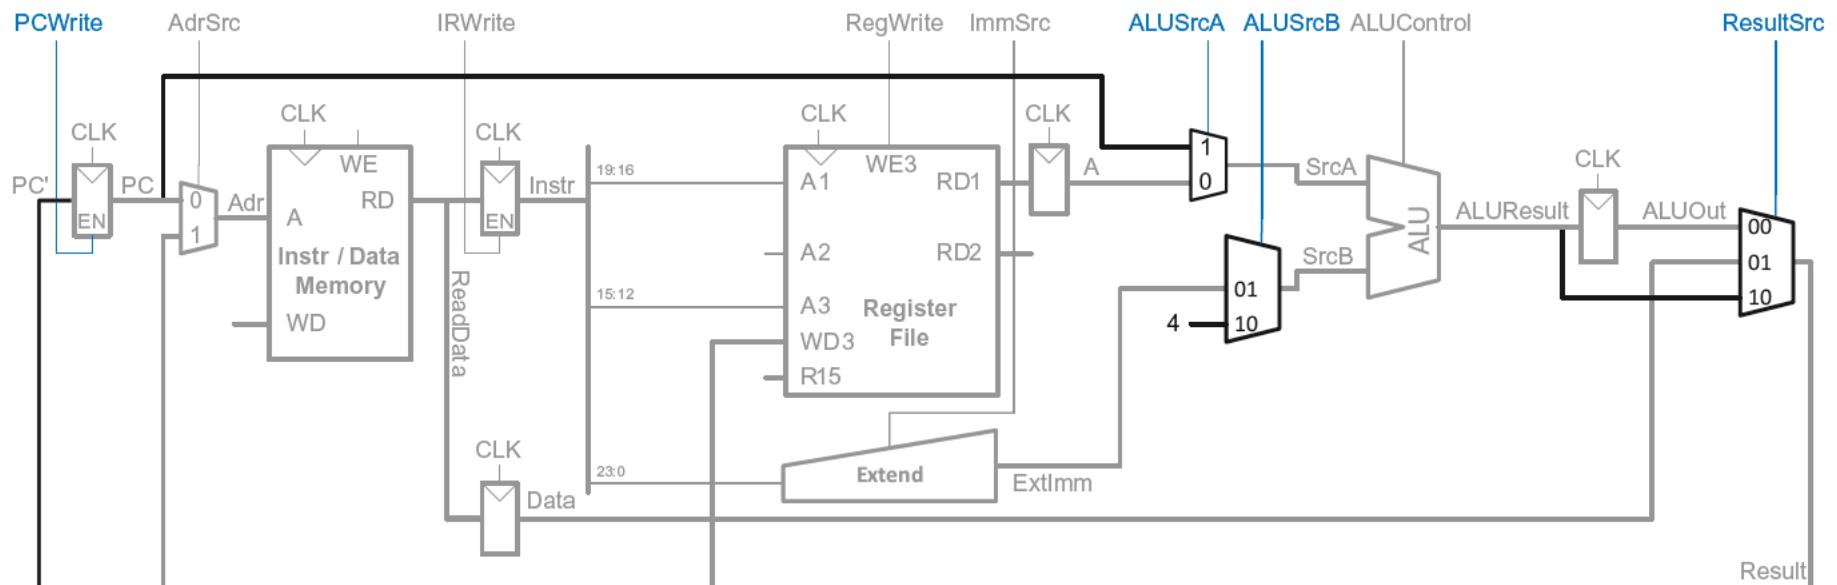
# Multicycle Datapath: LDR Write Register

## STEP 5: Write data back to register file



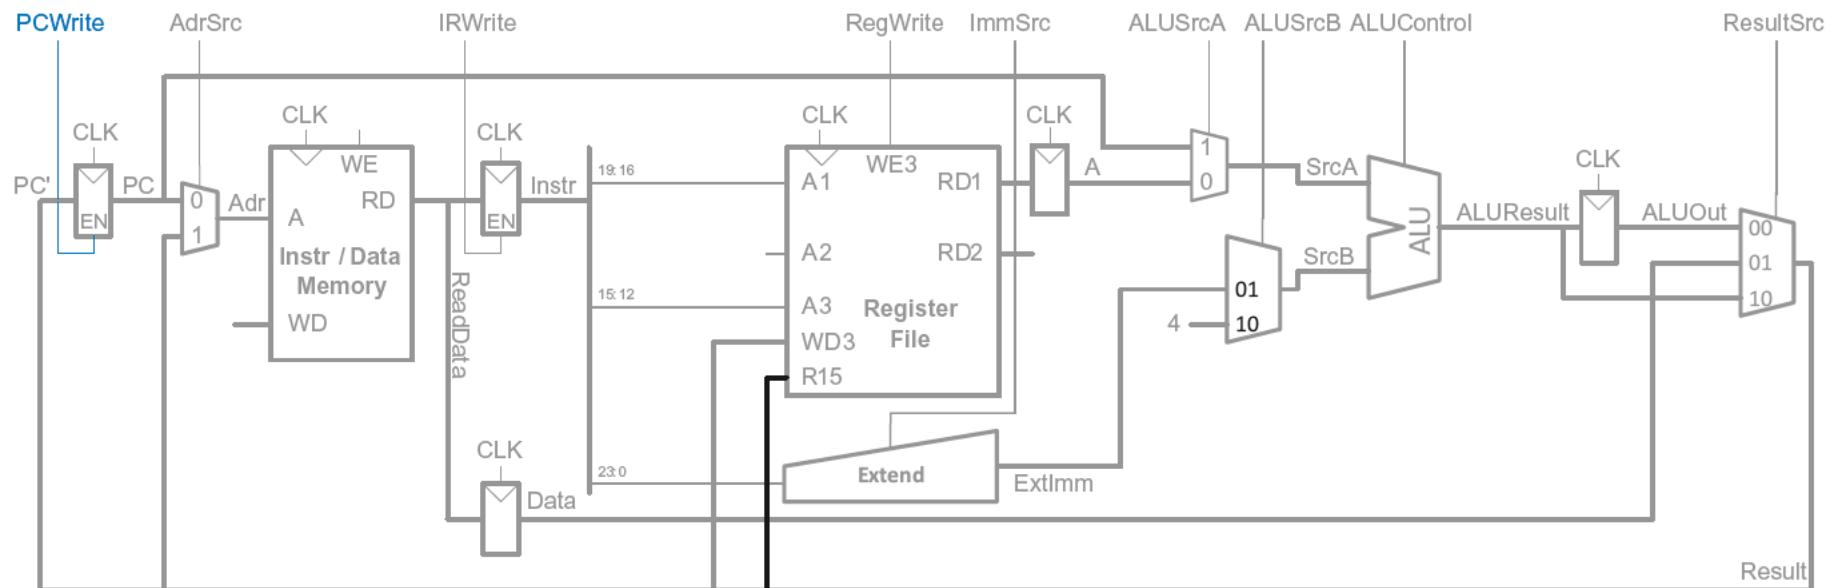
# Multicycle Datapath: Increment PC

## STEP 6: Increment PC



# Multicycle Datapath: Access to PC

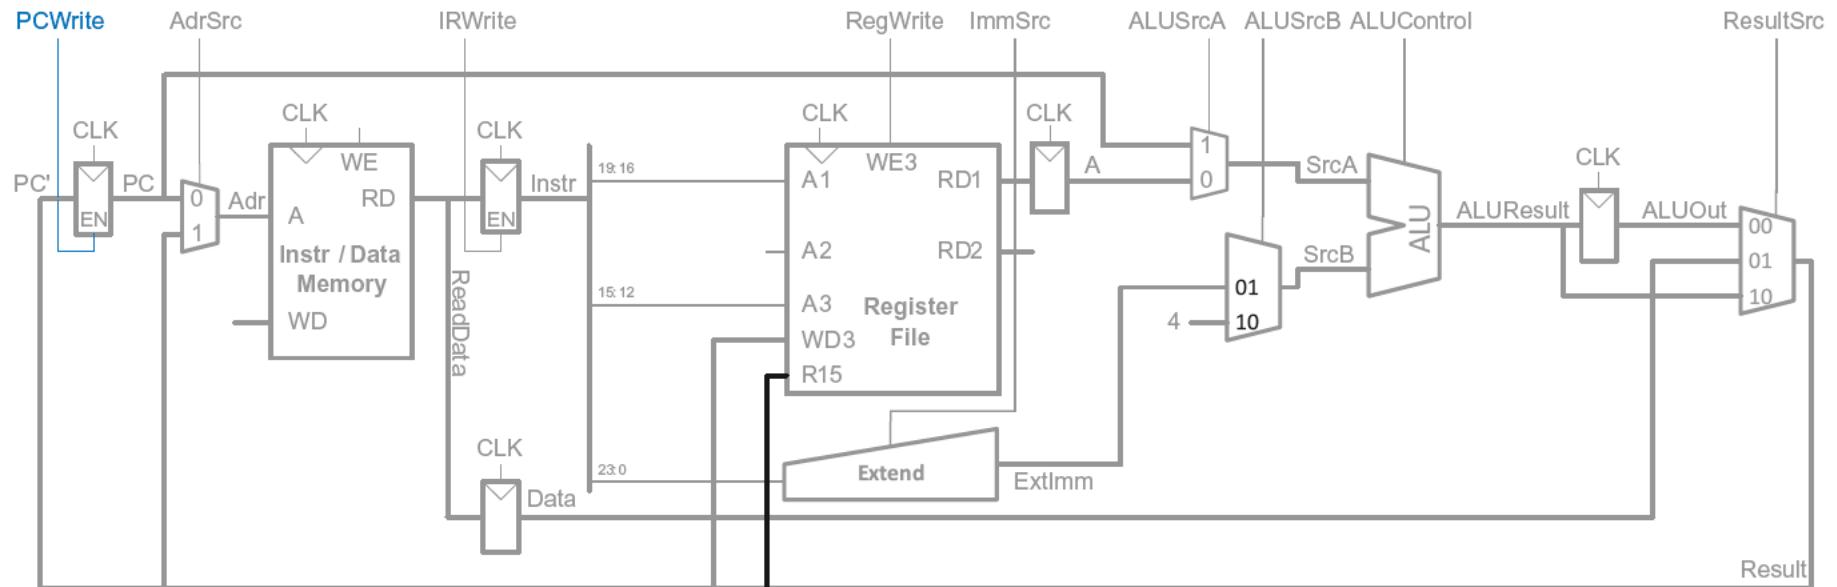
PC can be read/written by instruction



# Multicycle Datapath: Access to PC

PC can be read/written by instruction

- **Read:** R15 (PC+8) available in Register File



# Multicycle Datapath: Read to PC (R15)

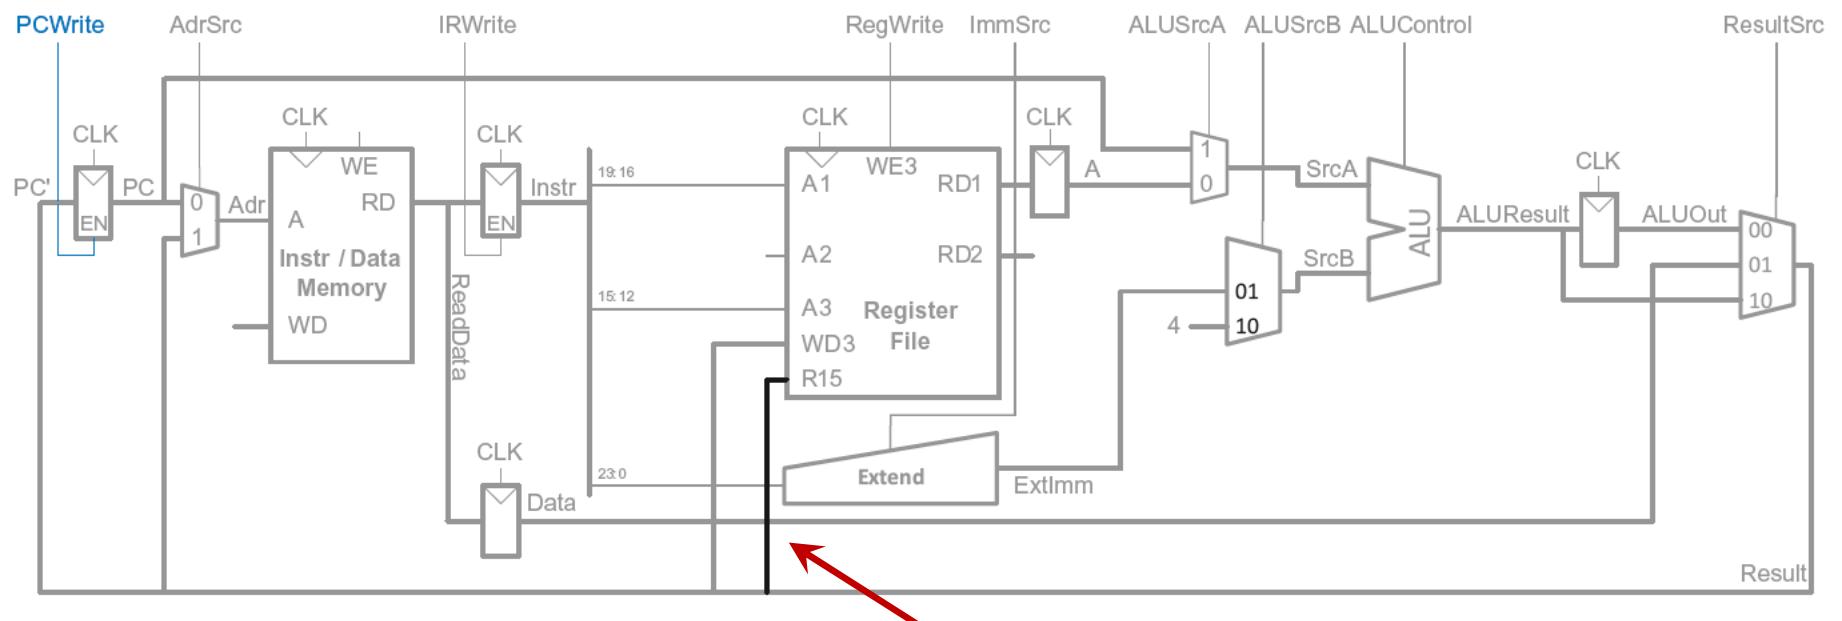
**Example:** ADD R1 , R15 , R2



# Multicycle Datapath: Read to PC (R15)

**Example:** ADD R1 , R15 , R2

- R15 needs to be read as PC+8 from Register File (RF) in 2<sup>nd</sup> step
- So (also in 2<sup>nd</sup> step) PC + 8 is produced by ALU and routed to R15 input of RF



# Multicycle Datapath: Read to PC (R15)

## Example: ADD R1 , R15 , R2

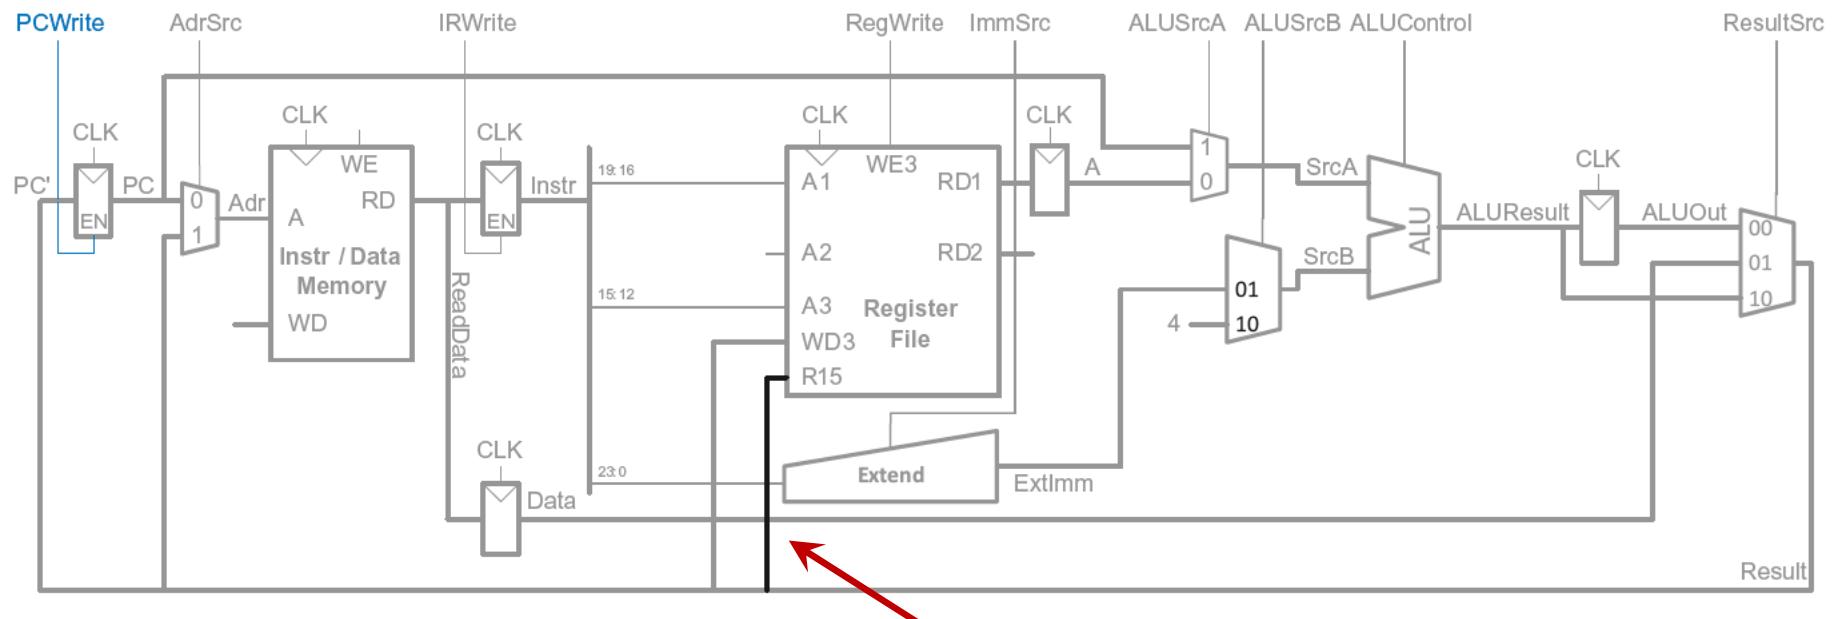
- R15 needs to be read as PC+8 from Register File (RF) in 2<sup>nd</sup> step
- So (also in 2<sup>nd</sup> step) PC + 8 is produced by ALU and routed to R15 input of RF
  - $SrcA = PC$  (which was already updated in step 1 to PC+4)
  - $SrcB = 4$
  - $ALUResult = PC + 8$
- ALUResult is fed to R15 input port of RF in 2<sup>nd</sup> step (which is then routed to RD1 output of RF)



# Multicycle Datapath: Read to PC (R15)

**Example:** ADD R1 , R15 , R2

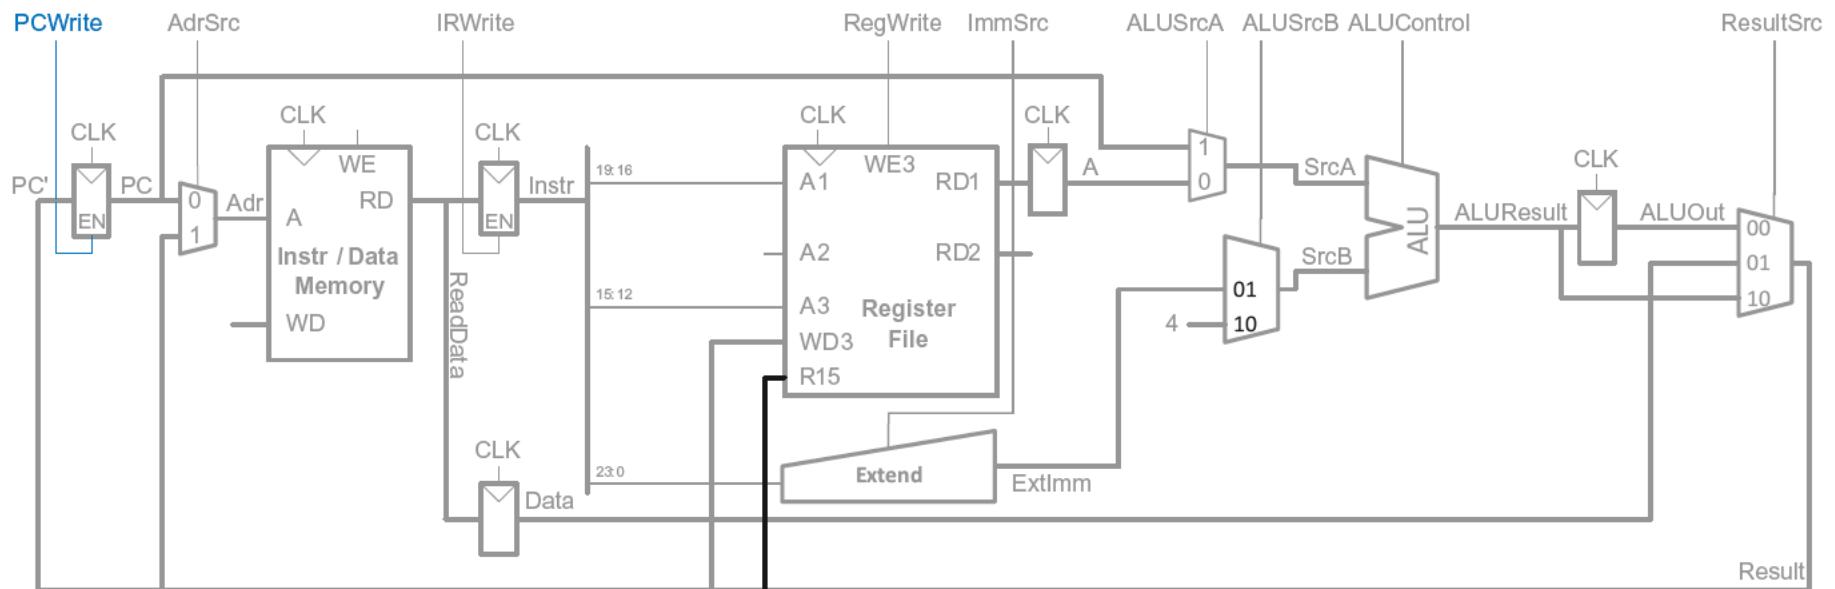
- R15 needs to be read as PC+8 from Register File (RF) in 2<sup>nd</sup> step
- So (also in 2<sup>nd</sup> step) PC + 8 is produced by ALU and routed to R15 input of RF



# Multicycle Datapath: Access to PC

PC can be read/written by instruction

- **Read:** R15 (PC+8) available in Register File
- **Write:** Be able to write result of instruction to PC



# Multicycle Datapath: Write to PC (R15)

**Example:** SUB R15, R8, R3



# Multicycle Datapath: Write to PC (R15)

**Example:** SUB R15, R8, R3

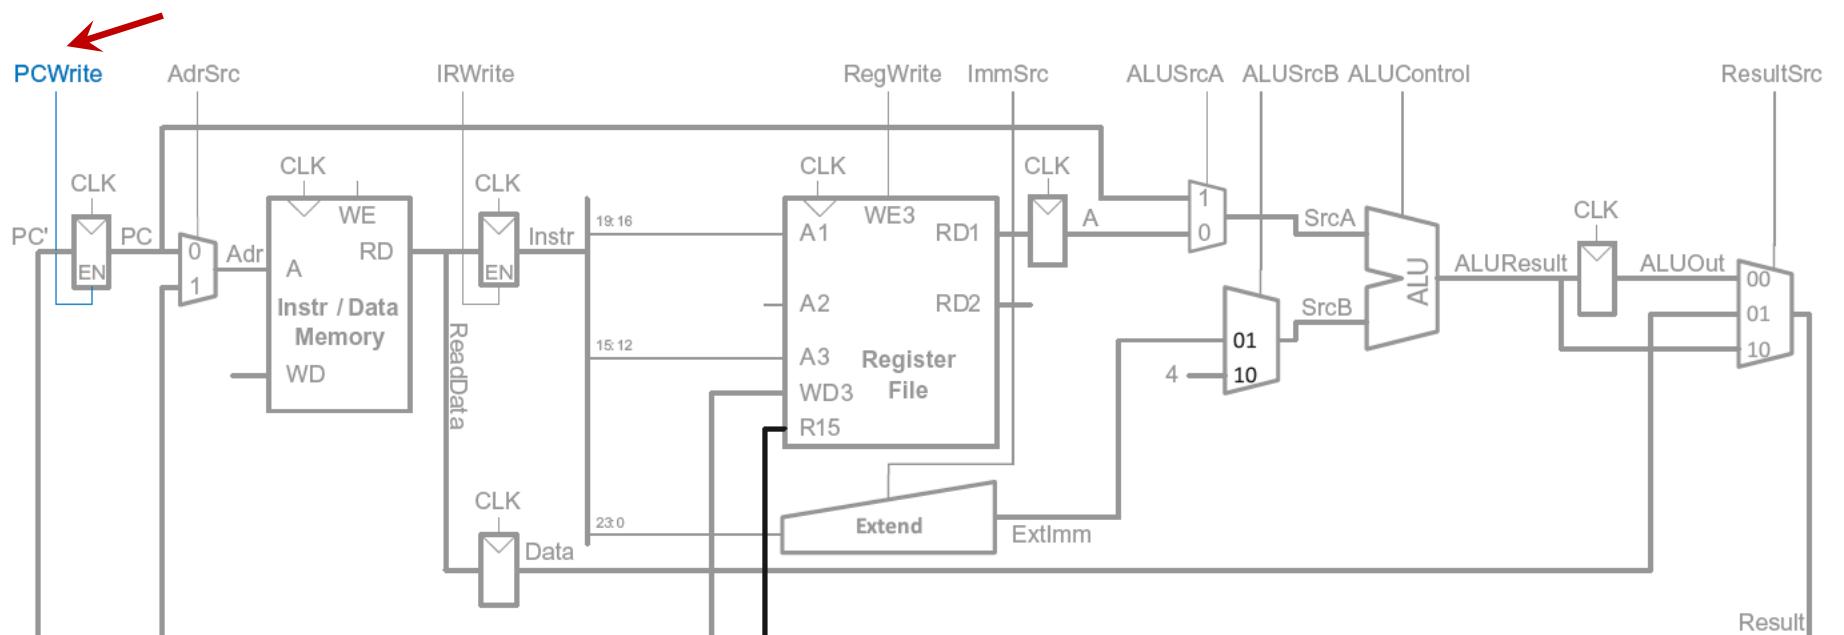
- Result of instruction needs to be written to the PC register
- ALUResult already routed to the PC register, just assert PCWrite



# Multicycle Datapath: Write to PC (R15)

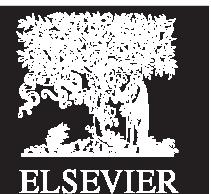
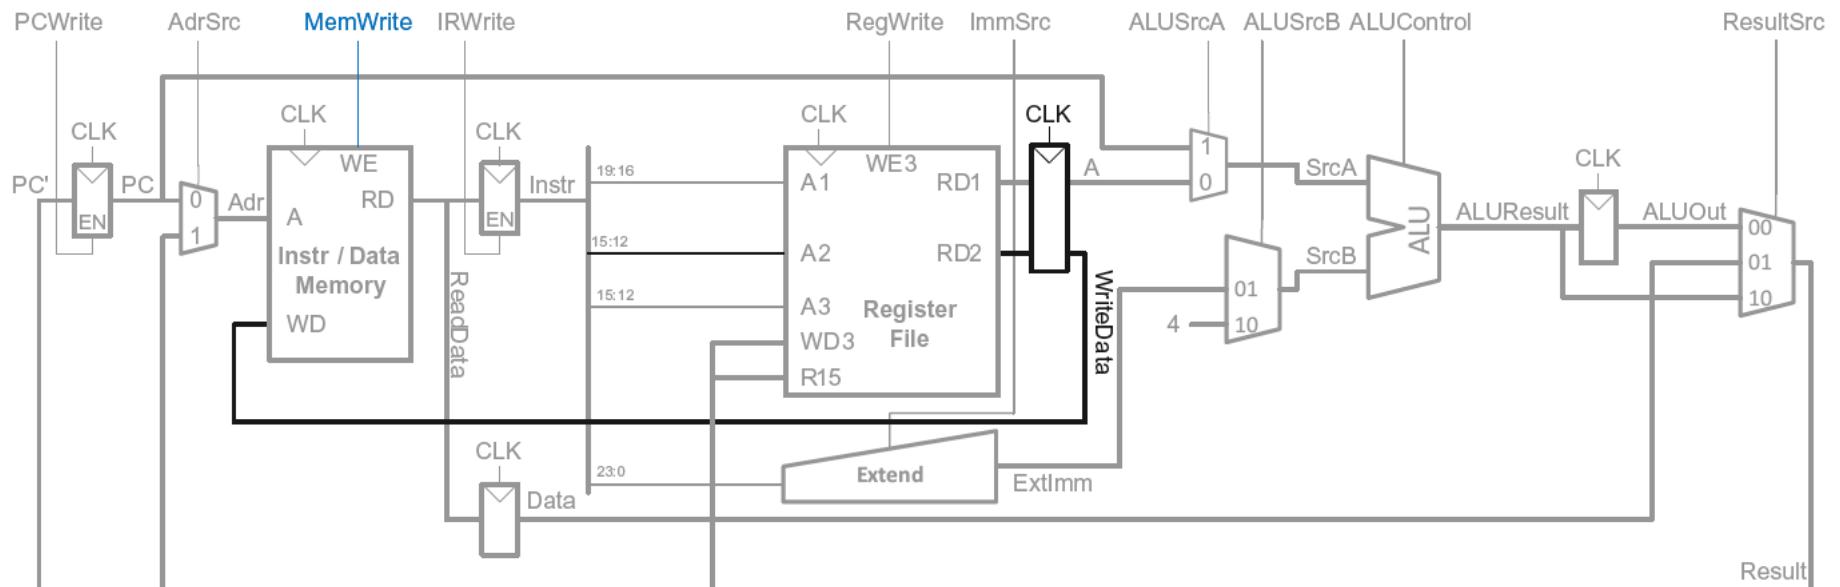
**Example:** SUB R15, R8 , R3

- Result of instruction needs to be written to the PC register
- ALUResult already routed to the PC register, just assert PCWrite



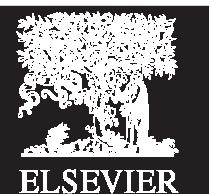
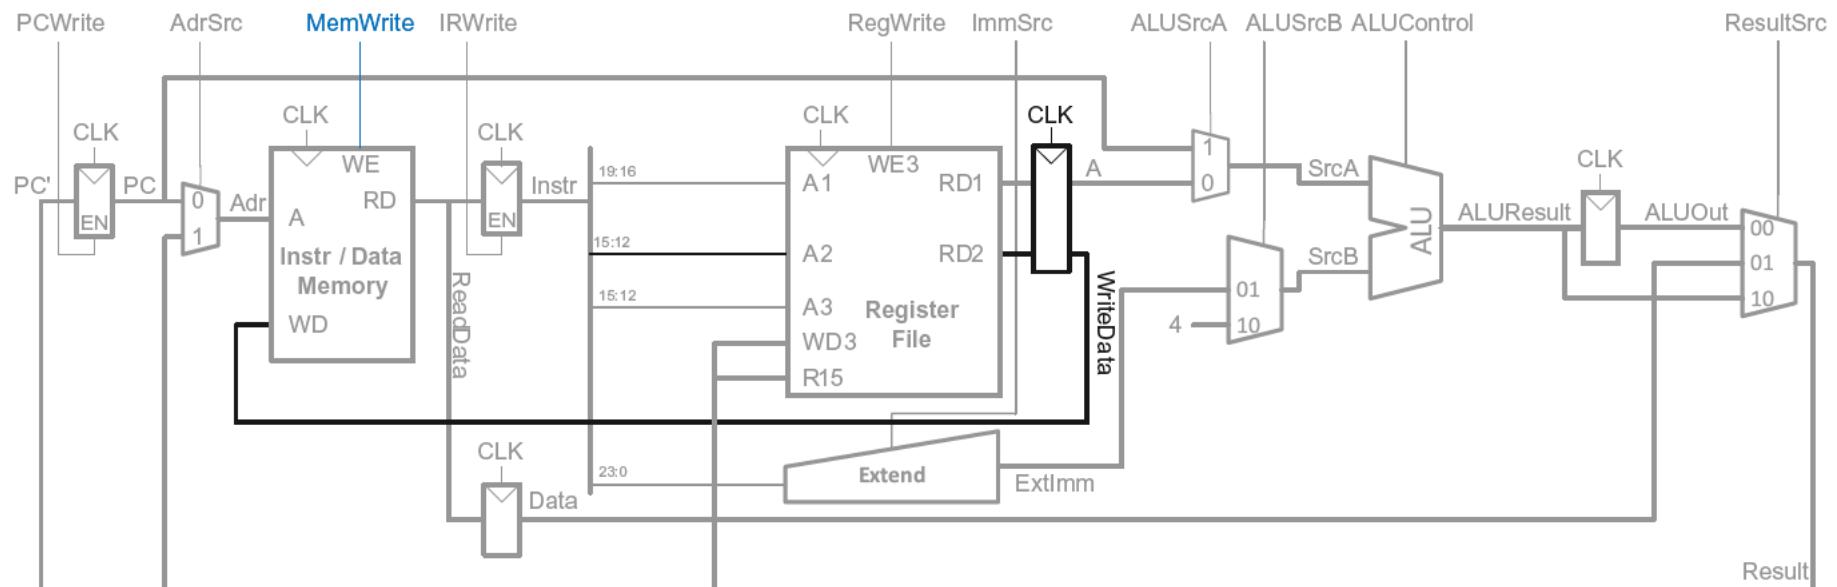
# Multicycle Datapath: STR

Write data in Rn to memory



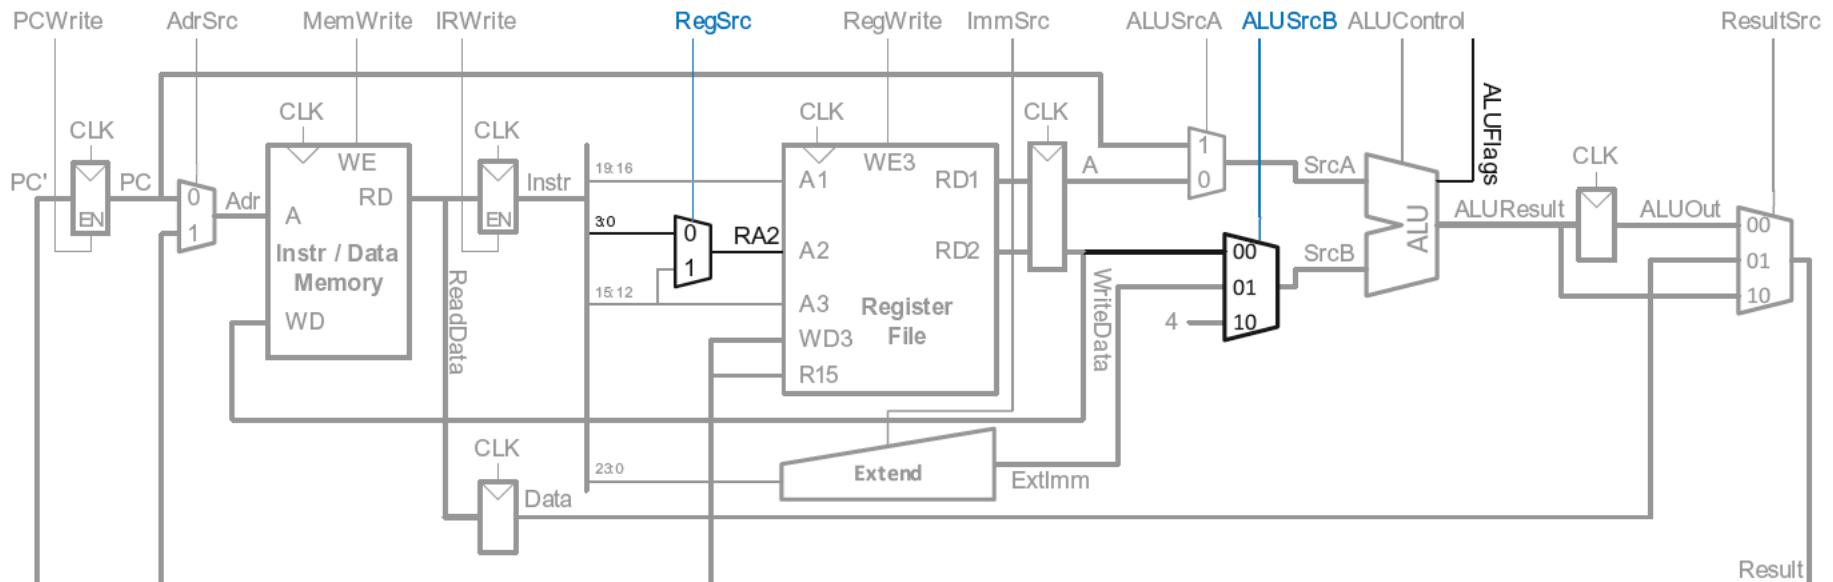
# Multicycle Datapath: Data-processing

With immediate addressing (i.e., an immediate *Src2*), no additional changes needed for datapath



# Multicycle Datapath: Data-processing

With register addressing (register *Src2*):  
Read from Rn and Rm

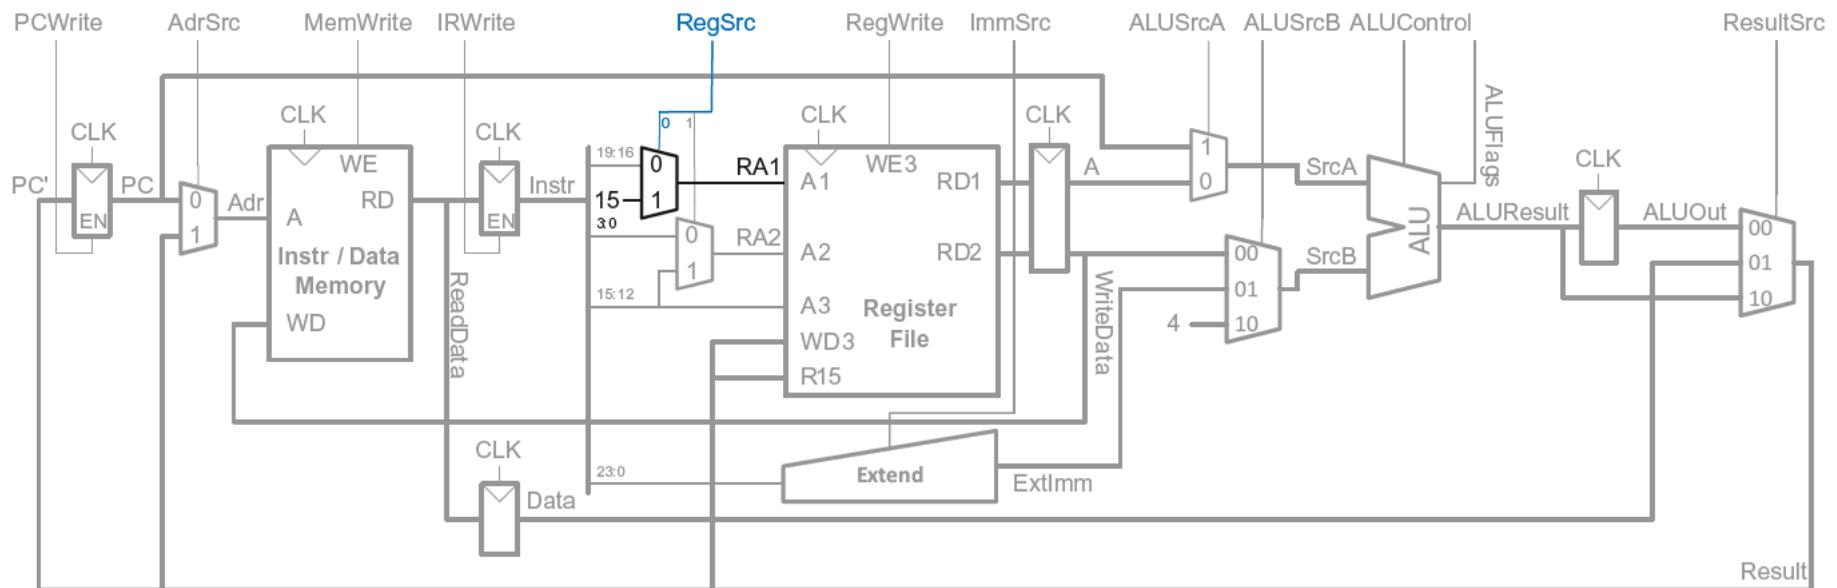


# Multicycle Datapath: B

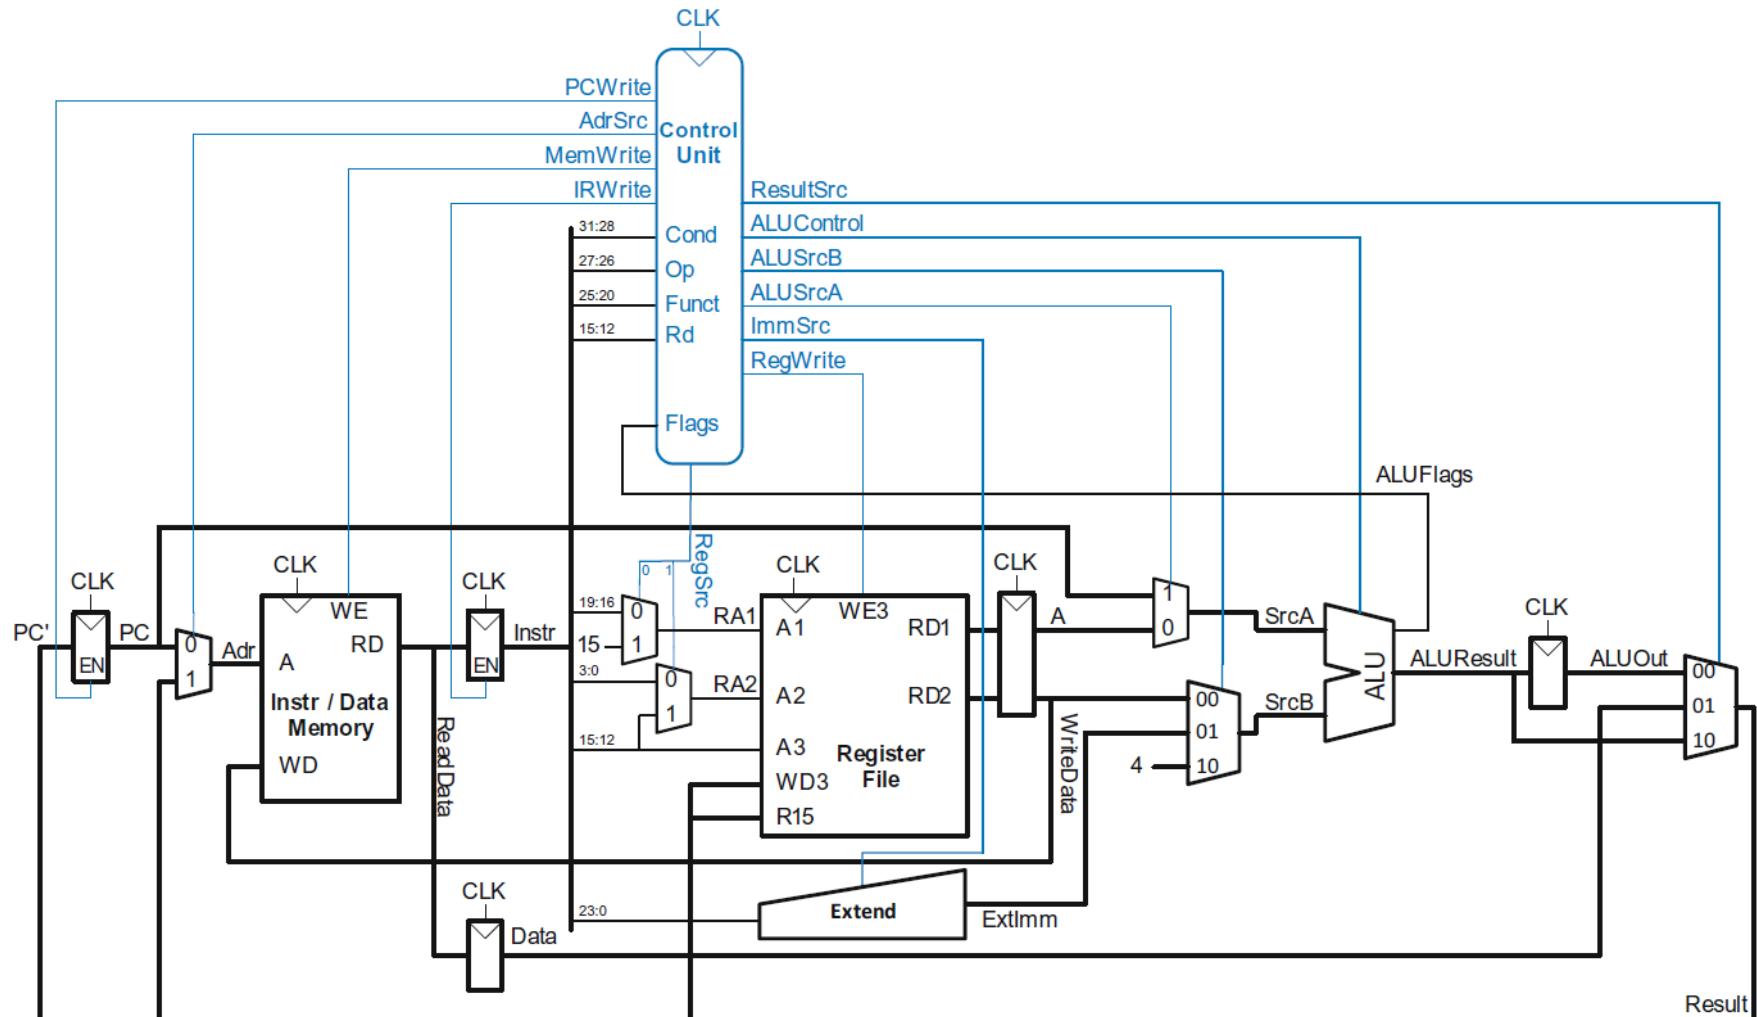
Calculate branch target address:

$$BTA = (ExtImm) + (PC+8)$$

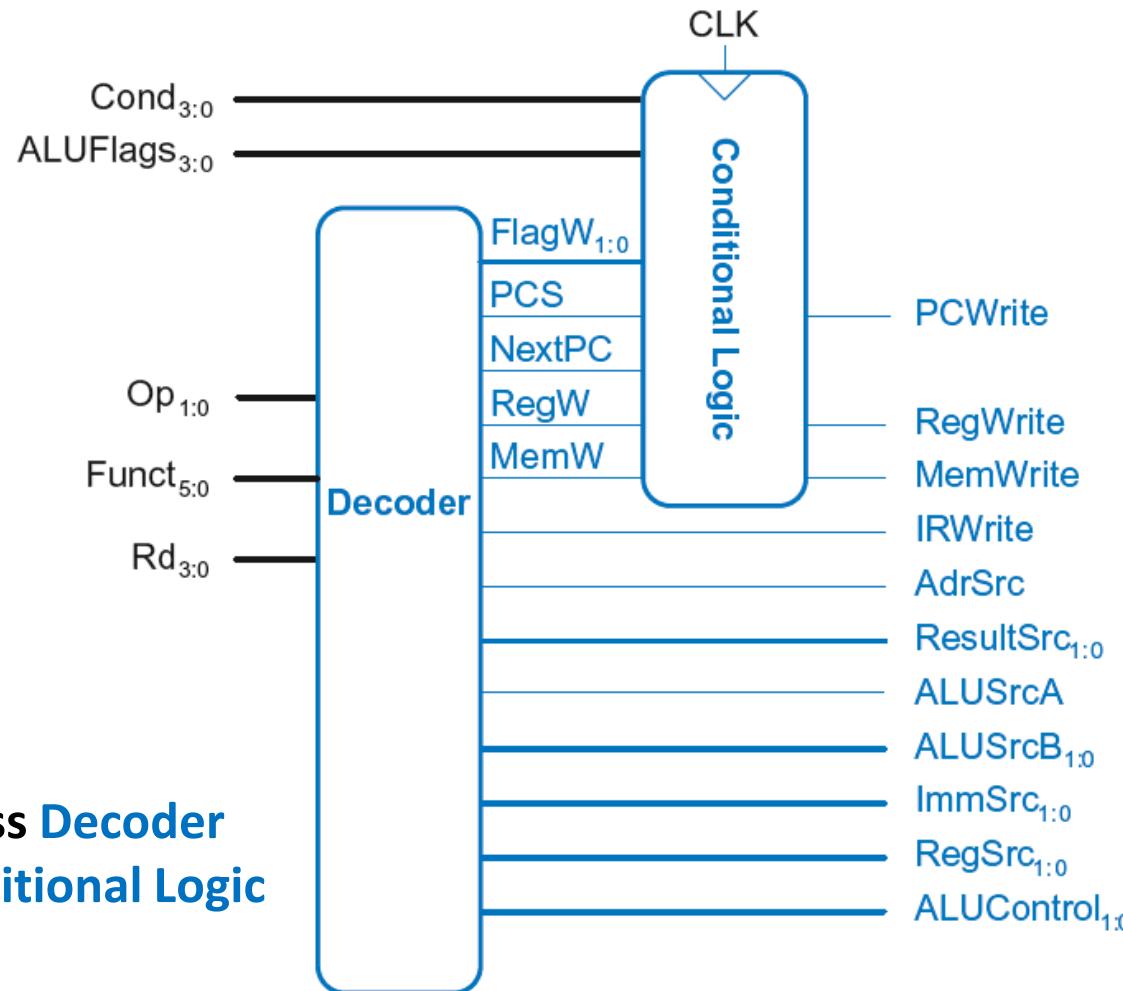
$ExtImm = Imm24 \ll 2$  and sign-extended



# Multicycle ARM Processor



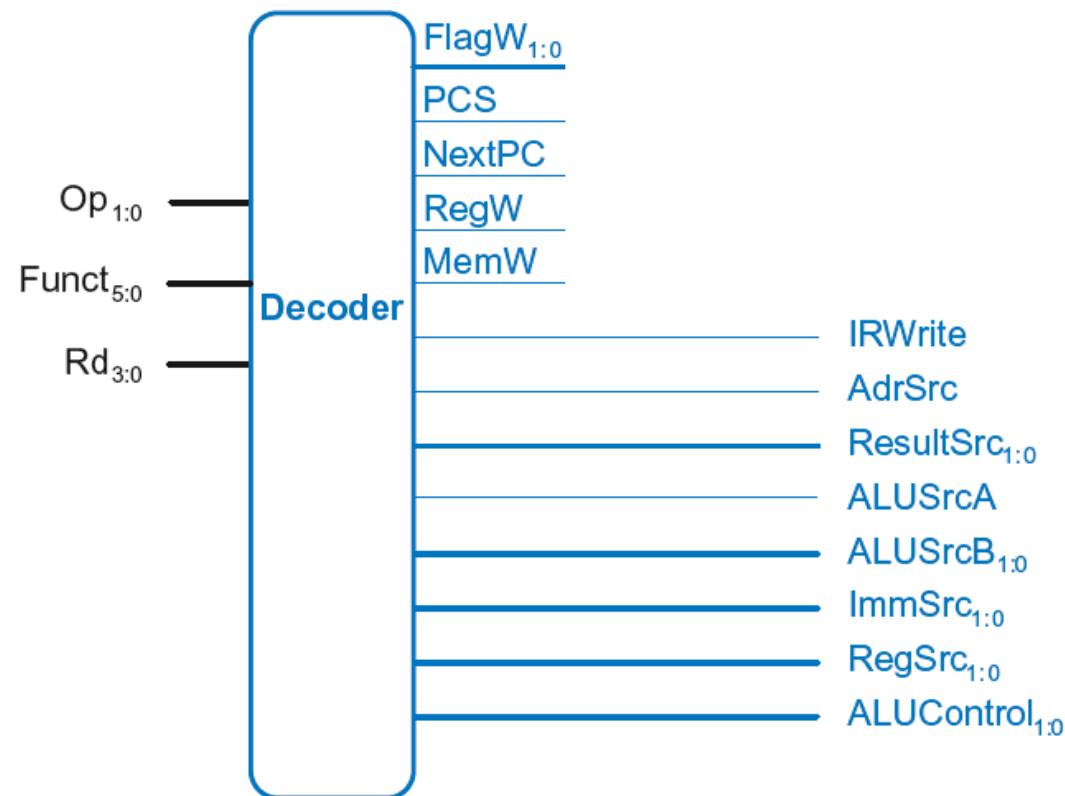
# Multicycle Control



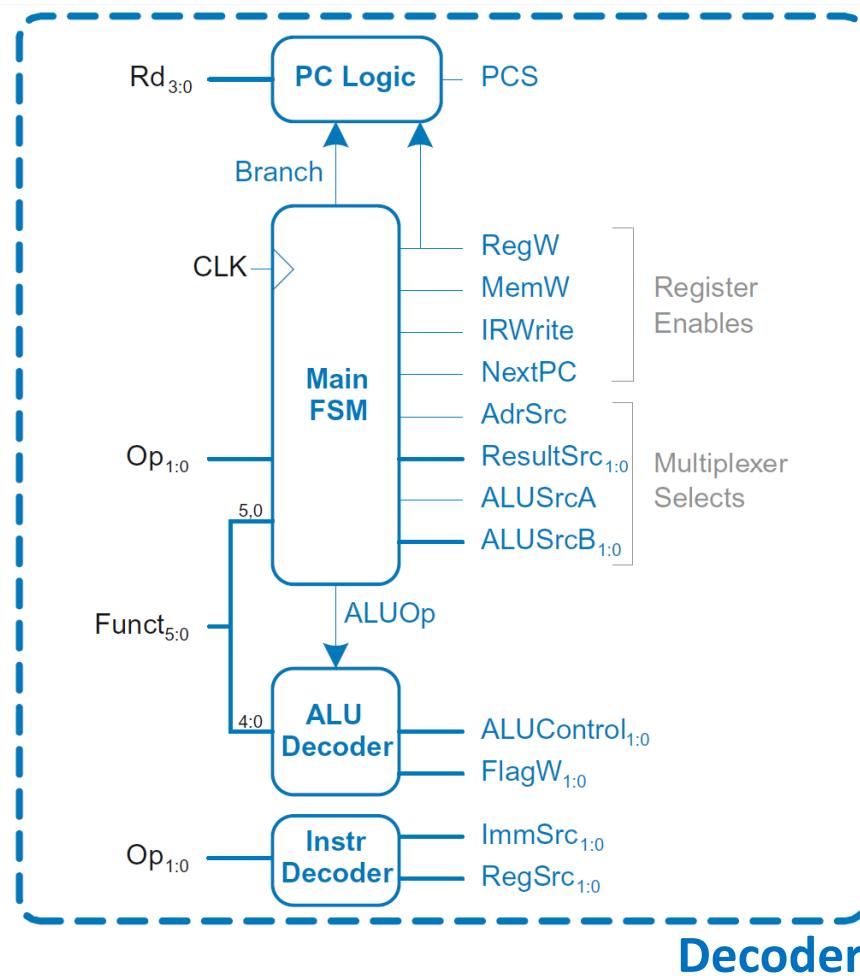
- First, discuss **Decoder**
- Then, **Conditional Logic**



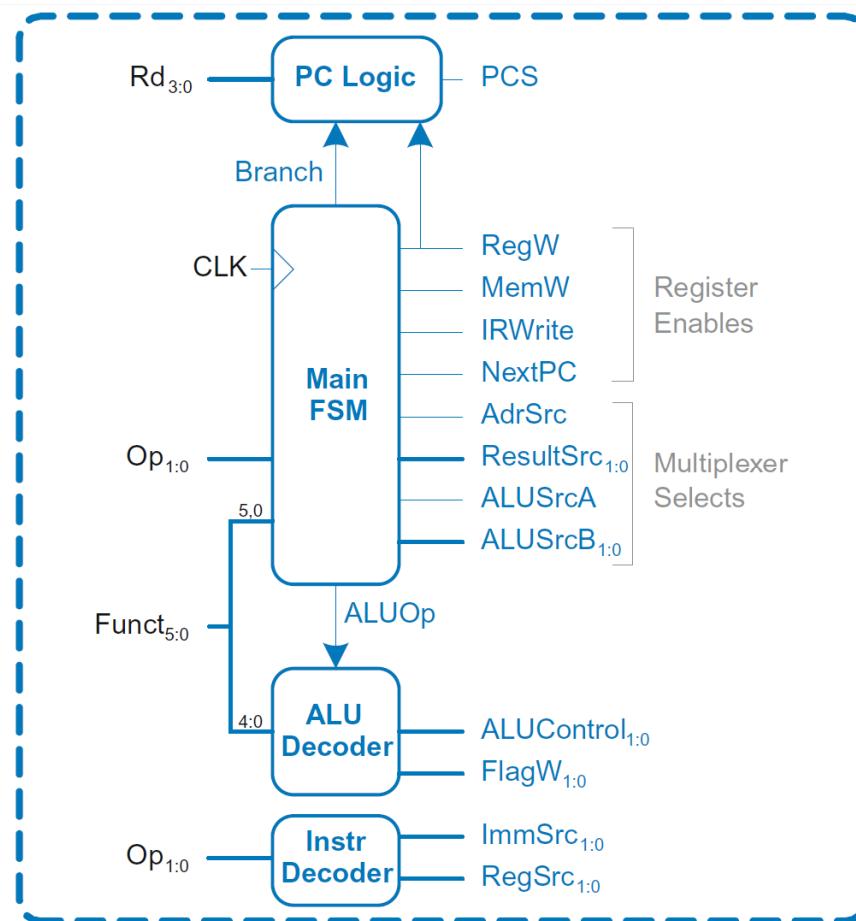
# Multicycle Control: Decoder



# Multicycle Control: Decoder



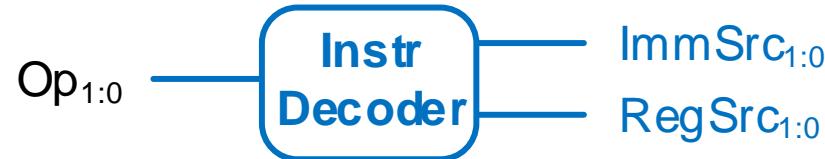
# Multicycle Control: Decoder



**ALU Decoder and PC Logic same as single-cycle**



# Multicycle Control: Instr Decoder



$$RegSrc_0 = (Op == 10_2)$$

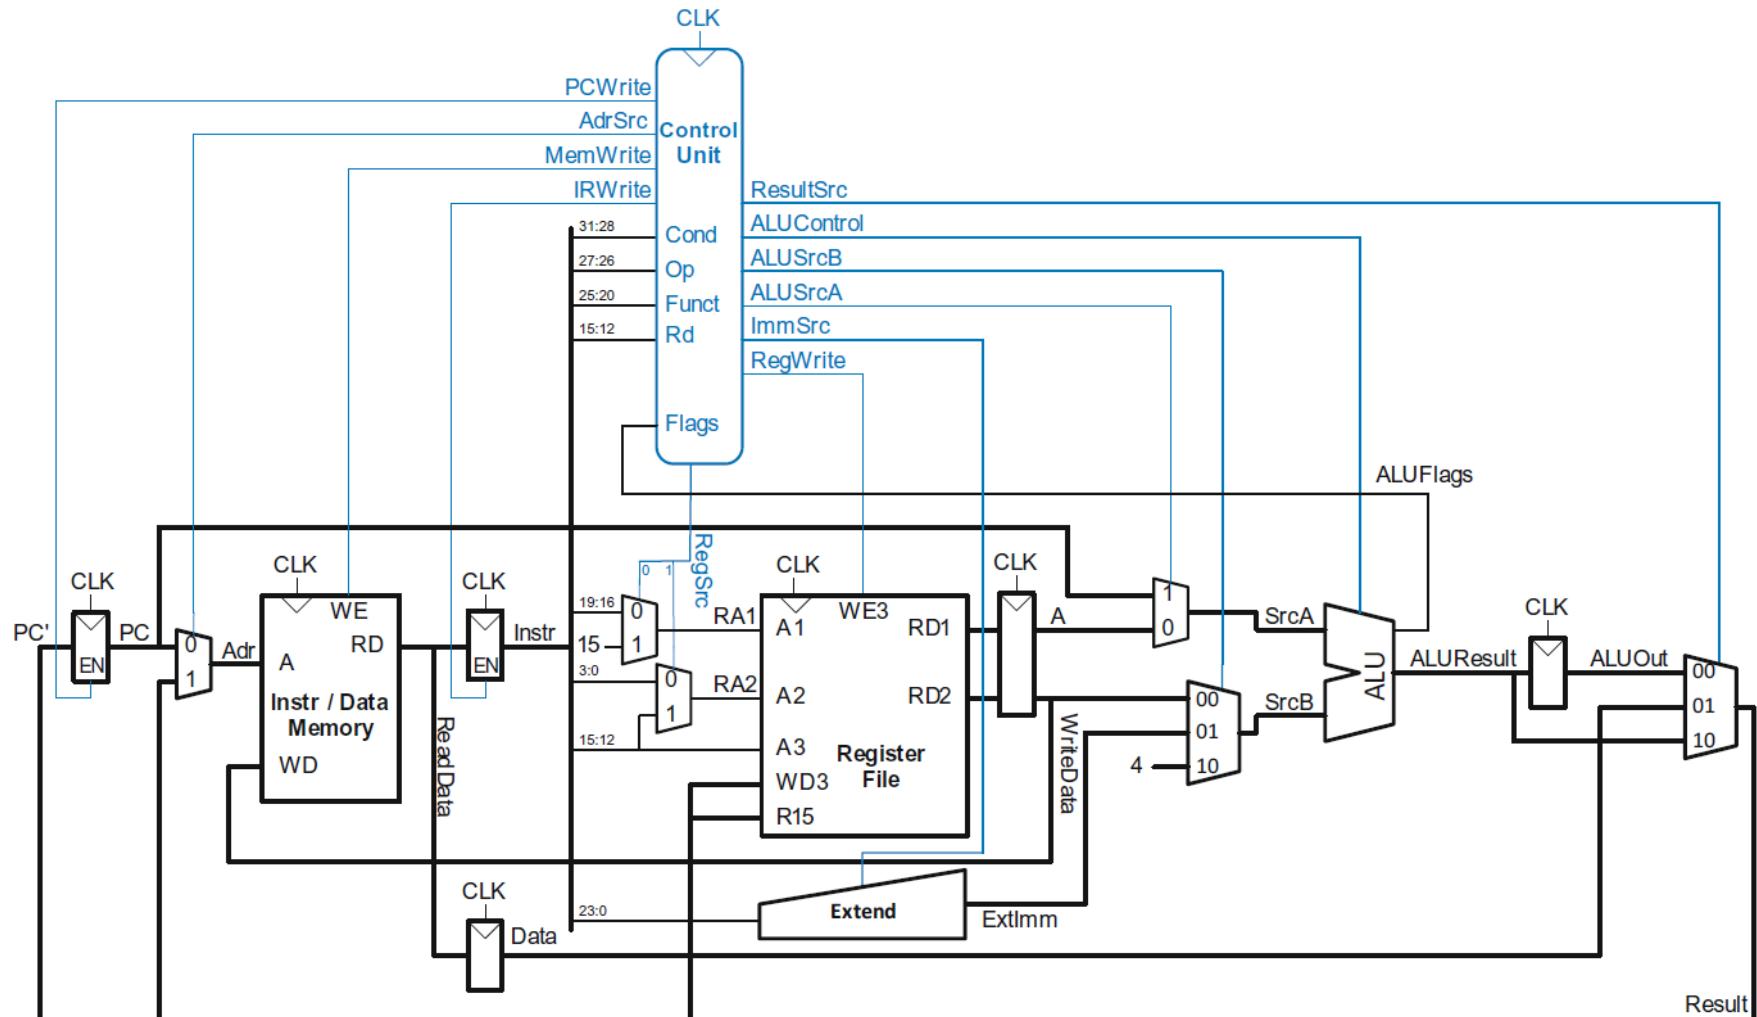
$$RegSrc_1 = (Op == 01_2)$$

$$ImmSrc_{1:0} = Op$$

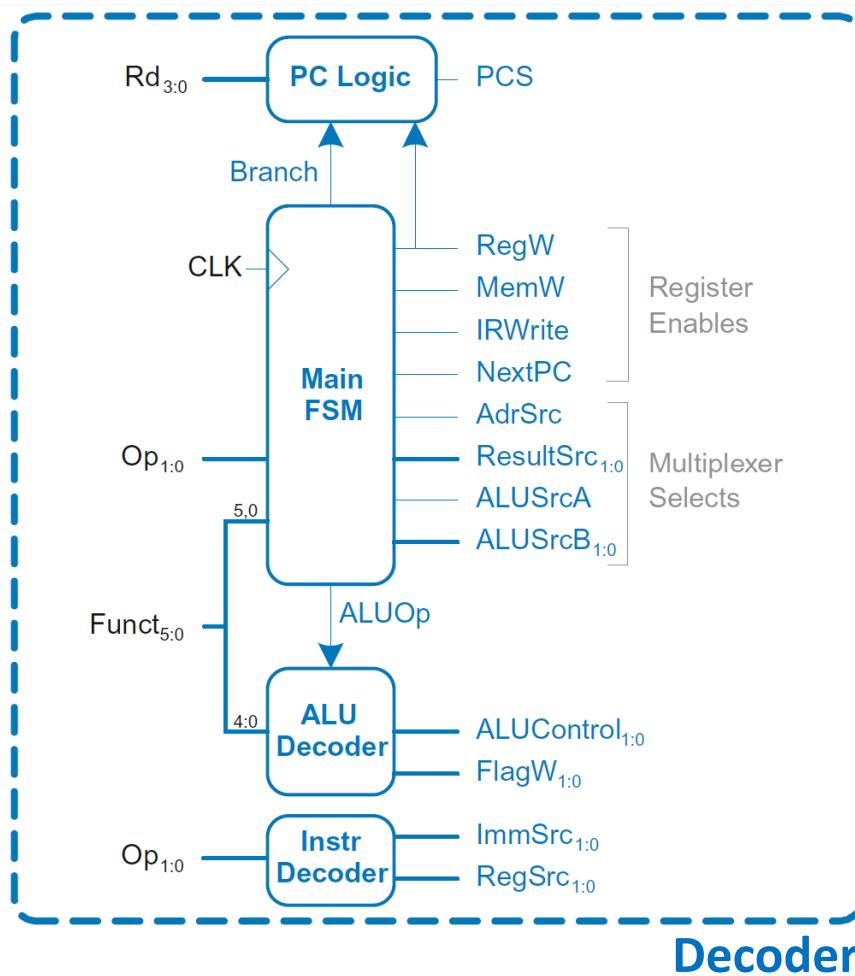
Instruction	$O_p$	$Funct_5$	$Funct_0$	$RegSrc_0$	$RegSrc_1$	$ImmSrc_{1:0}$
LDR	01	X	1	0	X	01
STR	01	X	0	0	1	01
DP immediate	00	1	X	0	X	00
DP register	00	0	X	0	0	00
B	10	X	X	1	X	10



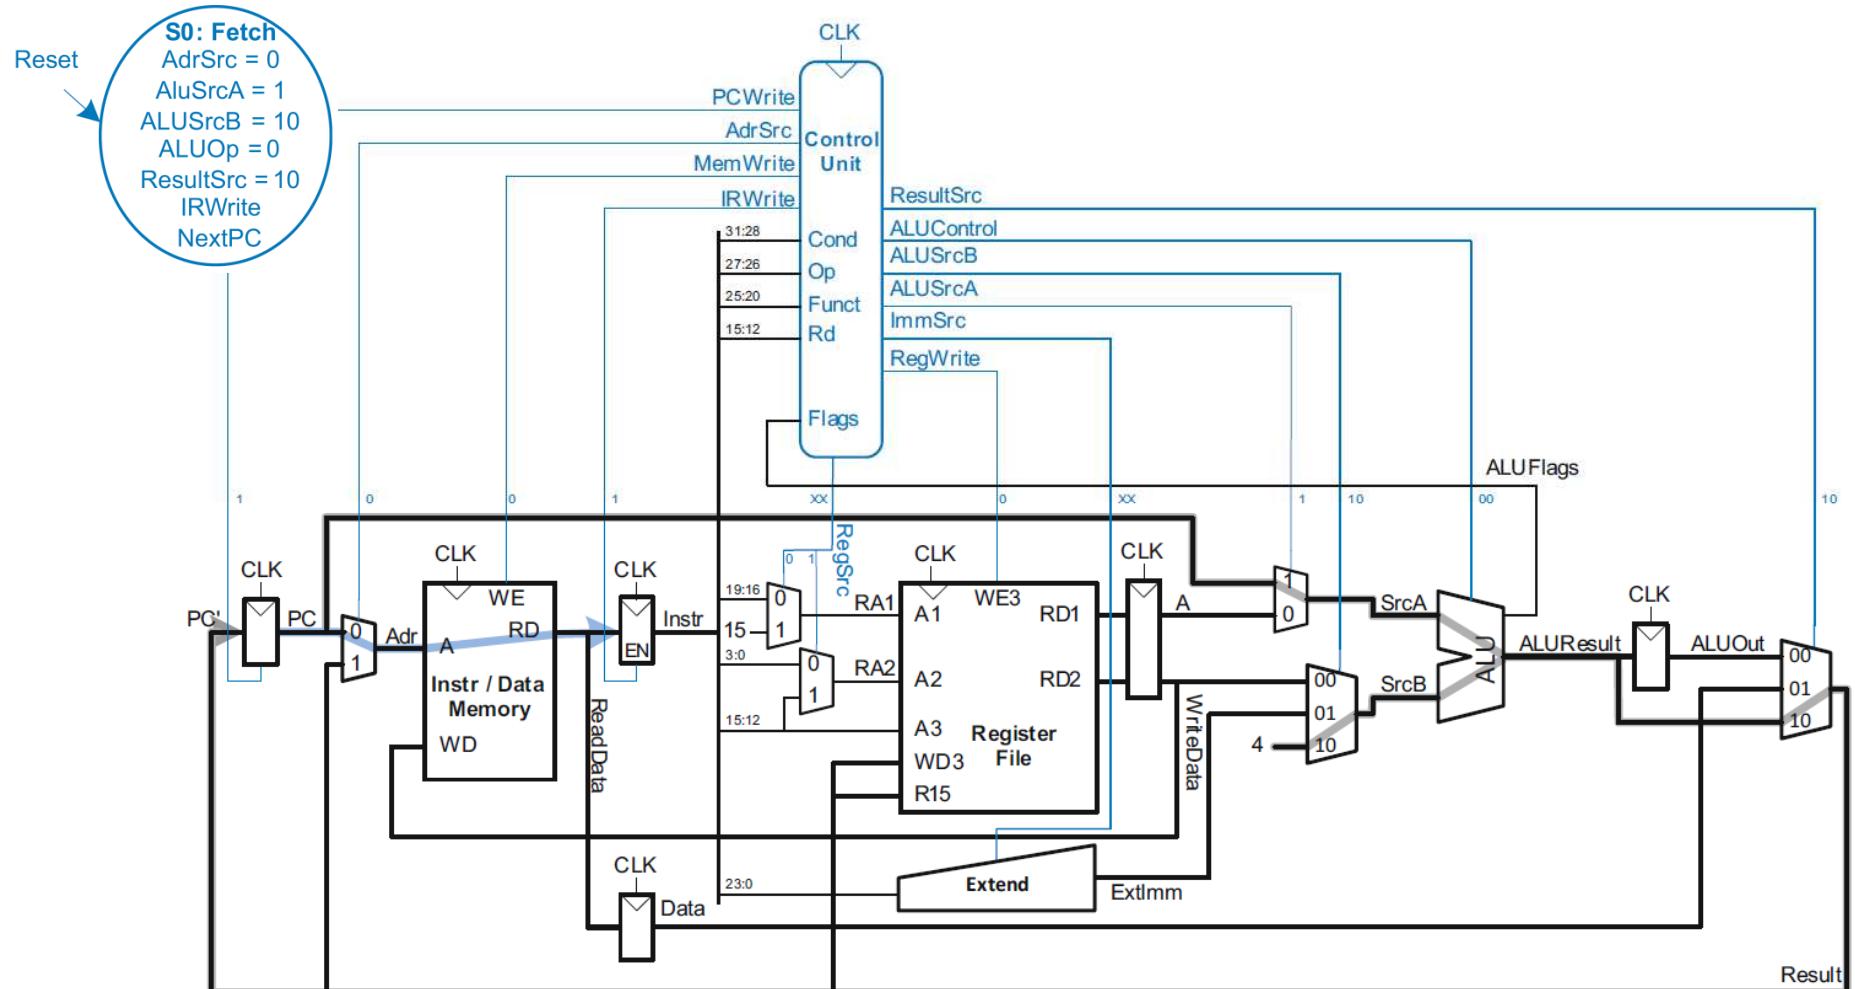
# Multicycle ARM Processor



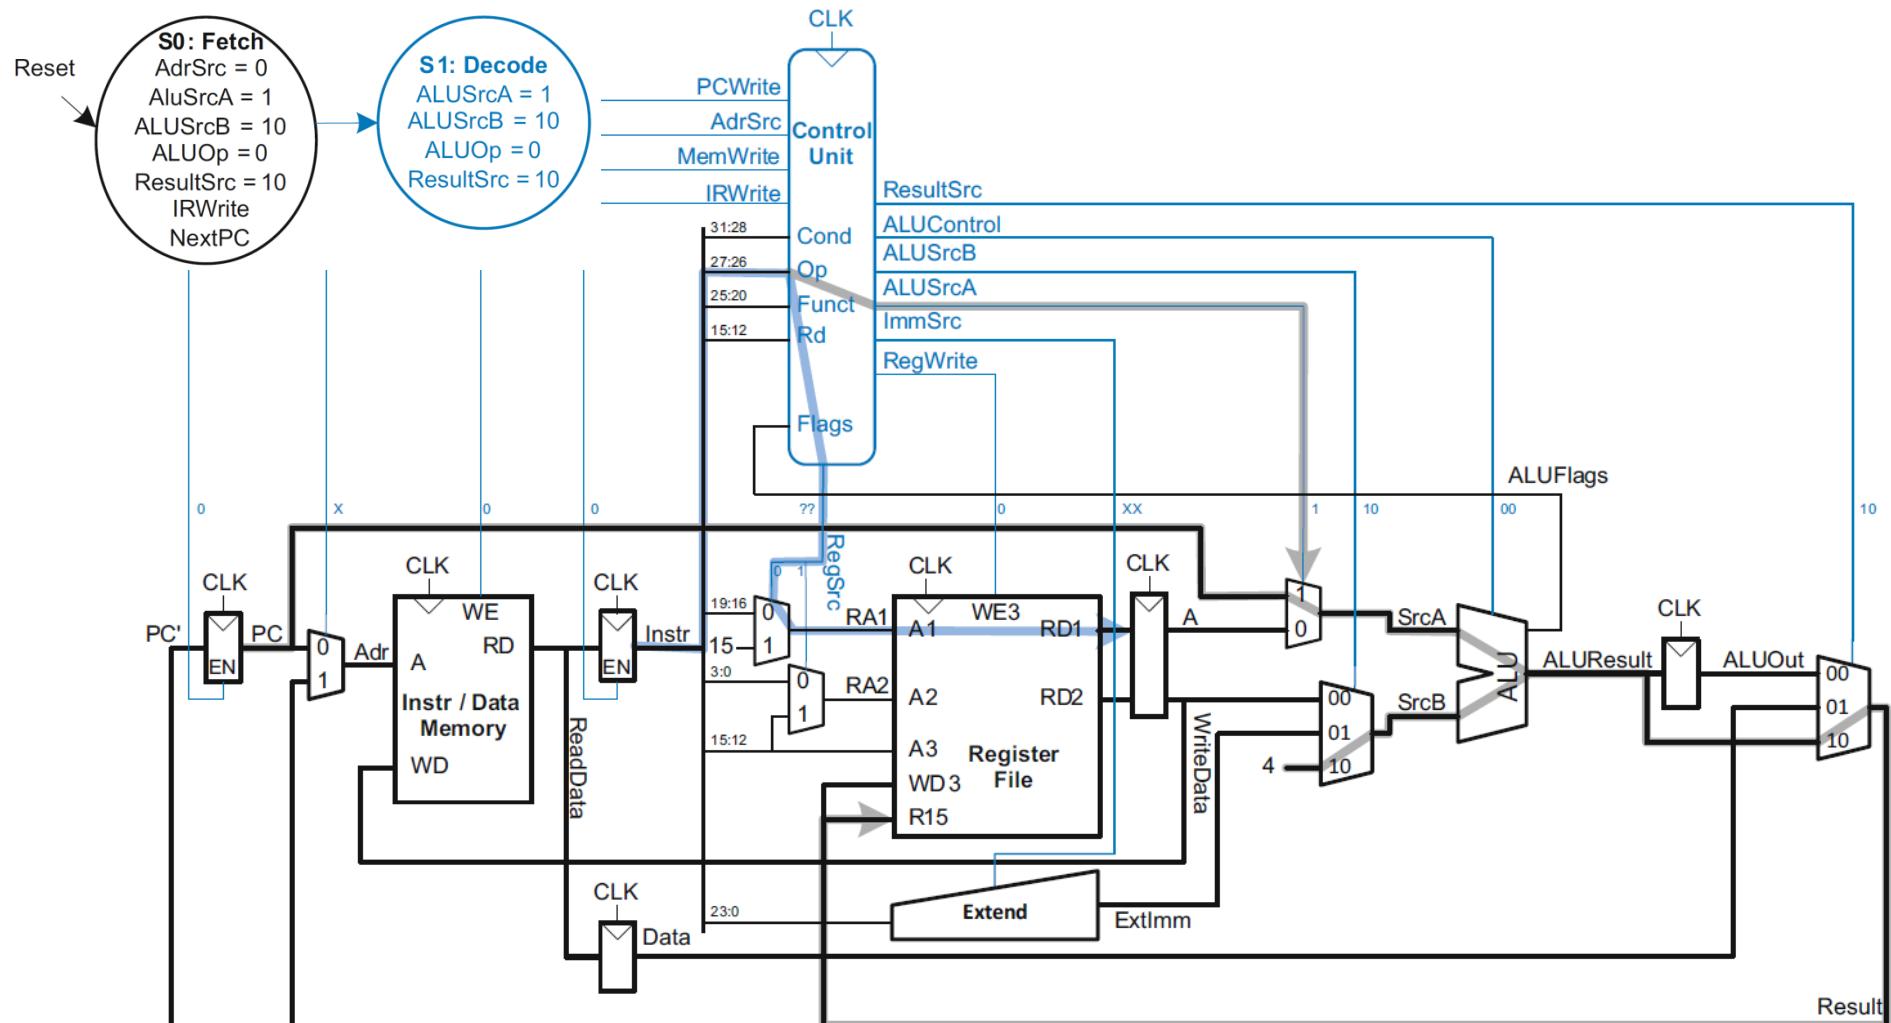
# Multicycle Control: Main FSM



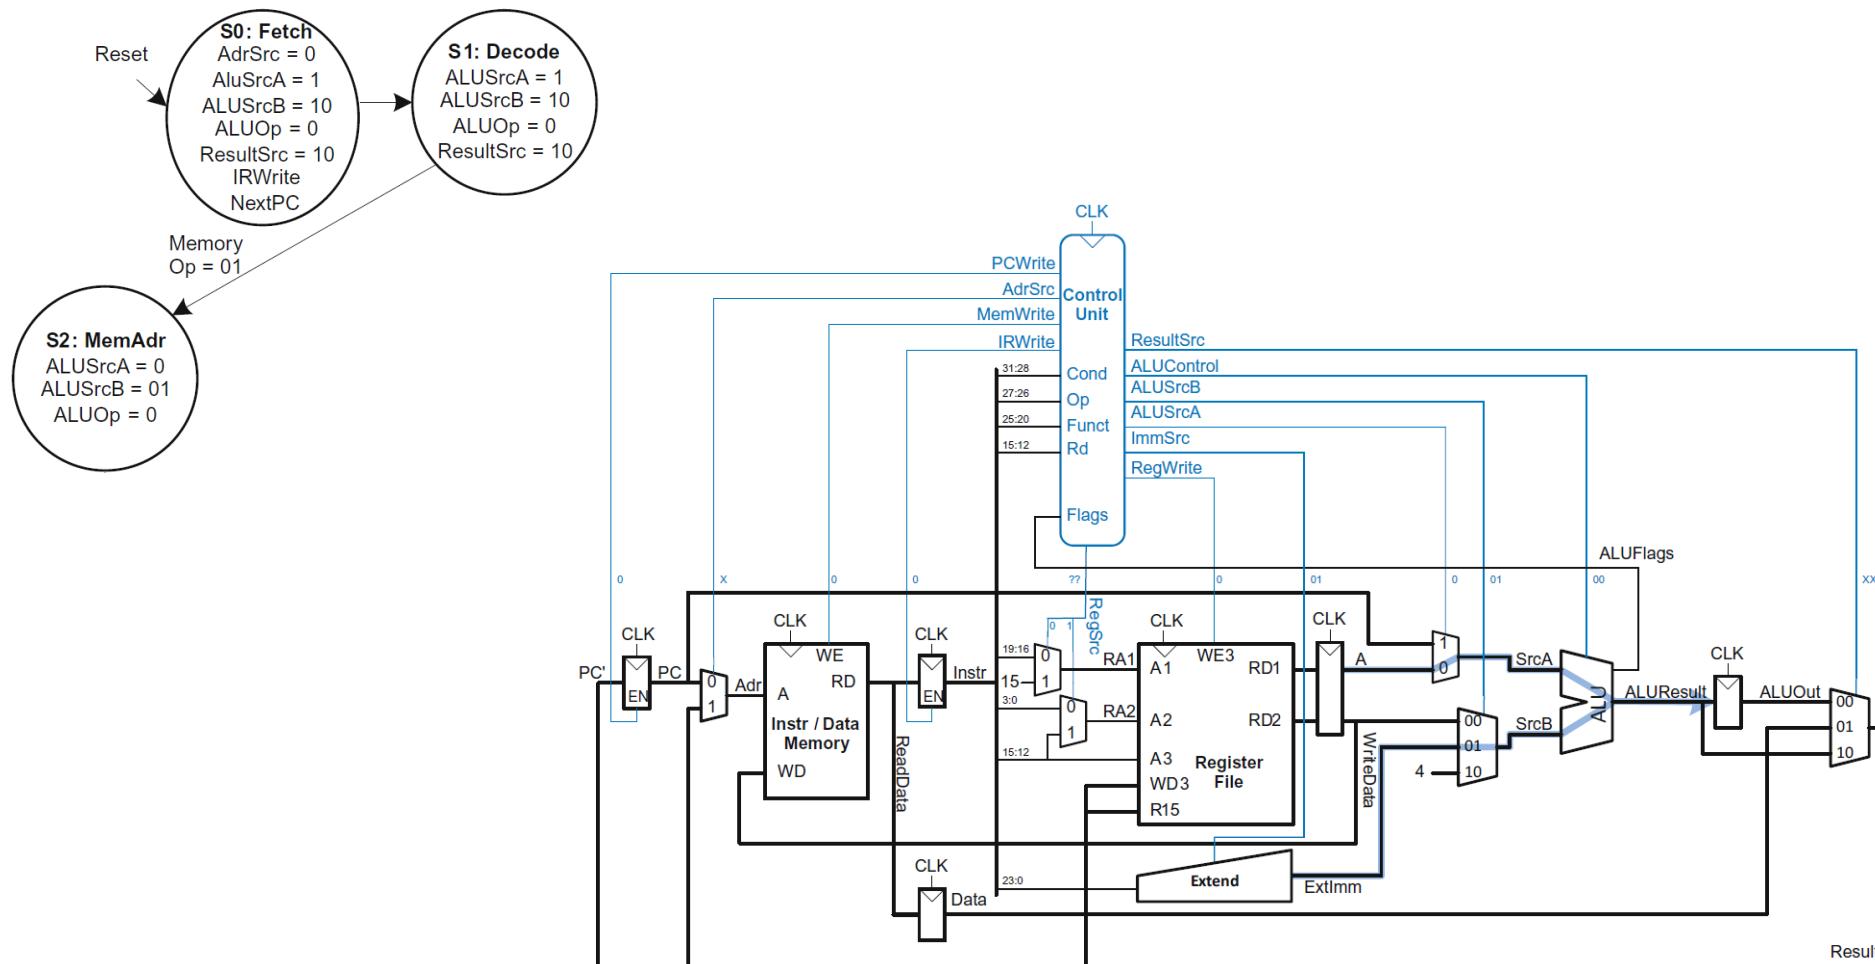
# Main Controller FSM: Fetch



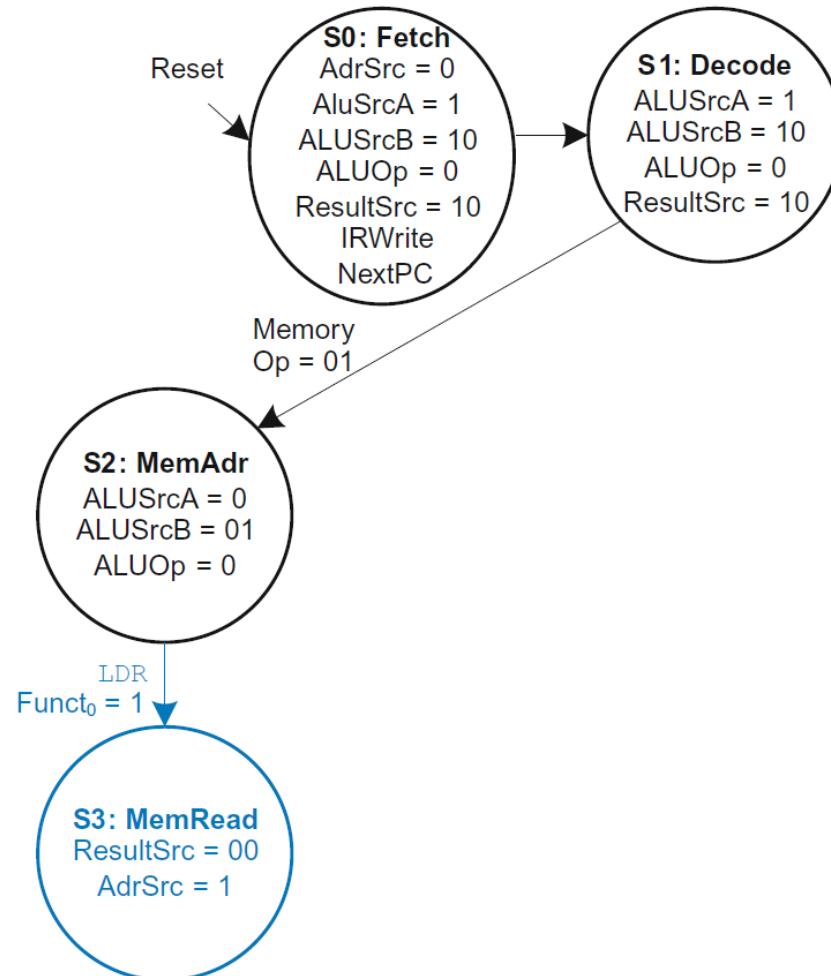
# Main Controller FSM: Decode



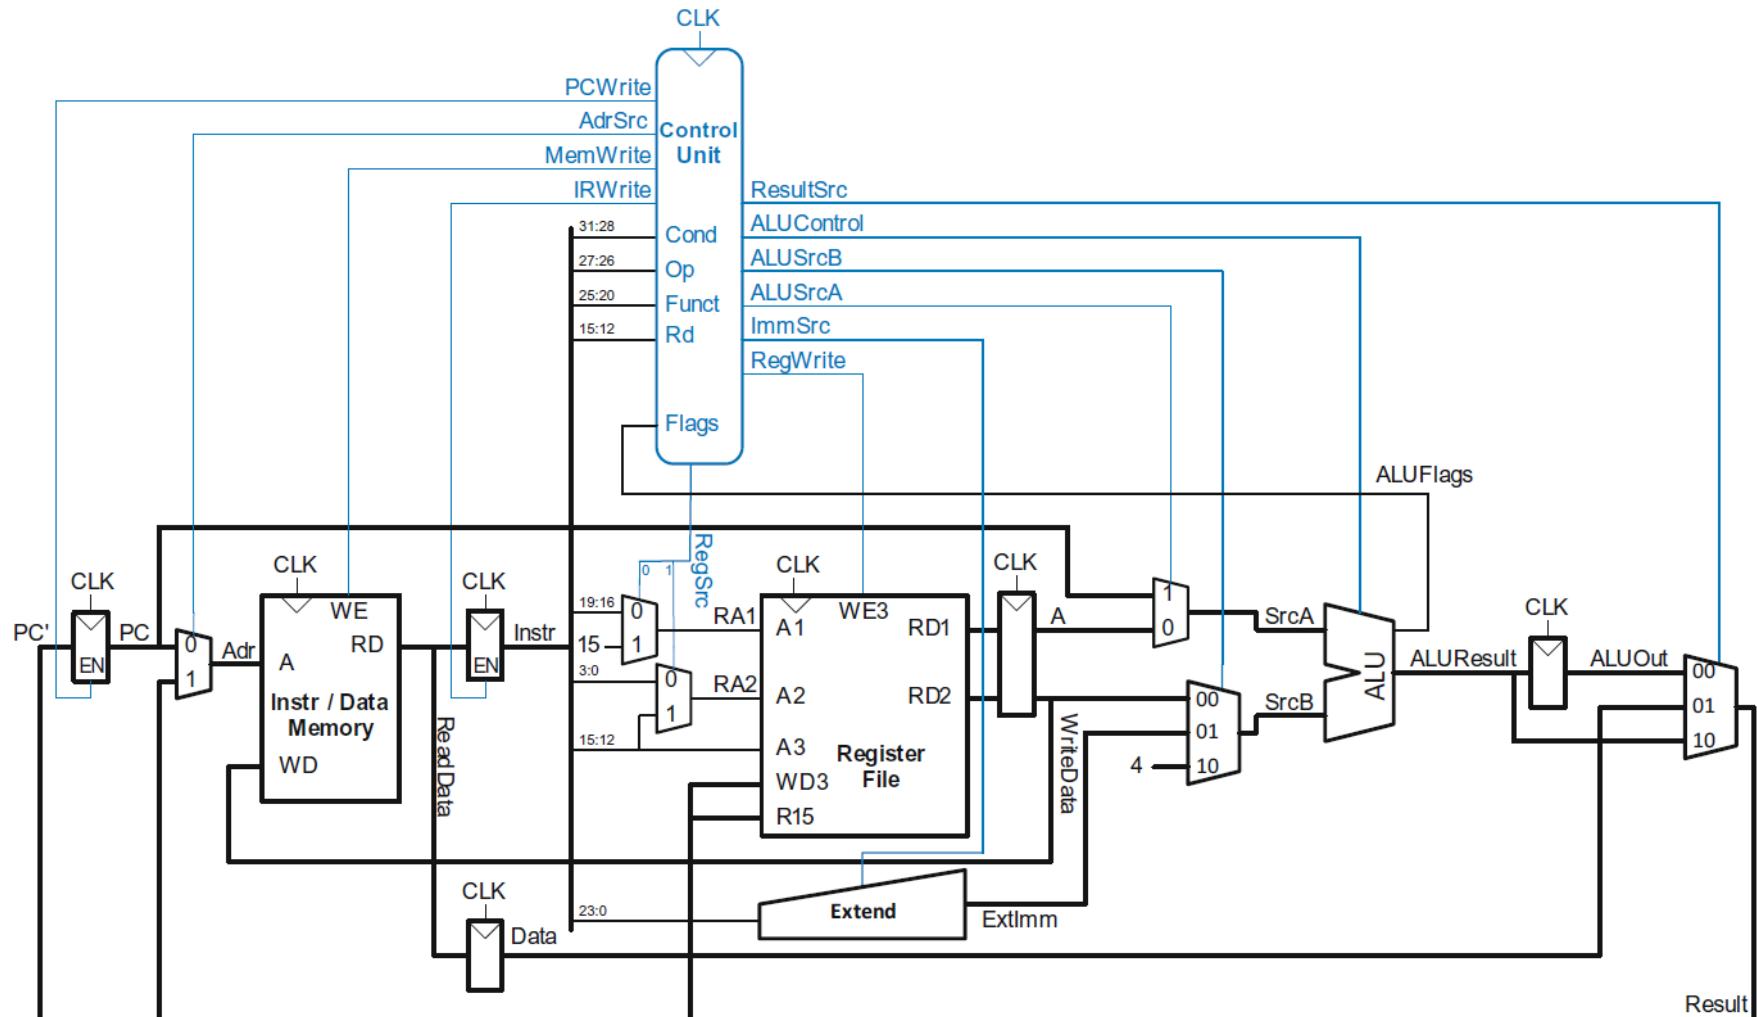
# Main Controller FSM: Address



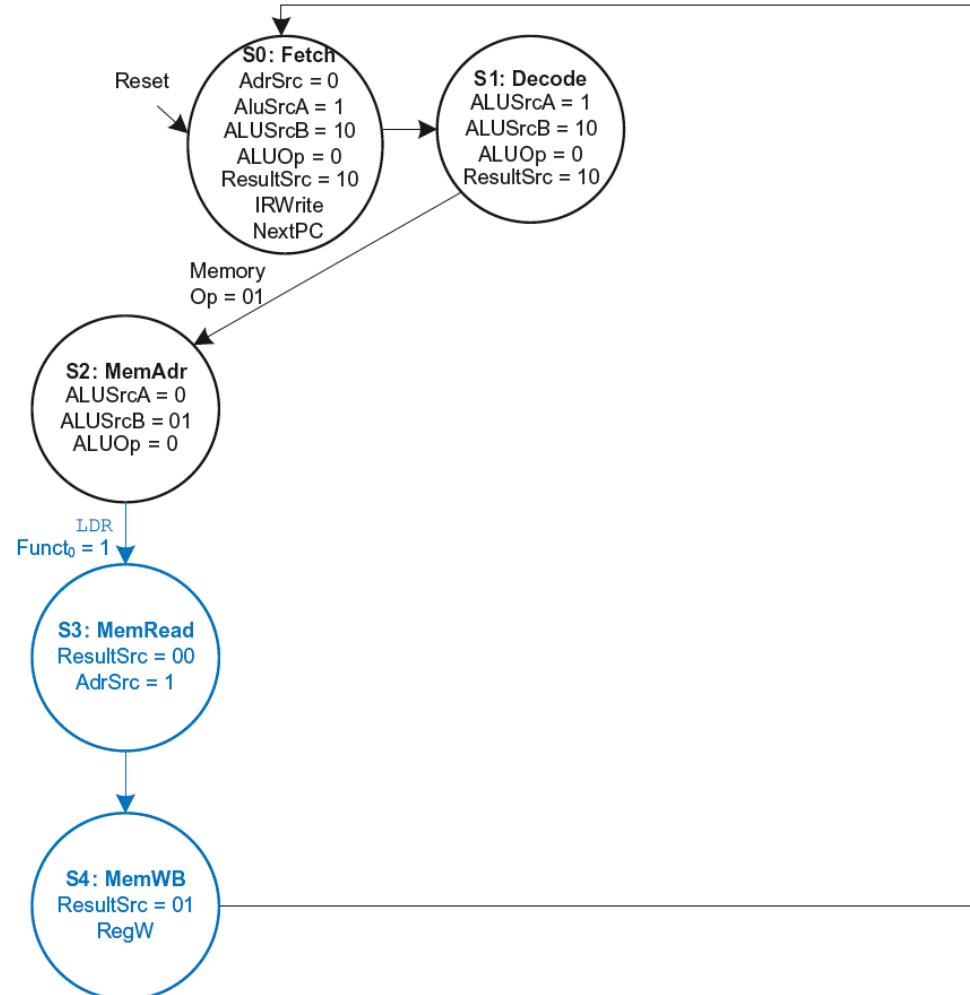
# Main Controller FSM: Read Memory



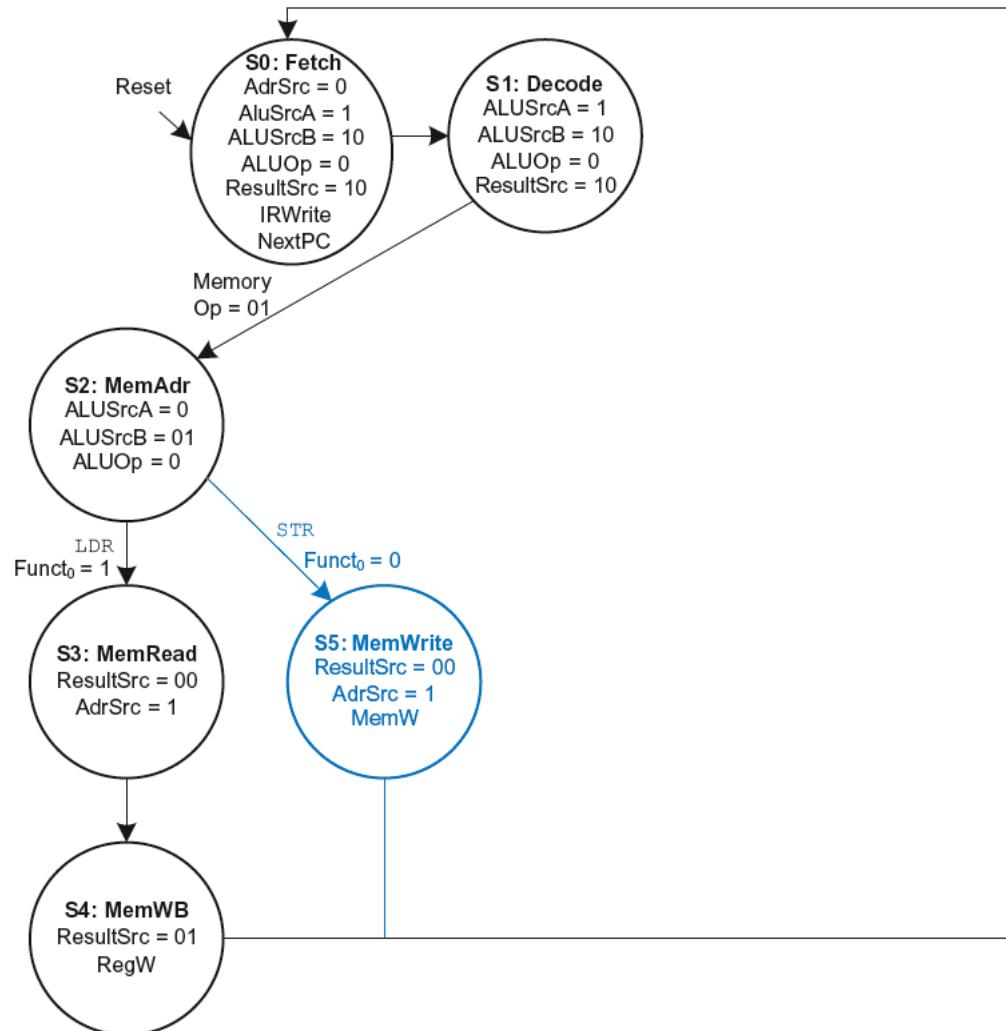
# Multicycle ARM Processor



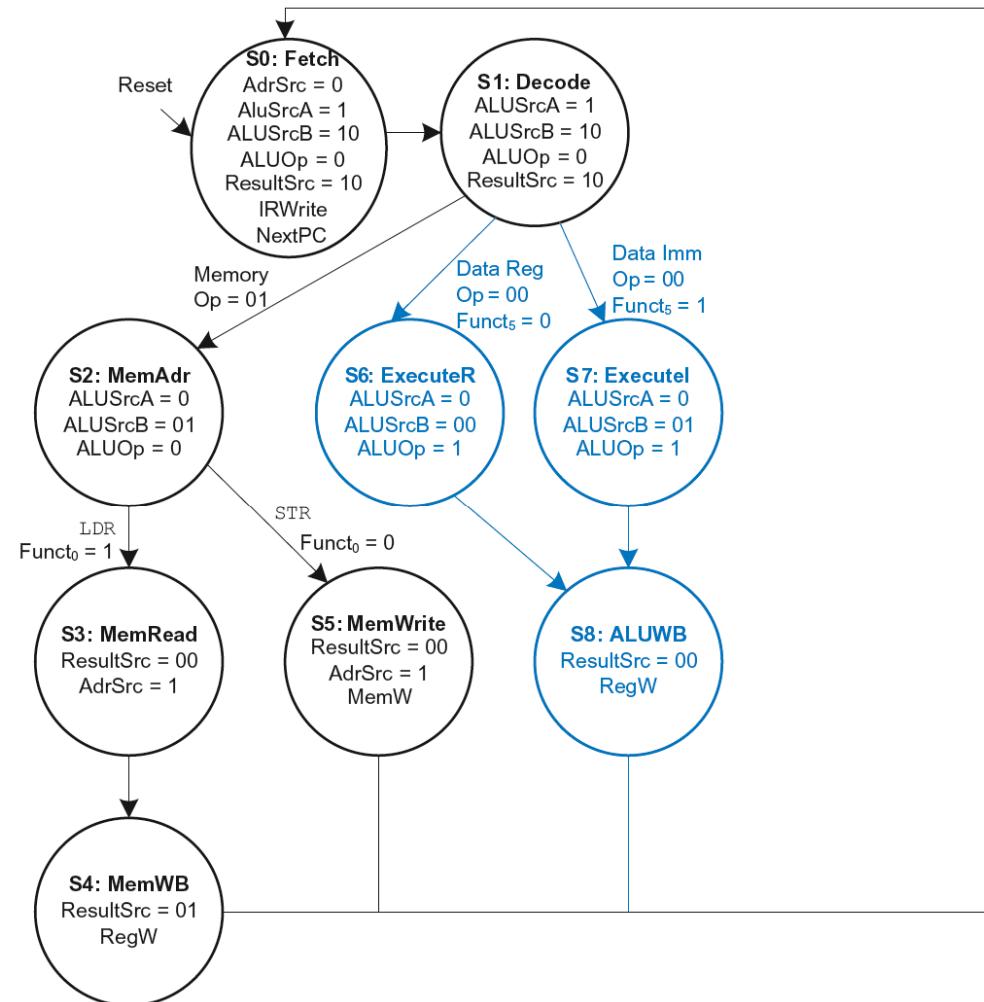
# Main Controller FSM: LDR



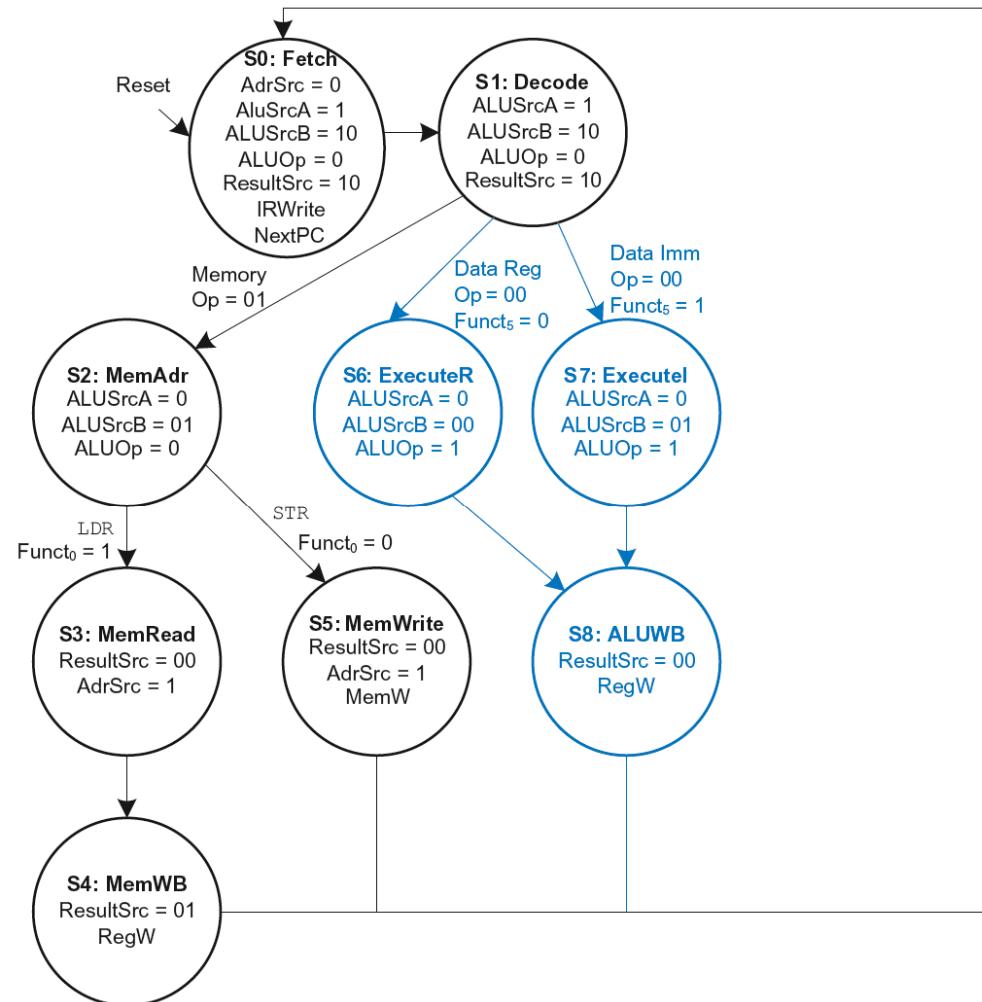
# Main Controller FSM: STR



# Main Controller FSM: Data-processing



# Main Controller FSM: Data-processing



# Multicycle Controller FSM

**State**

Fetch

Decode

MemAddr

MemRead

MemWB

MemWrite

ExecuteR

Executel

ALUWB

Branch

**Datapath μOp**

Instr  $\leftarrow$  Mem[PC]; PC  $\leftarrow$  PC+4

ALUOut  $\leftarrow$  PC+4

ALUOut  $\leftarrow$  Rn + Imm

Data  $\leftarrow$  Mem[ALUOut]

Rd  $\leftarrow$  Data

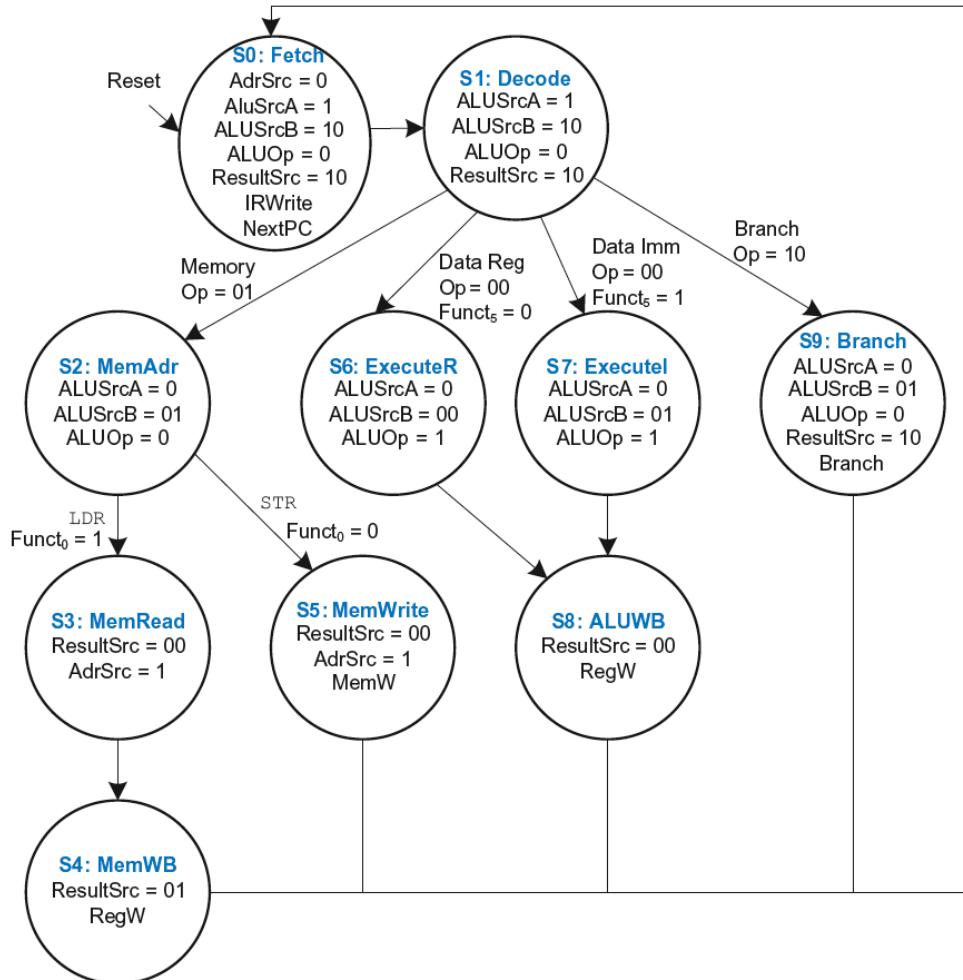
Mem[ALUOut]  $\leftarrow$  Rd

ALUOut  $\leftarrow$  Rn op Rm

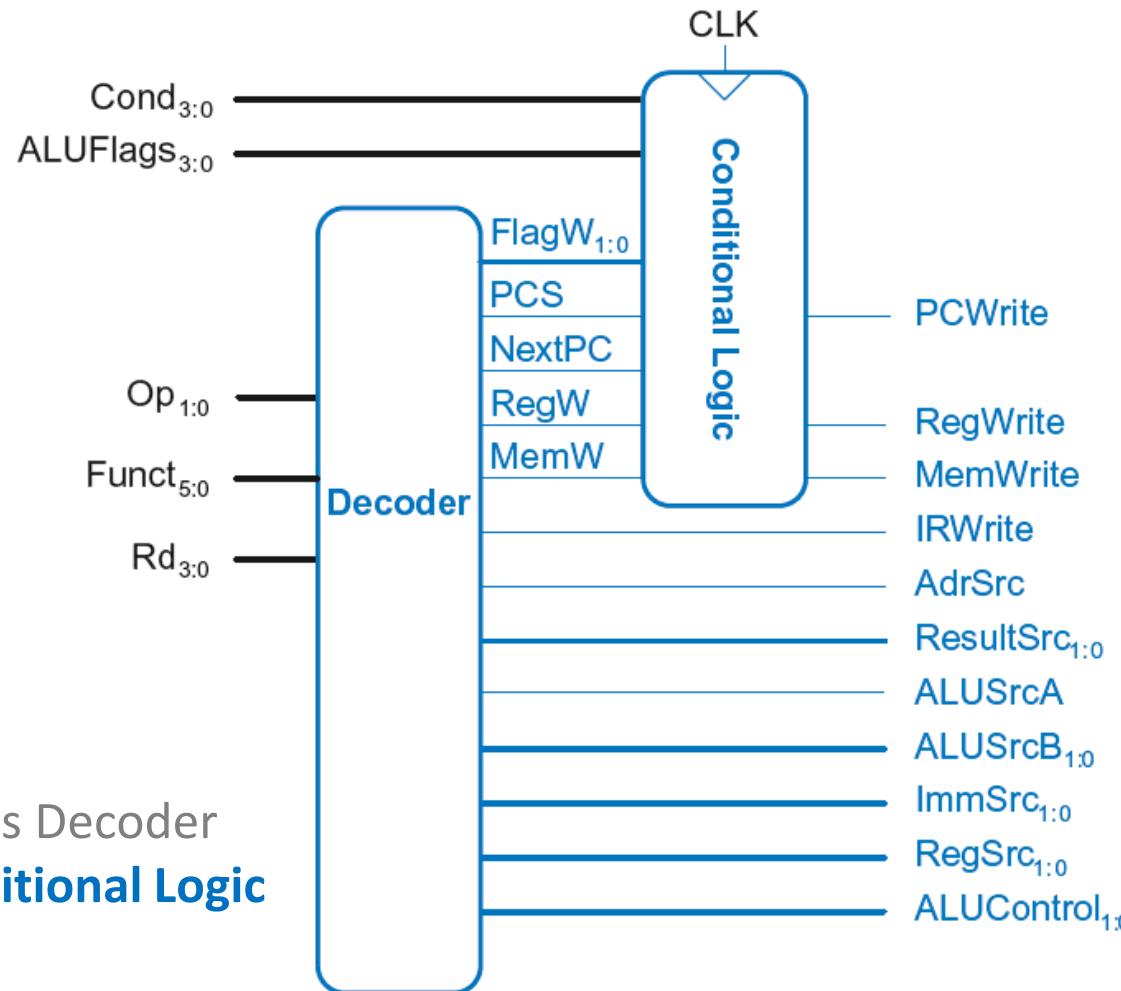
ALUOut  $\leftarrow$  Rn op Imm

Rd  $\leftarrow$  ALUOut

PC  $\leftarrow$  R15 + offset



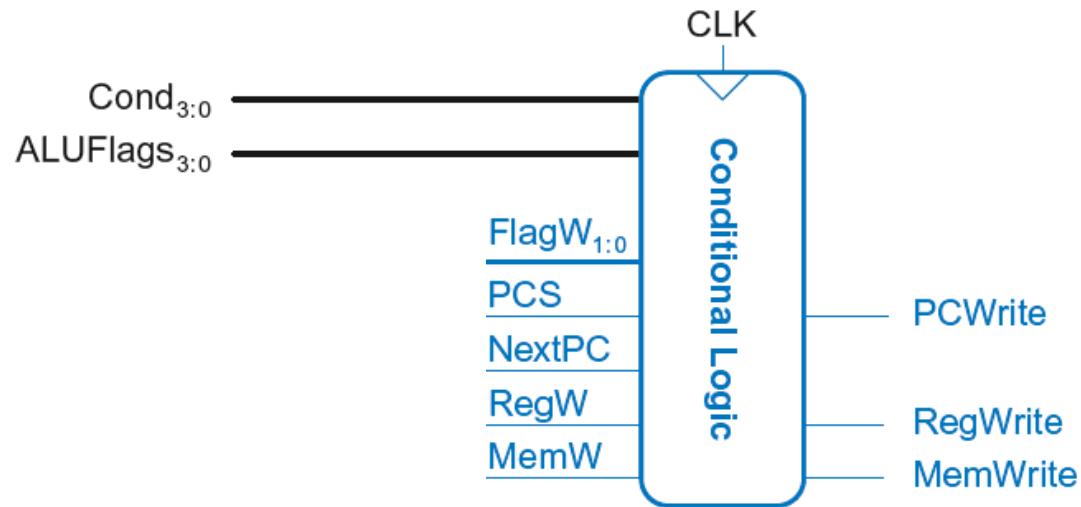
# Multicycle Control



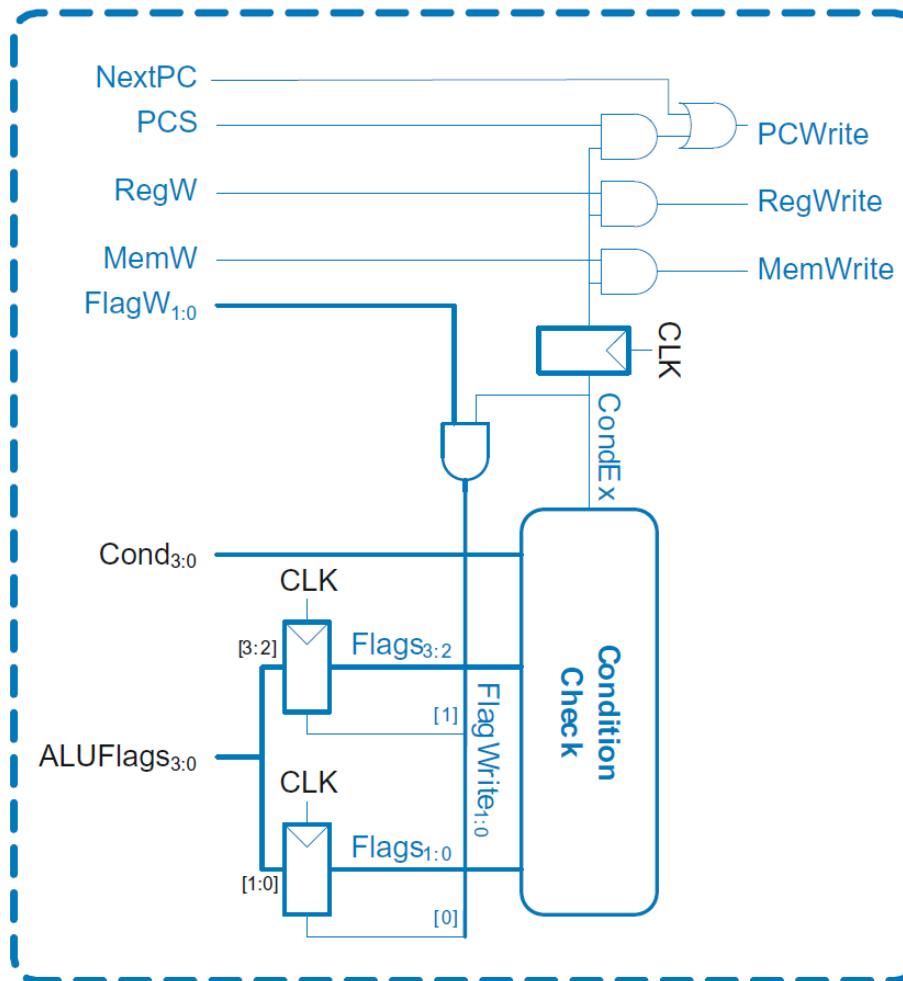
- First, discuss Decoder
- Then, **Conditional Logic**



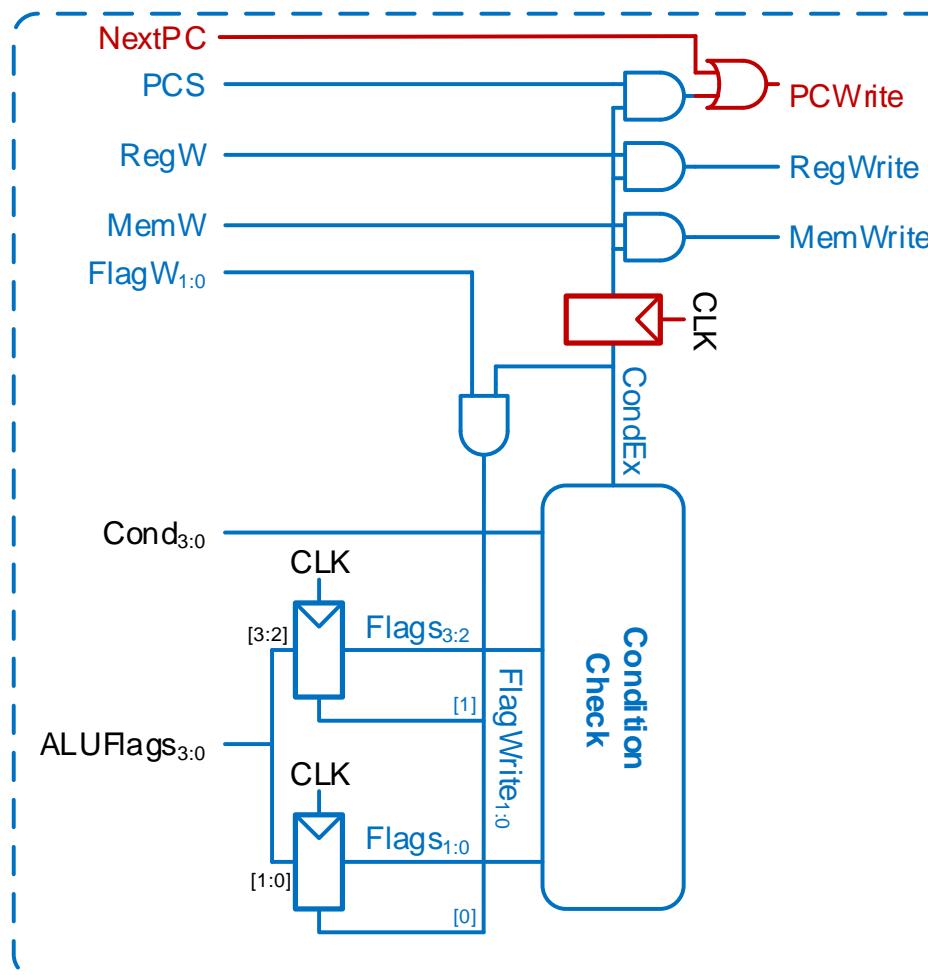
# Multicycle Control: Cond. Logic



# Single-Cycle Conditional Logic



# Multicycle Conditional Logic



- **PCWrite** asserted in Fetch state
- **ExecuteL/ExecuteR** state:  
*CondEx* asserts  
*ALUFlags* generated
- **ALUWB** state:  
*Flags* updated  
*CondEx* changes  
*PCWrite*, *RegWrite*, and  
*MemWrite* don't see  
change till new  
instruction (Fetch state)



# Multicycle Processor Performance

- Instructions take different number of cycles.

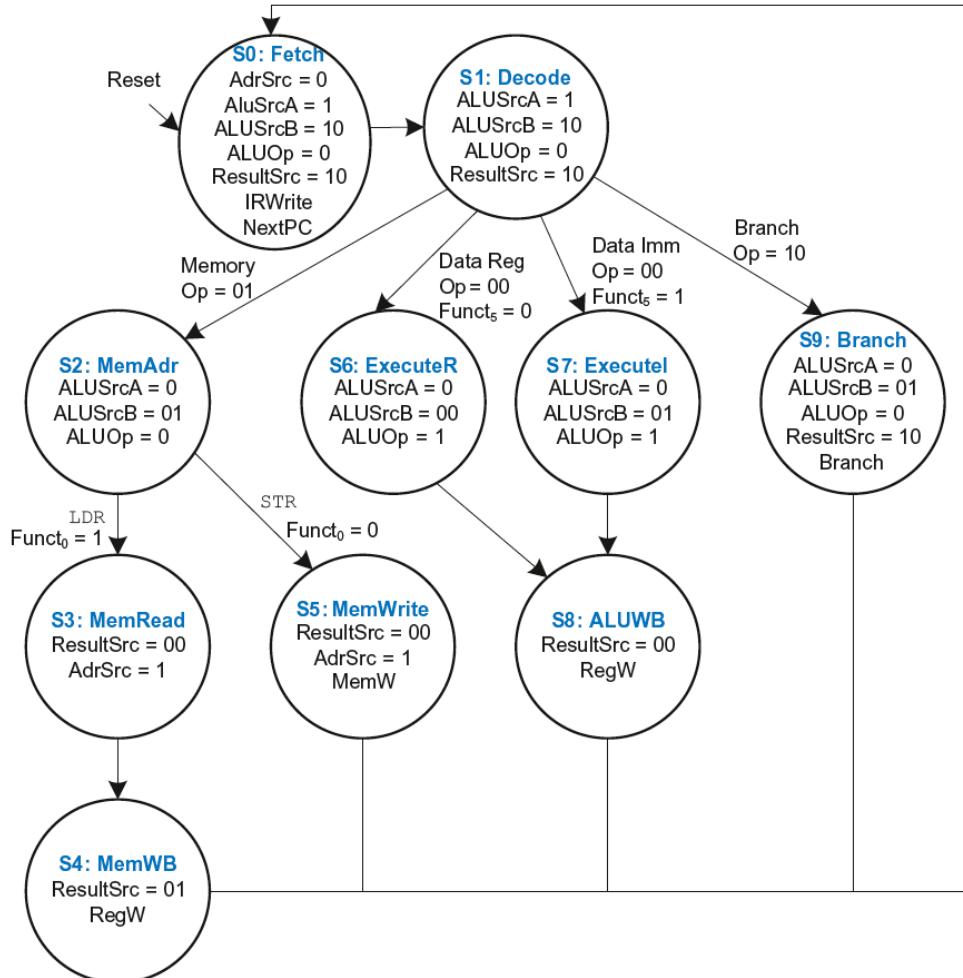


# Multicycle Controller FSM

State
Fetch
Decode
MemAddr
MemRead
MemWB
MemWrite
ExecuteR
Executel
ALUWB
Branch

## Datapath $\mu$ Op

Instr  $\leftarrow$  Mem[PC]; PC  $\leftarrow$  PC+4  
 ALUOut  $\leftarrow$  PC+4  
 ALUOut  $\leftarrow$  Rn + Imm  
 Data  $\leftarrow$  Mem[ALUOut]  
 Rd  $\leftarrow$  Data  
 Mem[ALUOut]  $\leftarrow$  Rd  
 ALUOut  $\leftarrow$  Rn op Rm  
 ALUOut  $\leftarrow$  Rn op Imm  
 Rd  $\leftarrow$  ALUOut  
 PC  $\leftarrow$  R15 + offset



# Multicycle Processor Performance

- Instructions take different number of cycles:
  - 3 cycles:
  - 4 cycles:
  - 5 cycles:



# Multicycle Processor Performance

- Instructions take different number of cycles:
  - 3 cycles: B
  - 4 cycles: DP, STR
  - 5 cycles: LDR



# Multicycle Processor Performance

- Instructions take different number of cycles:
  - 3 cycles: B
  - 4 cycles: DP, STR
  - 5 cycles: LDR
- CPI is weighted average
- SPECINT2000 benchmark:
  - 25% loads
  - 10% stores
  - 13% branches
  - 52% R-type



# Multicycle Processor Performance

- Instructions take different number of cycles:
  - 3 cycles: B
  - 4 cycles: DP, STR
  - 5 cycles: LDR
- CPI is weighted average
- SPECINT2000 benchmark:
  - **25%** loads
  - **10%** stores
  - **13%** branches
  - **52%** R-type

$$\text{Average CPI} = (0.13)(3) + (0.52 + 0.10)(4) + (0.25)(5) = 4.12$$



# Multicycle Processor Performance

Multicycle critical path:

- Assumptions:
  - RF is faster than memory
  - writing memory is faster than reading memory

$$T_{c2} = t_{pcq} + 2t_{\text{mux}} + \max(t_{\text{ALU}} + t_{\text{mux}}, t_{\text{mem}}) + t_{\text{setup}}$$



# Multicycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	40
Register setup	$t_{\text{setup}}$	50
Multiplexer	$t_{\text{mux}}$	25
ALU	$t_{\text{ALU}}$	120
Decoder	$t_{\text{dec}}$	70
Memory read	$t_{\text{mem}}$	200
Register file read	$t_{RF\text{read}}$	100
Register file setup	$t_{RF\text{setup}}$	60

$$T_{c2} = ?$$



# Multicycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	40
Register setup	$t_{\text{setup}}$	50
Multiplexer	$t_{\text{mux}}$	25
ALU	$t_{\text{ALU}}$	120
Decoder	$t_{\text{dec}}$	70
Memory read	$t_{\text{mem}}$	200
Register file read	$t_{RF\text{read}}$	100
Register file setup	$t_{RF\text{setup}}$	60

$$\begin{aligned}T_{c2} &= t_{pcq} + 2t_{\text{mux}} + \max[t_{\text{ALU}} + t_{\text{mux}}, t_{\text{mem}}] + t_{\text{setup}} \\&= [40 + 2(25) + 200 + 50] \text{ ps} = 340 \text{ ps}\end{aligned}$$



# Multicycle Performance Example

For a program with **100 billion** instructions executing on a **multicycle** ARM processor

- CPI = 4.12 cycles/instruction
- Clock cycle time:  $T_{c2} = 340 \text{ ps}$

**Execution Time = ?**



# Multicycle Performance Example

For a program with **100 billion** instructions executing on a **multicycle** ARM processor

- **CPI** = 4.12 cycles/instruction
- **Clock cycle time:**  $T_{c2} = 340 \text{ ps}$

$$\begin{aligned}\text{Execution Time} &= (\# \text{ instructions}) \times \text{CPI} \times T_c \\ &= (100 \times 10^9)(4.12)(340 \times 10^{-12}) \\ &= \mathbf{140 \text{ seconds}}\end{aligned}$$



# Multicycle Performance Example

For a program with **100 billion** instructions executing on a **multicycle** ARM processor

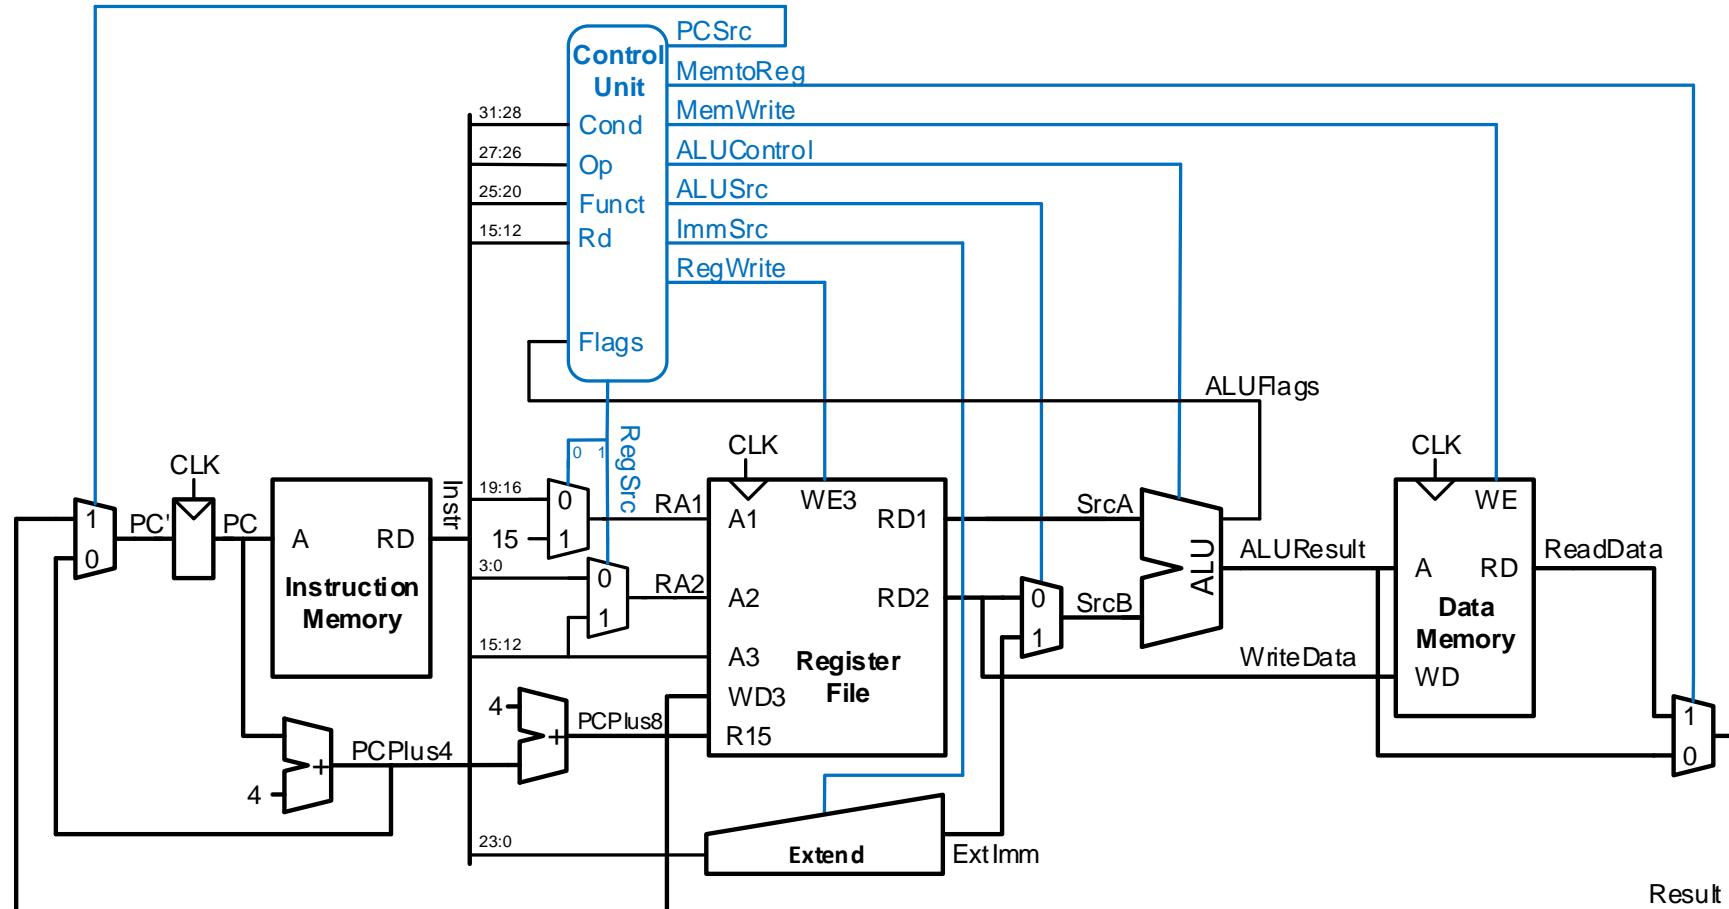
- CPI = 4.12 cycles/instruction
- Clock cycle time:  $T_{c2} = 340 \text{ ps}$

$$\begin{aligned}\text{Execution Time} &= (\# \text{ instructions}) \times \text{CPI} \times T_c \\ &= (100 \times 10^9)(4.12)(340 \times 10^{-12}) \\ &= \mathbf{140 \text{ seconds}}\end{aligned}$$

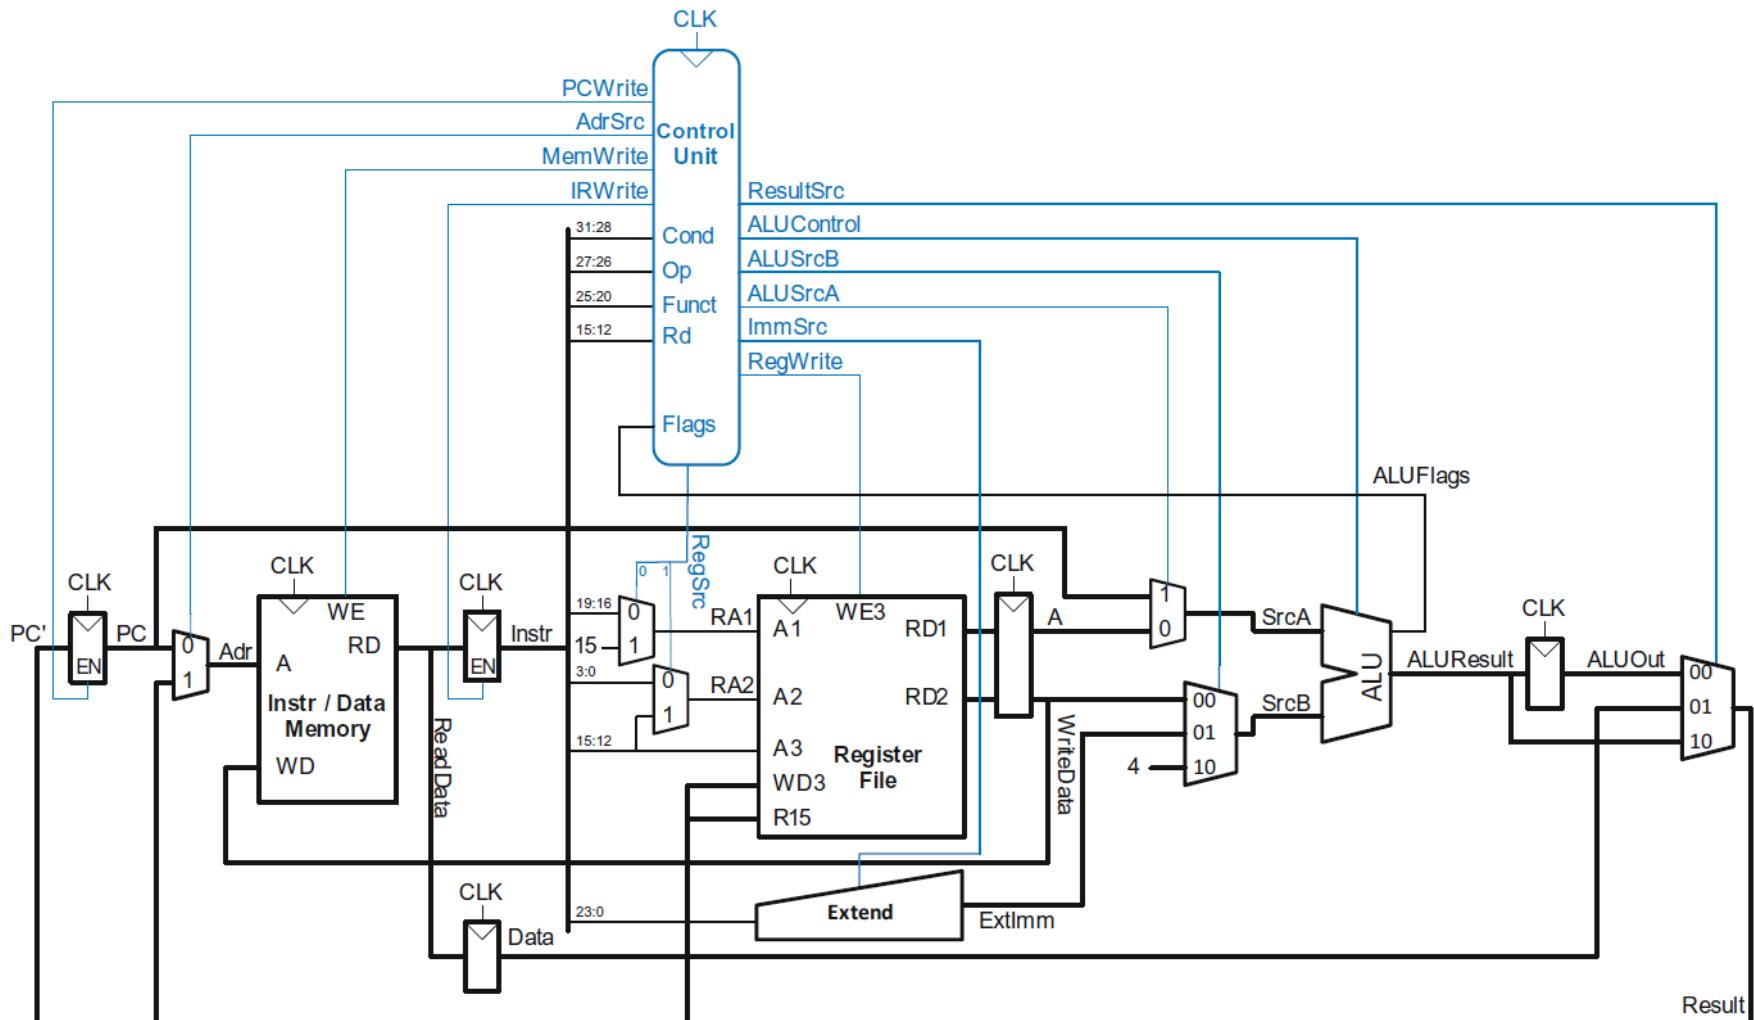
This is **slower** than the single-cycle processor (84 sec.)



# Review: Single-Cycle ARM Processor



# Review: Multicycle ARM Processor



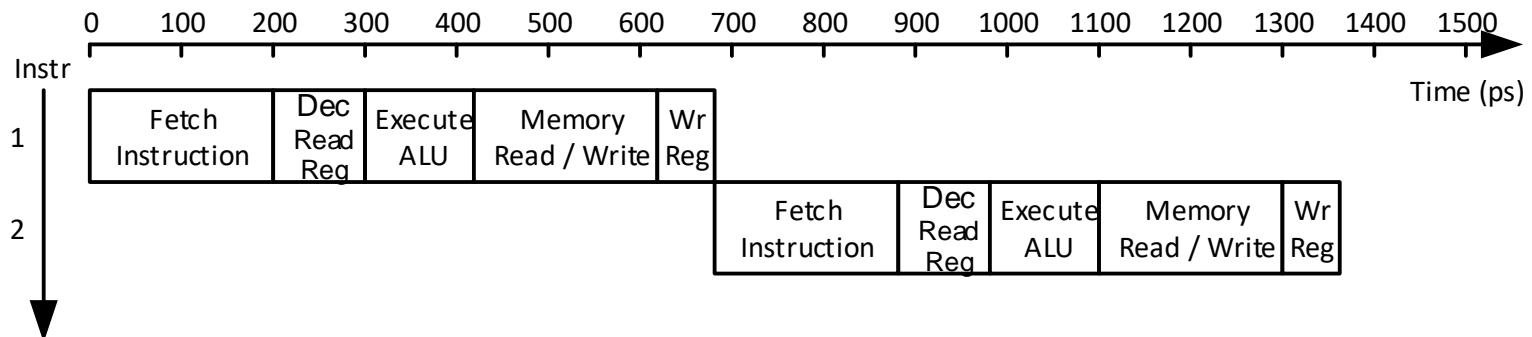
# Pipelined ARM Processor

- Temporal parallelism
- Divide single-cycle processor into 5 stages:
  - Fetch
  - Decode
  - Execute
  - Memory
  - Writeback
- Add pipeline registers between stages

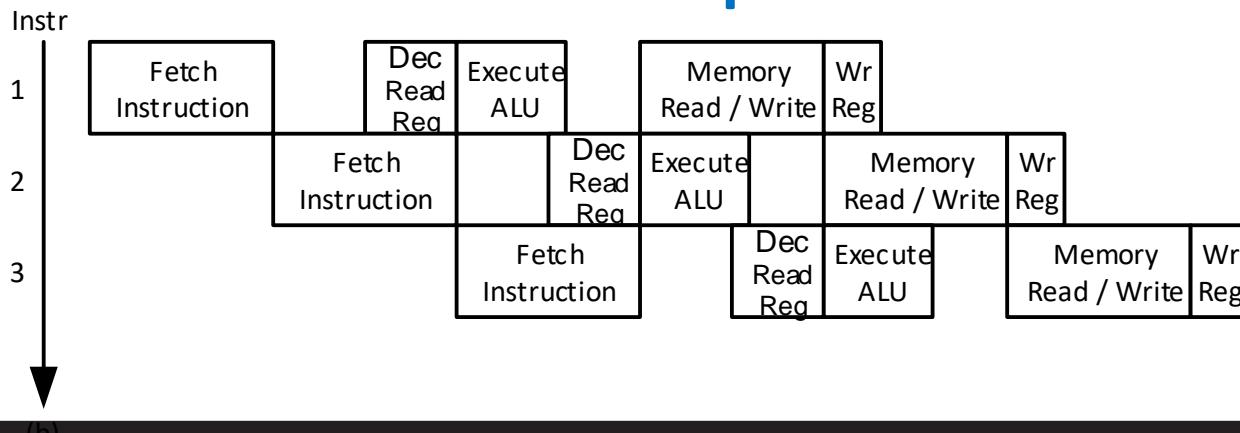


# Single-Cycle vs. Pipelined

## Single-Cycle



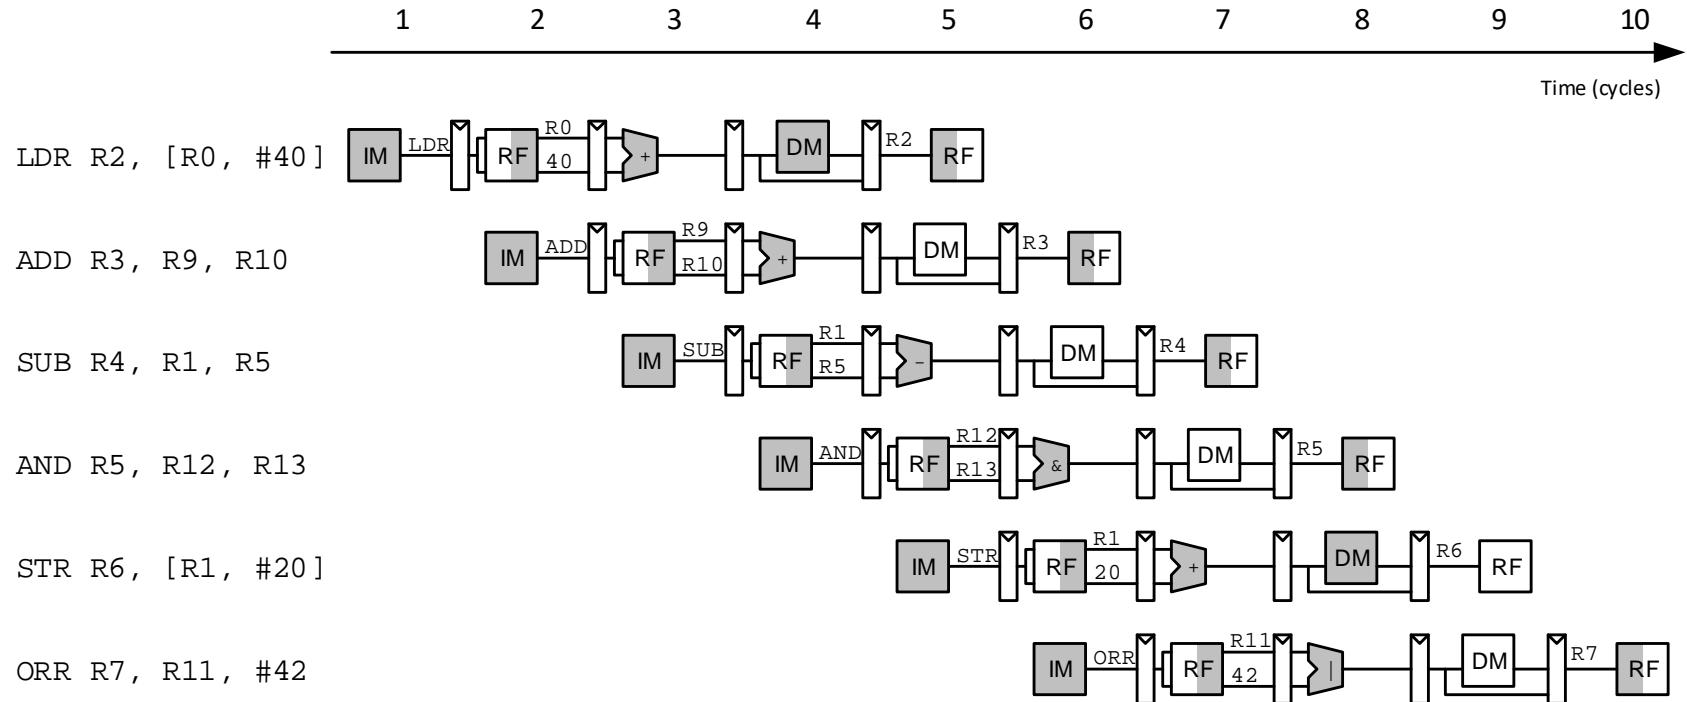
## Pipelined



(b)

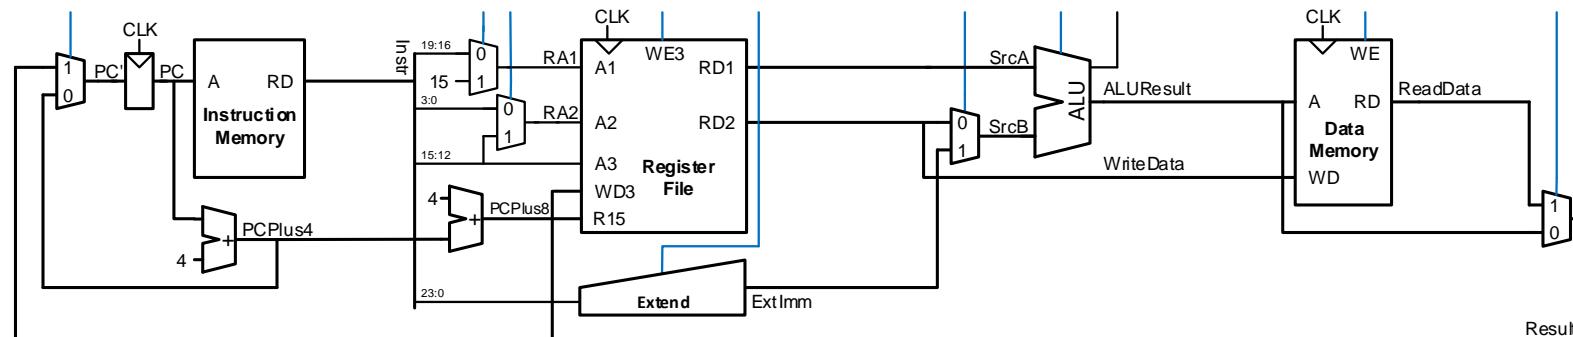


# Pipelined Processor Abstraction

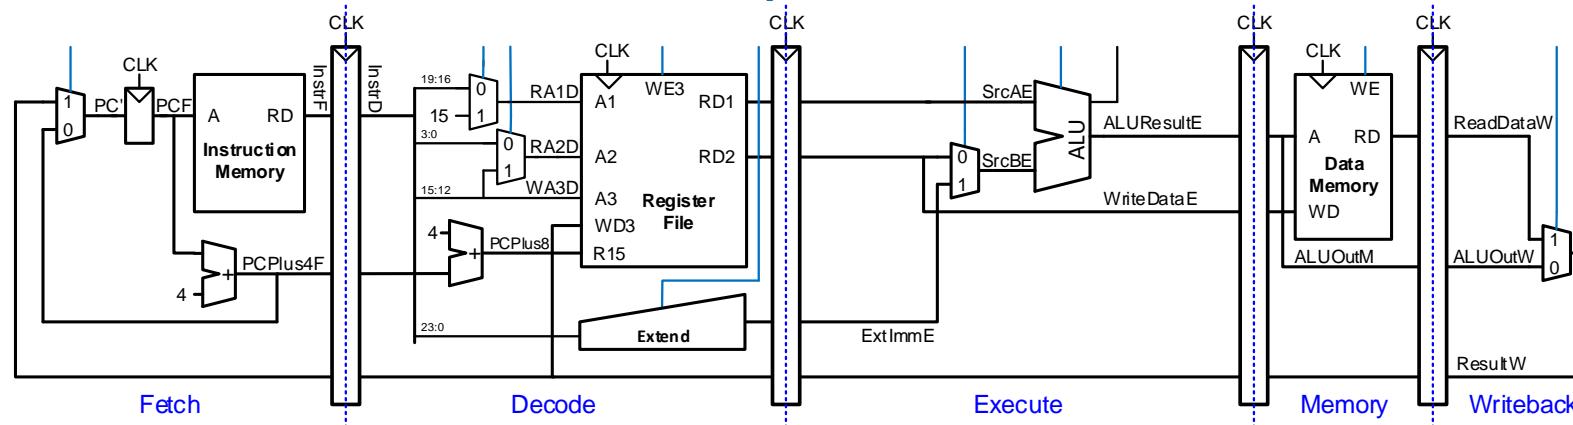


# Single-Cycle & Pipelined Datapath

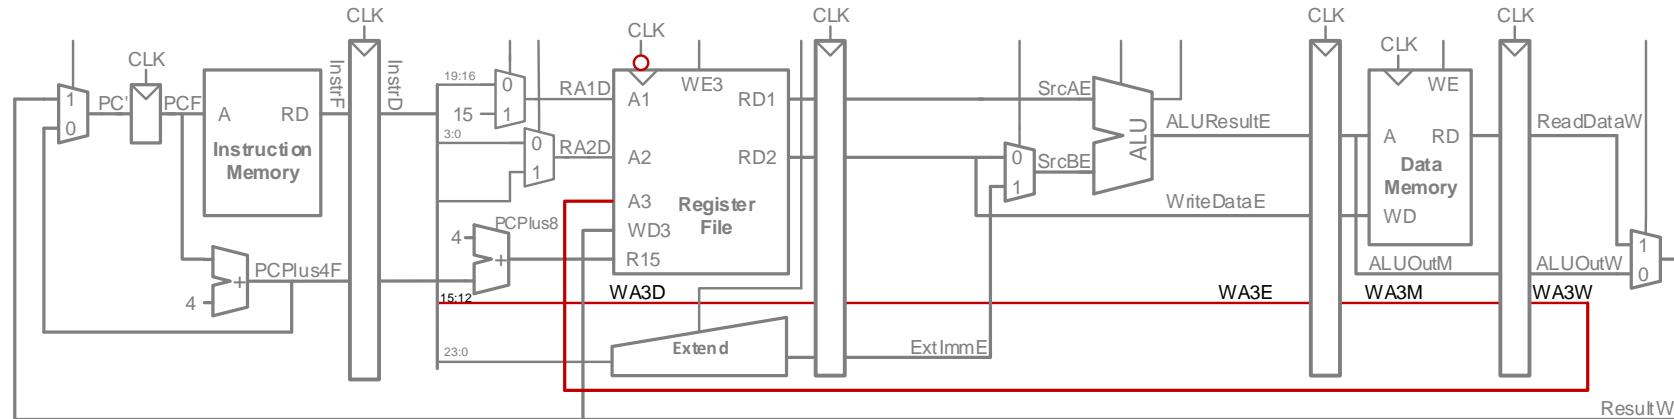
**Single-Cycle**



**Pipelined**



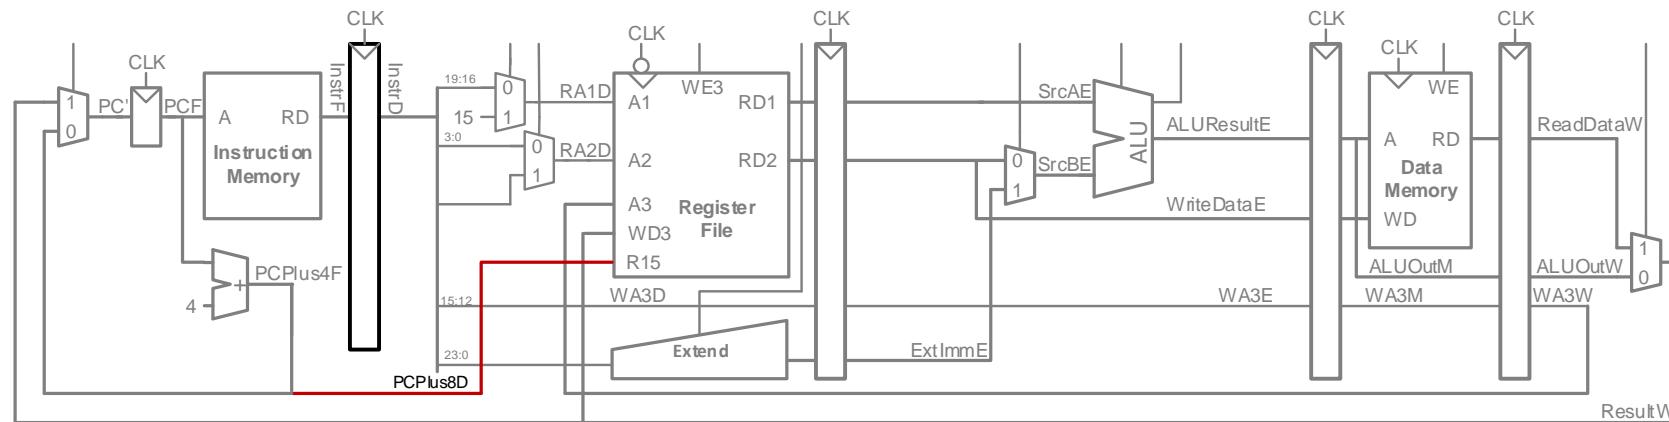
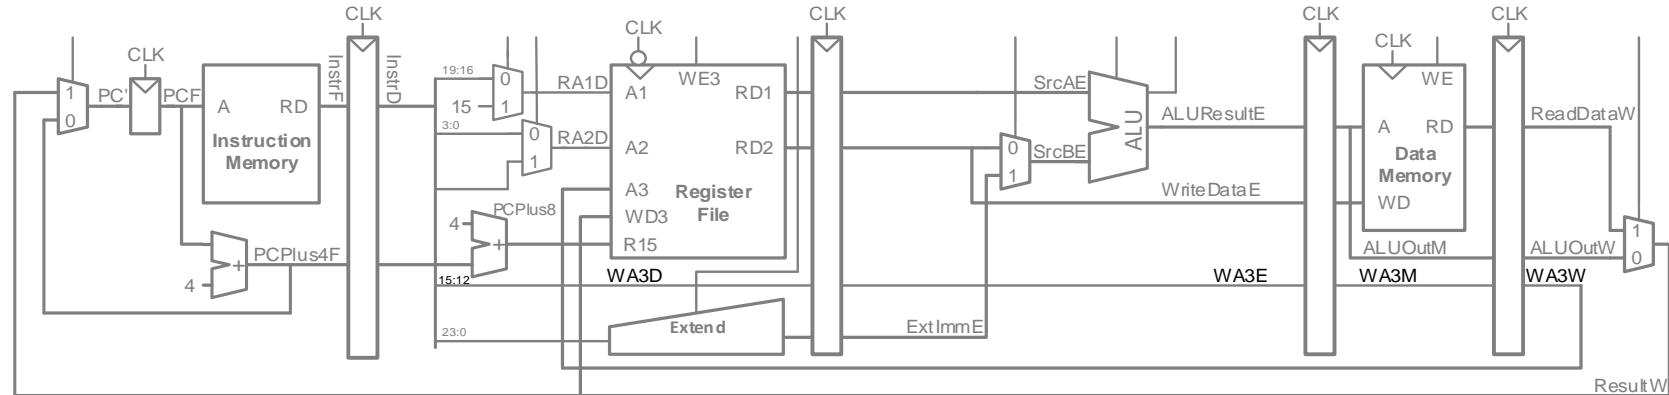
# Corrected Pipelined Datapath



- **WA3 must arrive at same time as *Result***
- **Register file written on falling edge of *CLK***



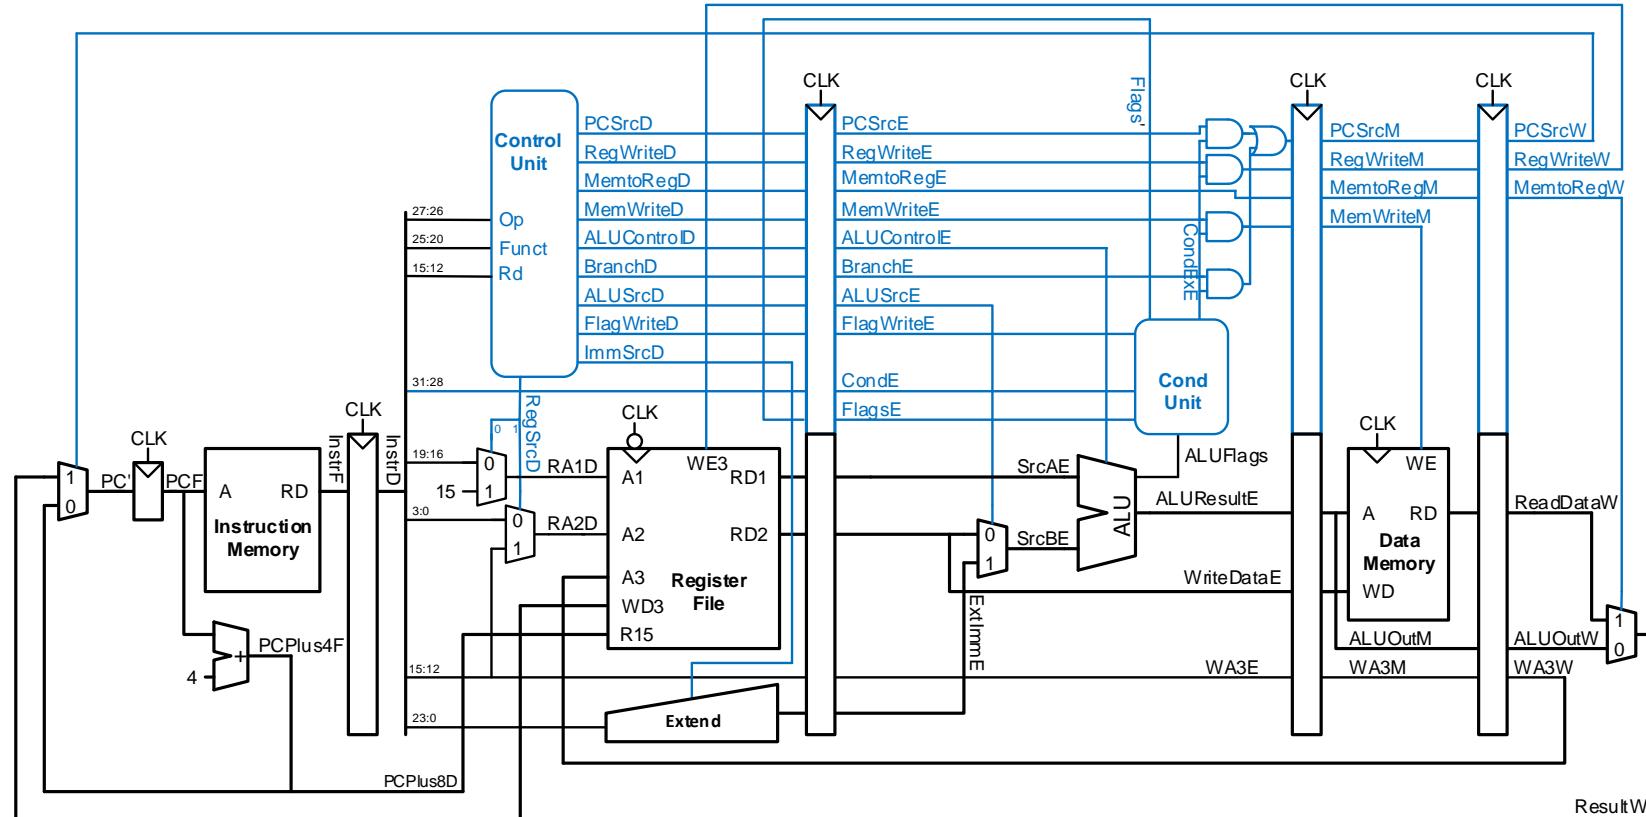
# Optimized Pipelined Datapath



**Remove adder by using  $PCPlus4F$  after  $PC$  has been updated to  $PC+4$**



# Pipelined Processor Control



- Same control unit as single-cycle processor
- Control delayed to proper pipeline stage

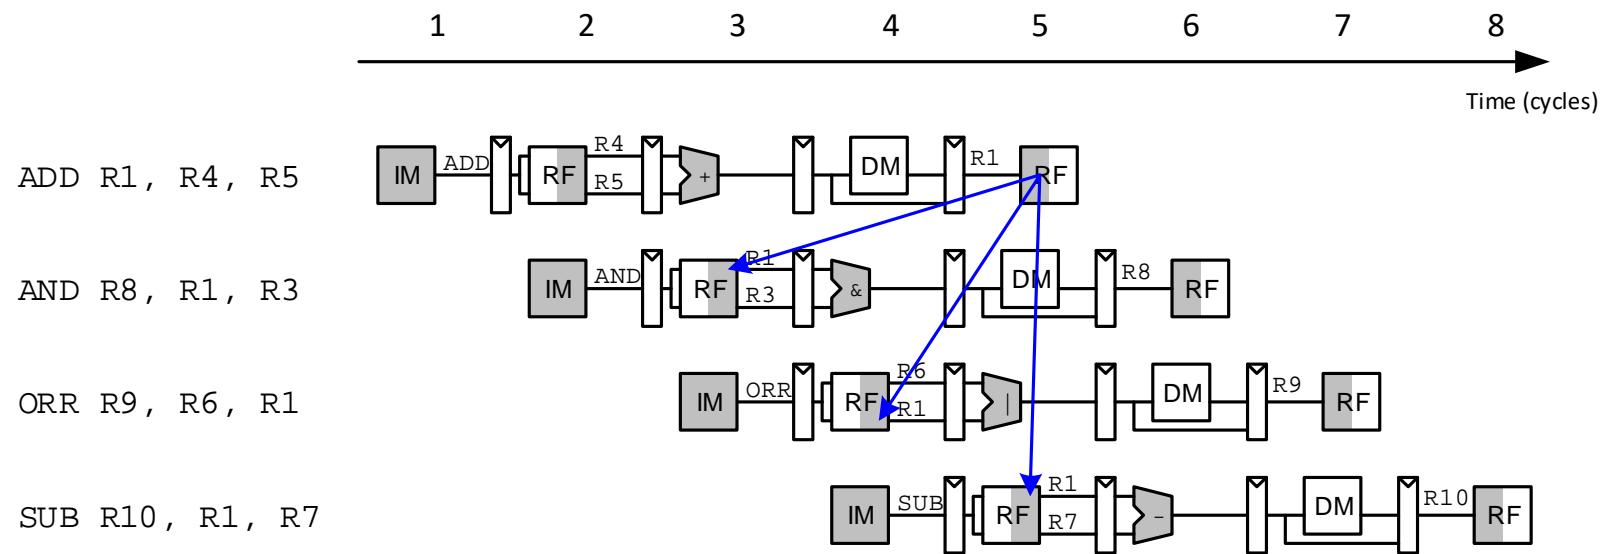


# Pipeline Hazards

- When an instruction depends on result from instruction that hasn't completed
- Types:
  - **Data hazard:** register value not yet written back to register file
  - **Control hazard:** next instruction not decided yet (caused by branch)



# Data Hazard



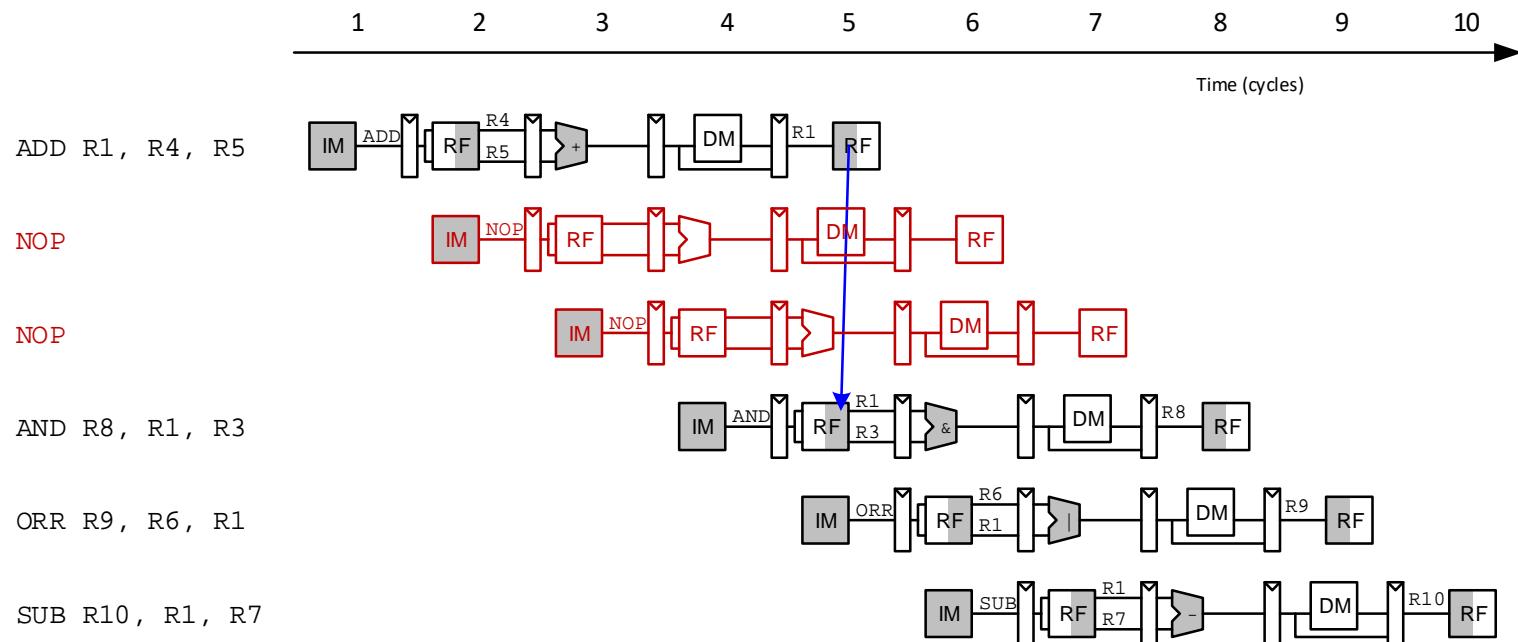
# Handling Data Hazards

- Insert NOPs in code at compile time
- Rearrange code at compile time
- Forward data at run time
- Stall the processor at run time

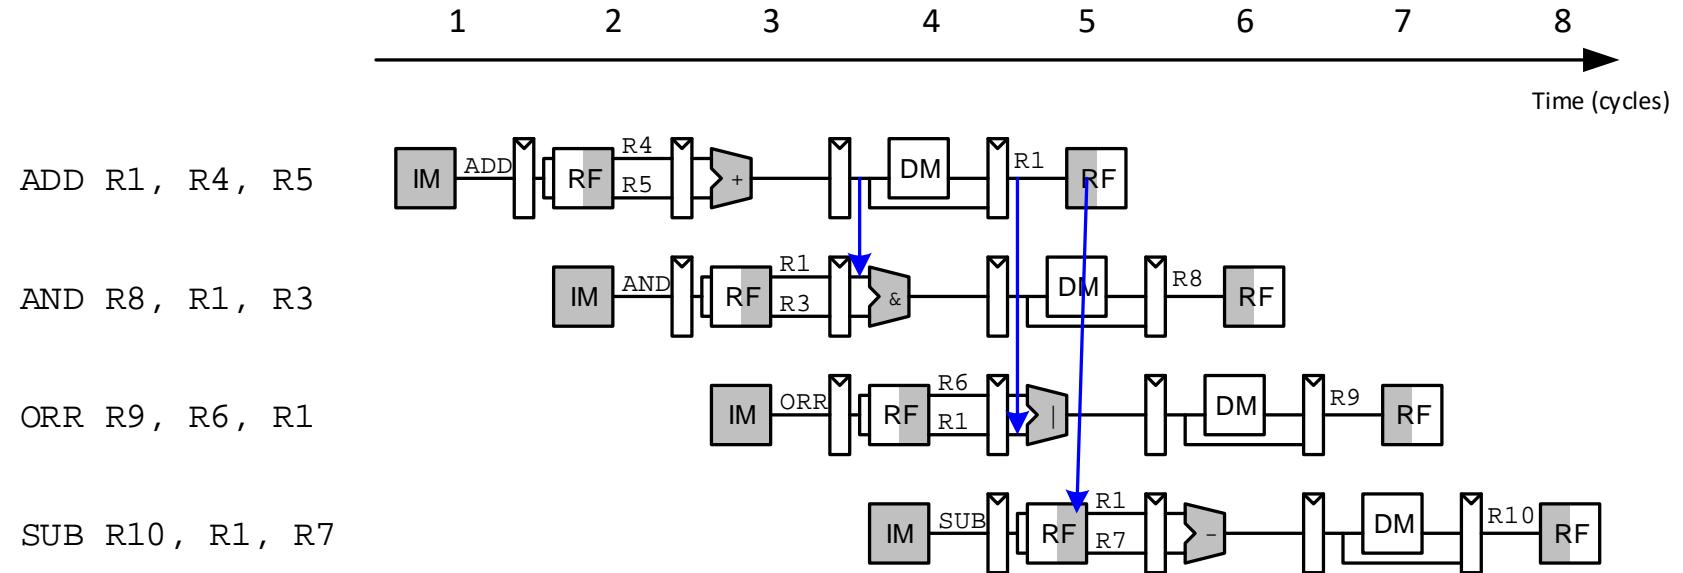


# Compile-Time Hazard Elimination

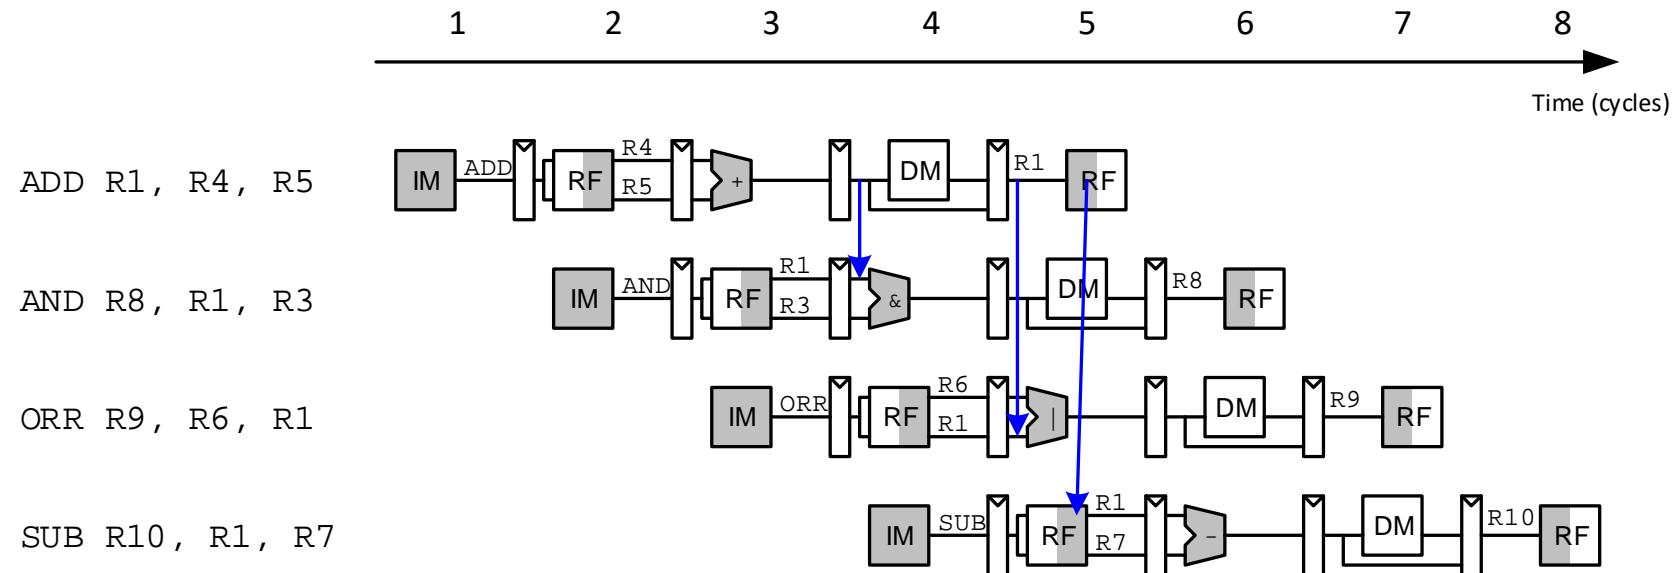
- Insert enough NOPs for result to be ready
- Or move independent useful instructions forward



# Data Forwarding



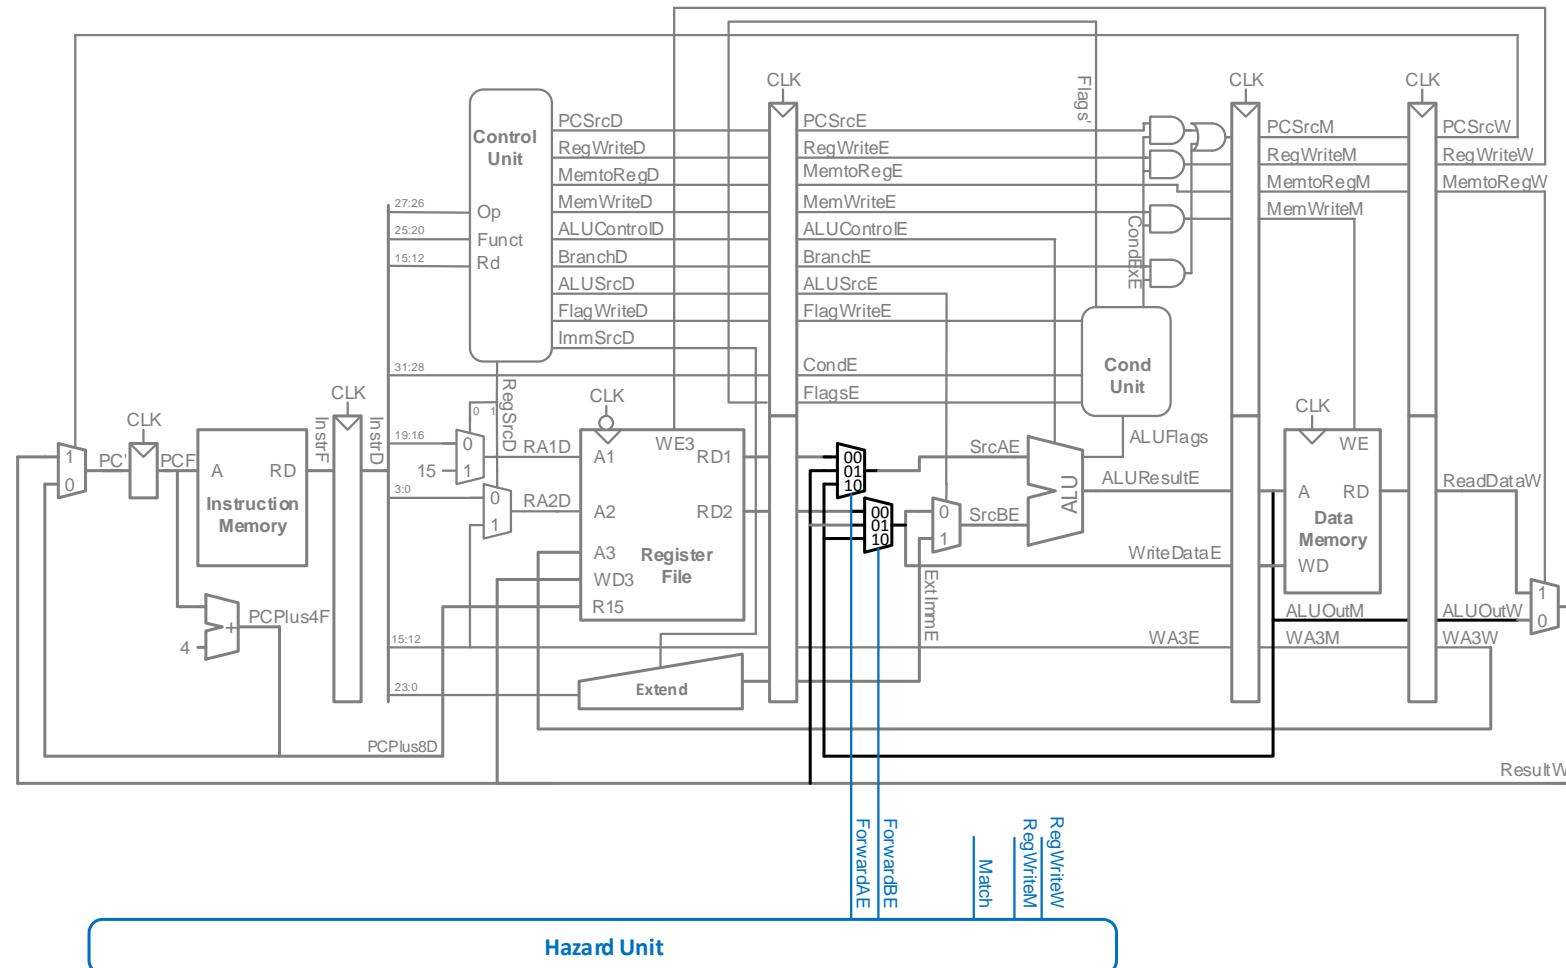
# Data Forwarding



- Check if register read in Execute stage matches register written in Memory or Writeback stage
- If so, forward result



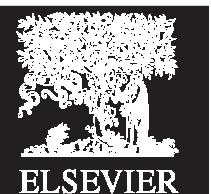
# Data Forwarding



# Data Forwarding

- Execute stage register matches **Memory** stage register?  
Match\_1E\_M = (RA1E == WA3M)  
Match\_2E\_M = (RA2E == WA3M)
- Execute stage register matches **Writeback** stage register?  
Match\_1E\_W = (RA1E == WA3W)  
Match\_2E\_W = (RA2E == WA3W)
- If it matches, forward result:

```
if      (Match_1E_M • RegWriteM)    ForwardAE = 10;  
else if (Match_1E_W • RegWriteW)    ForwardAE = 01;  
else                                ForwardAE = 00;
```



# Data Forwarding

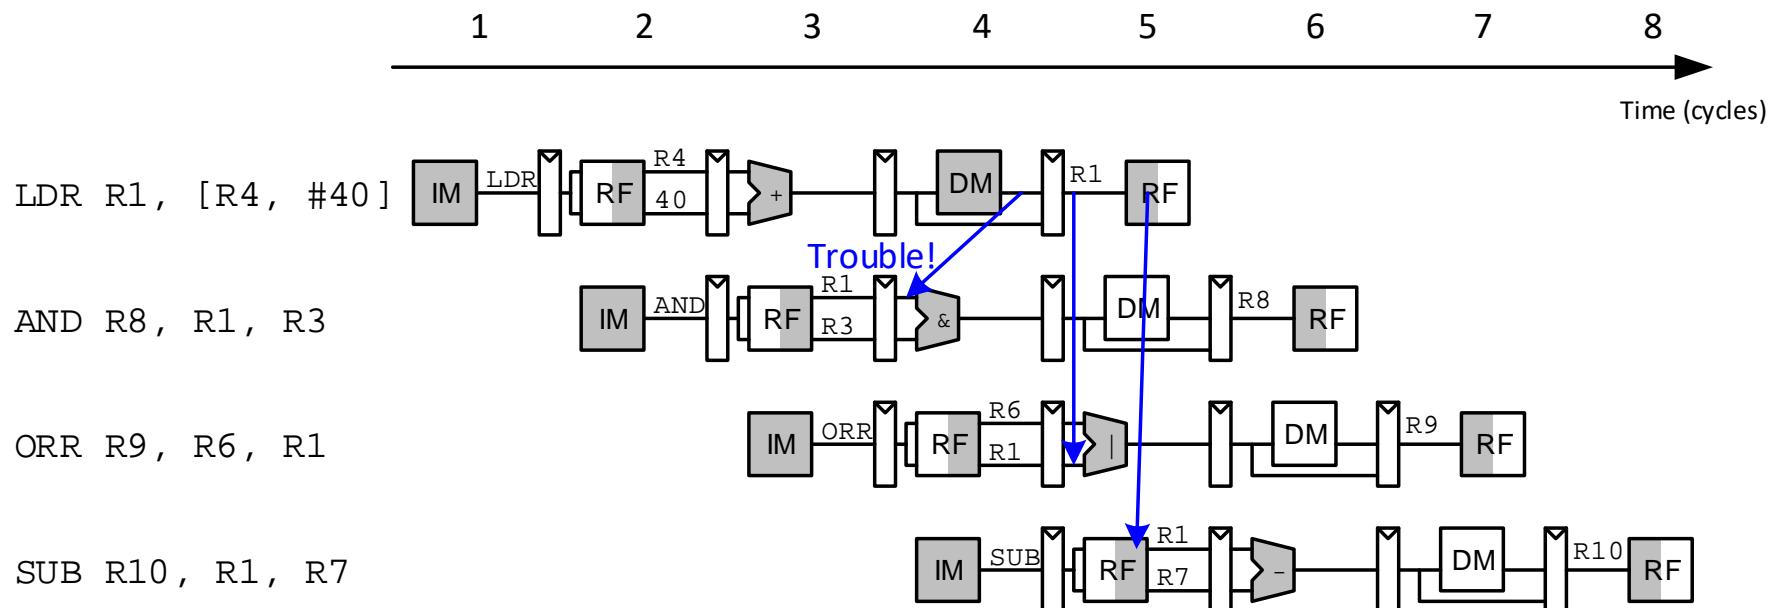
- Execute stage register matches **Memory** stage register?  
Match\_1E\_M = (RA1E == WA3M)  
Match\_2E\_M = (RA2E == WA3M)
- Execute stage register matches **Writeback** stage register?  
Match\_1E\_W = (RA1E == WA3W)  
Match\_2E\_W = (RA2E == WA3W)
- If it matches, forward result:

```
if      (Match_1E_M • RegWriteM)    ForwardAE = 10;  
else if (Match_1E_W • RegWriteW)    ForwardAE = 01;  
else                                ForwardAE = 00;
```

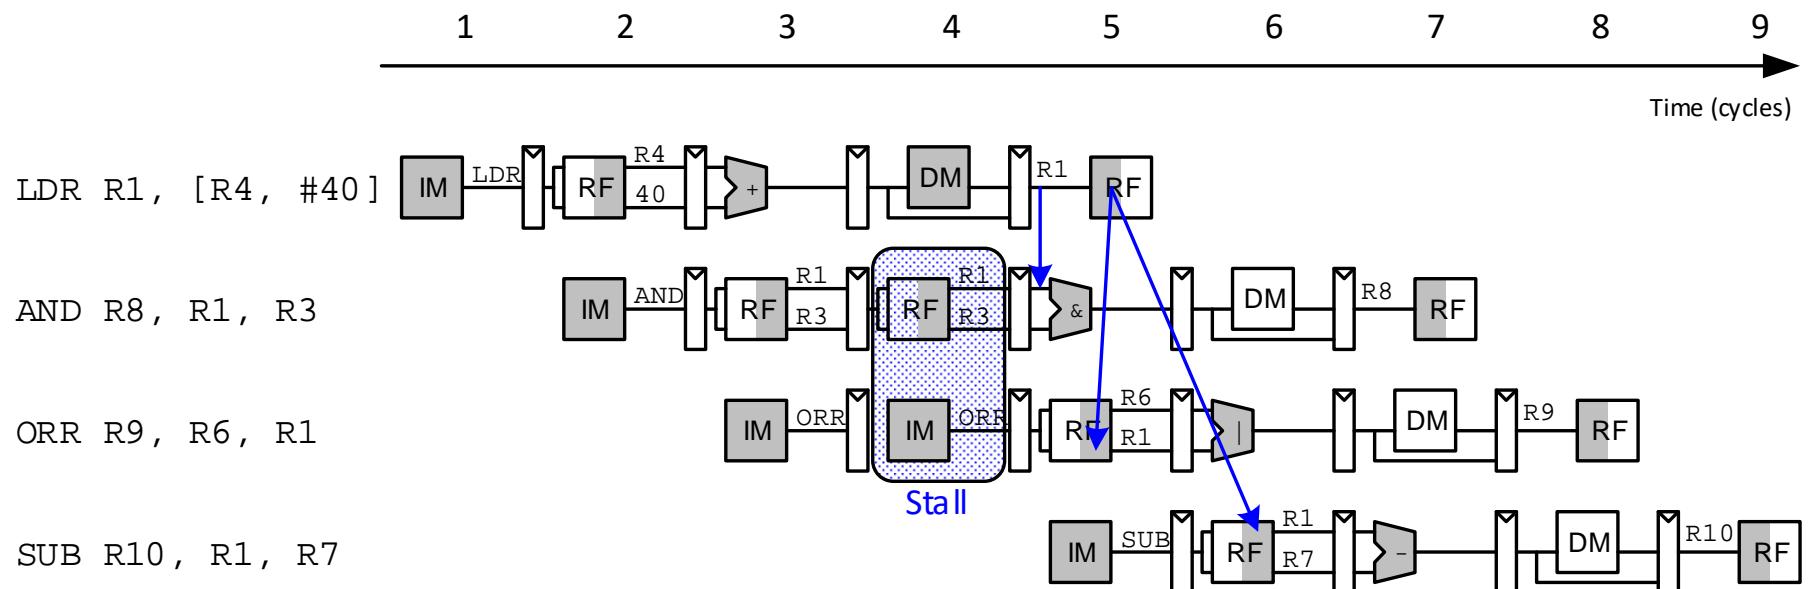
**ForwardBE same but with Match2E**



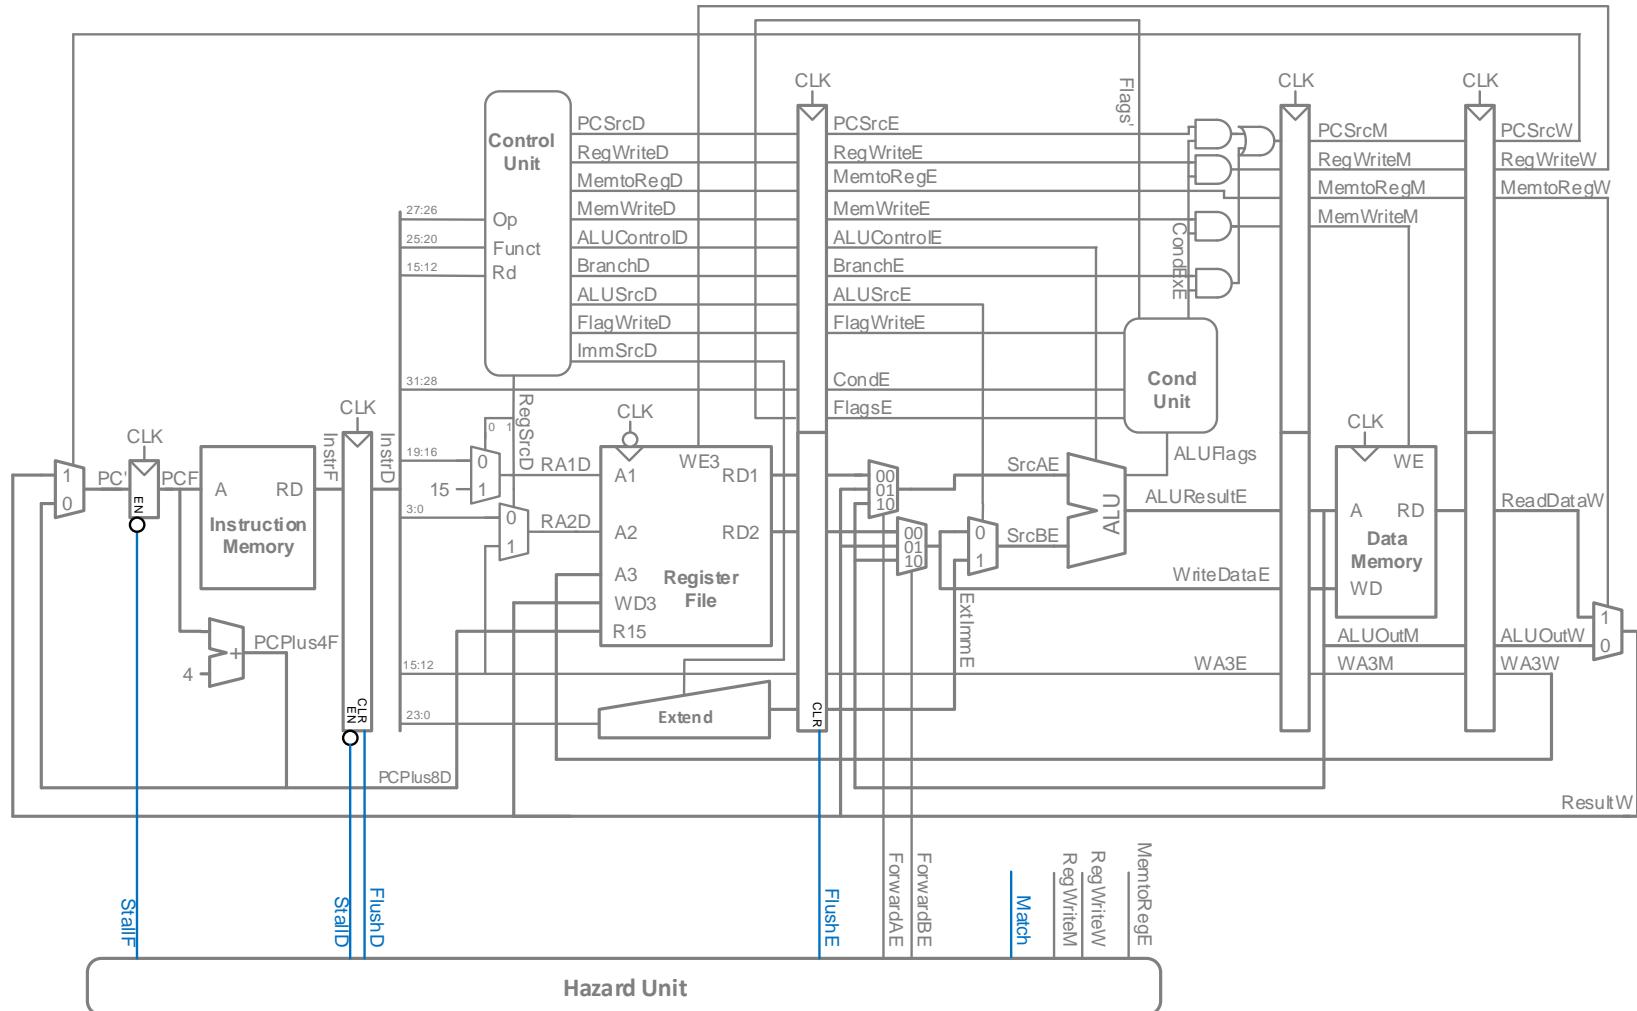
# Stalling



# Stalling



# Stalling Hardware



# Stalling Logic

- Is either source register in the Decode stage the same as the one being written in the Execute stage?

$$Match_{12D\_E} = (RA1D == WA3E) + (RA2D == WA3E)$$

- Is a LDR in the Execute stage AND  $Match_{12D\_E}$ ?

$$ldrstall = Match_{12D\_E} \bullet MemtoRegE$$

$$StallF = StallD = FlushE = ldrstall$$

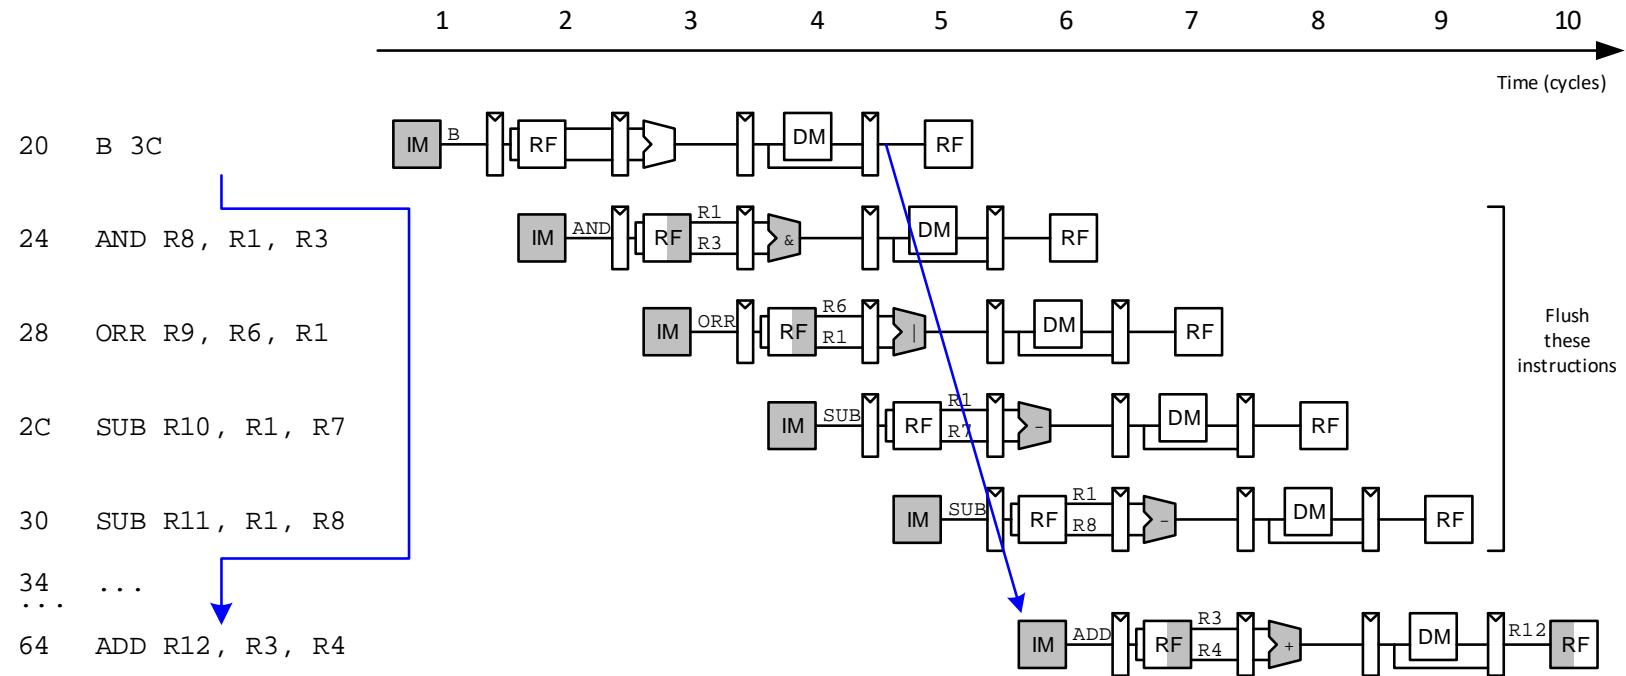


# Control Hazards

- **B:**
  - branch not determined until the Writeback stage of pipeline
  - Instructions after branch fetched before branch occurs
  - These 4 instructions must be flushed if branch happens
- **Writes to PC (R15) similar**



# Control Hazards



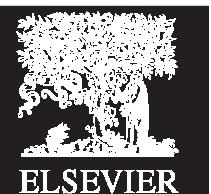
## Branch misprediction penalty

- number of instruction flushed when branch is taken (4)
- May be reduced by determining BTA earlier

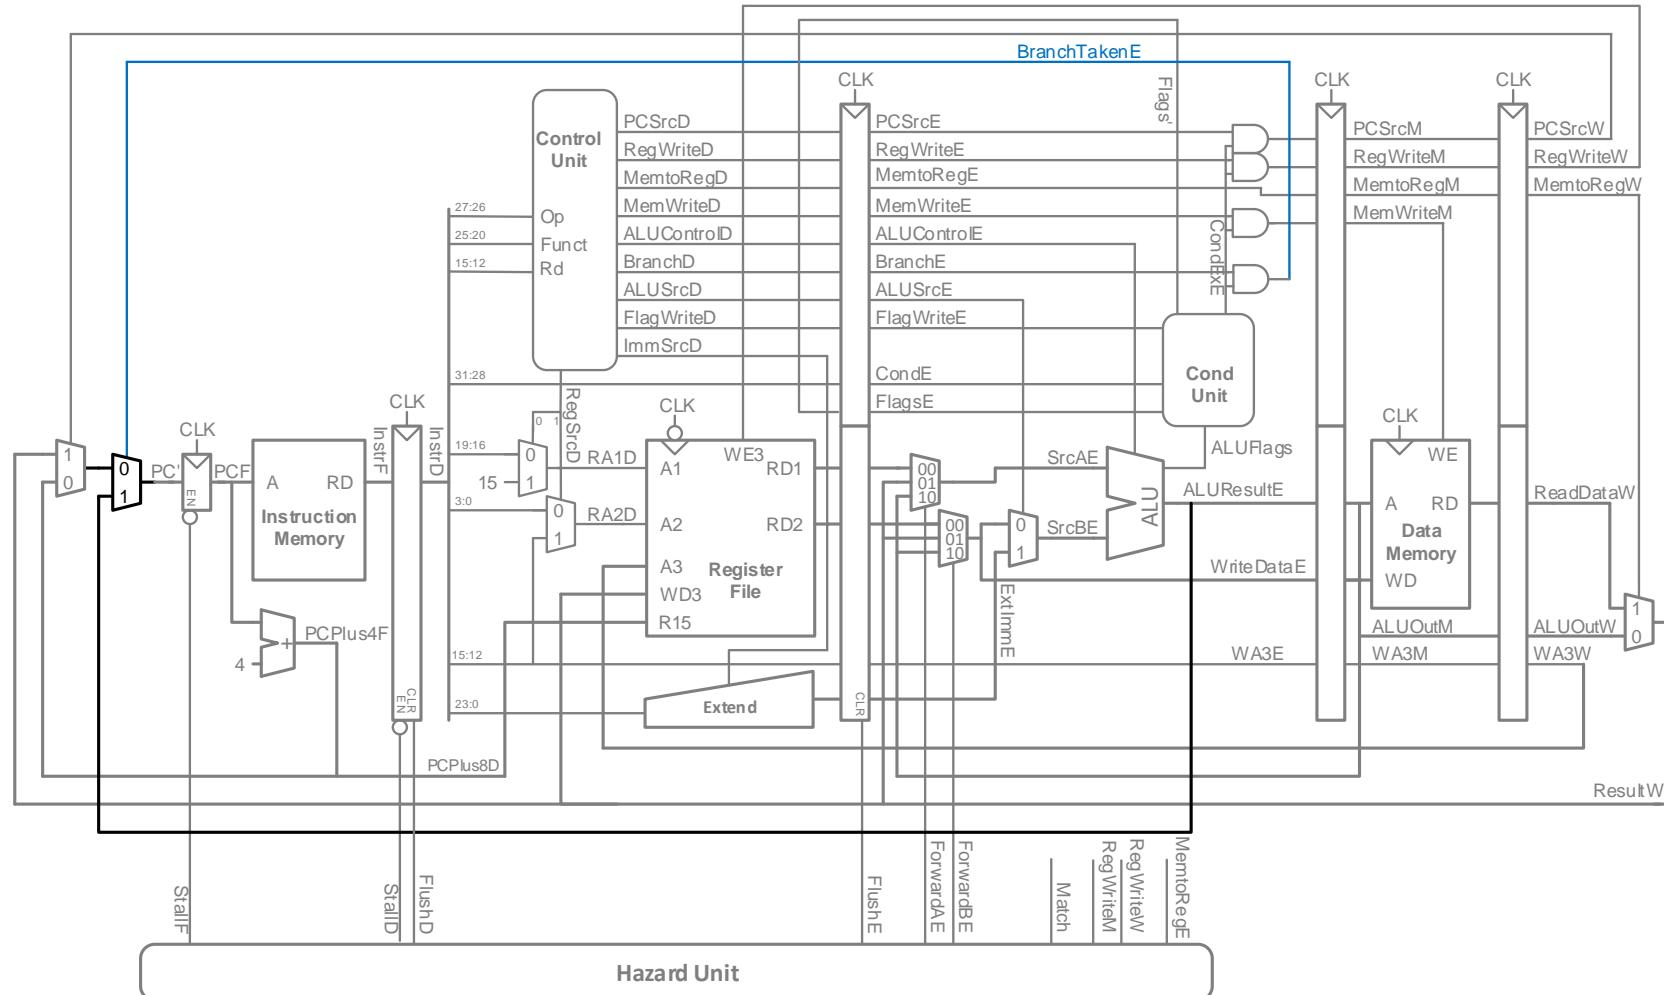


# Early Branch Resolution

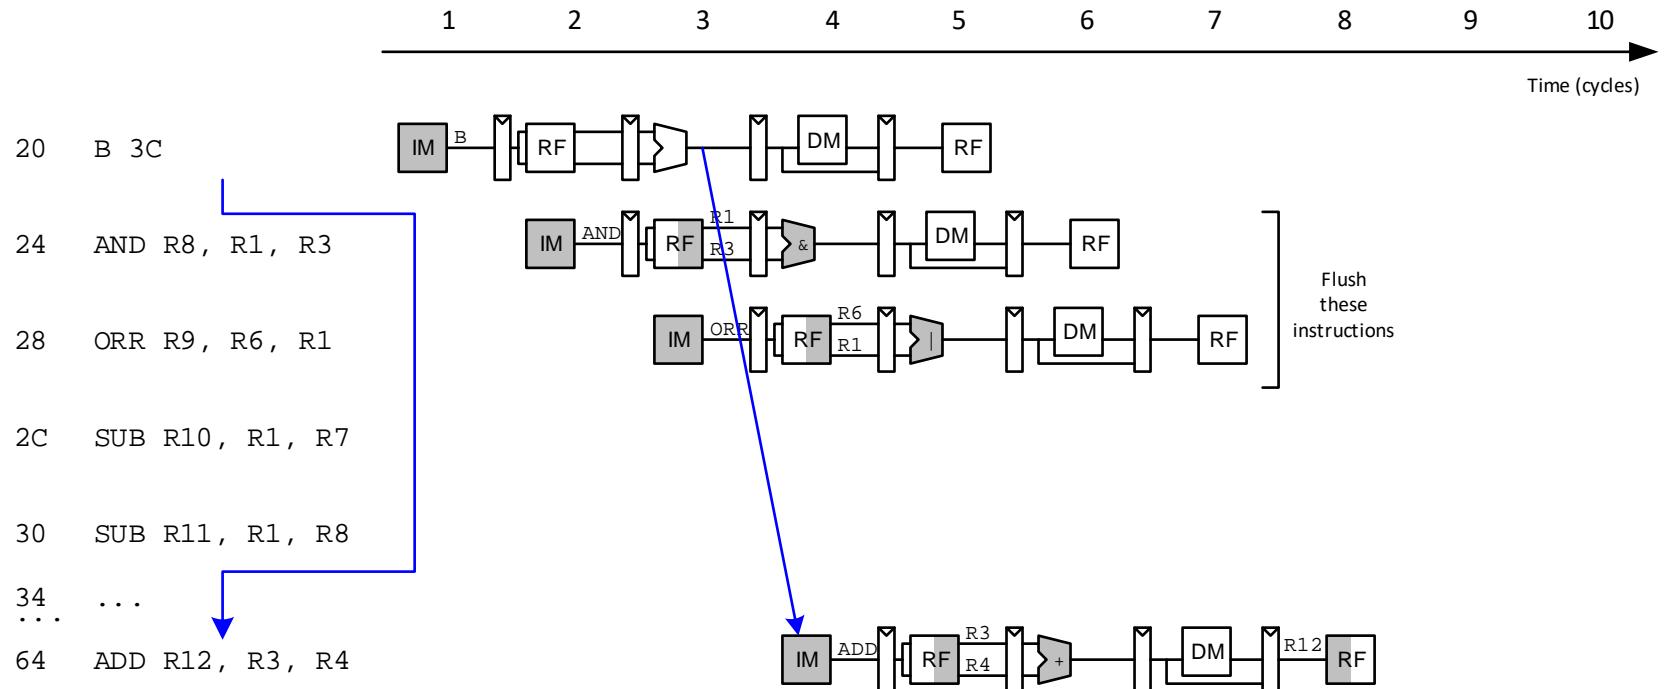
- **Determine BTA in Execute stage**
  - Branch misprediction penalty = 2 cycles
- **Hardware changes**
  - Add a branch multiplexer before *PC* register to select BTA from *ALUResultE*
  - Add *BranchTakenE* select signal for this multiplexer (only asserted if branch condition satisfied)
  - *PCSrcW* now only asserted for writes to *PC*



# Pipelined processor with Early BTA



# Control Hazards with Early BTA



# Control Stalling Logic

- $PCWrPendingF = 1$  if write to  $PC$  in Decode, Execute or Memory

$$PCWrPendingF = PCSrcD + PCSrcE + PCSrcM$$

- **Stall Fetch** if  $PCWrPendingF$

$$StallF = IdrStallD + PCWrPendingF$$

- **Flush Decode** if  $PCWrPendingF$  OR  $PC$  is written in Writeback OR branch is taken

$$FlushD = PCWrPendingF + PCSrcW + BranchTakenE$$

- **Flush Execute** if branch is taken

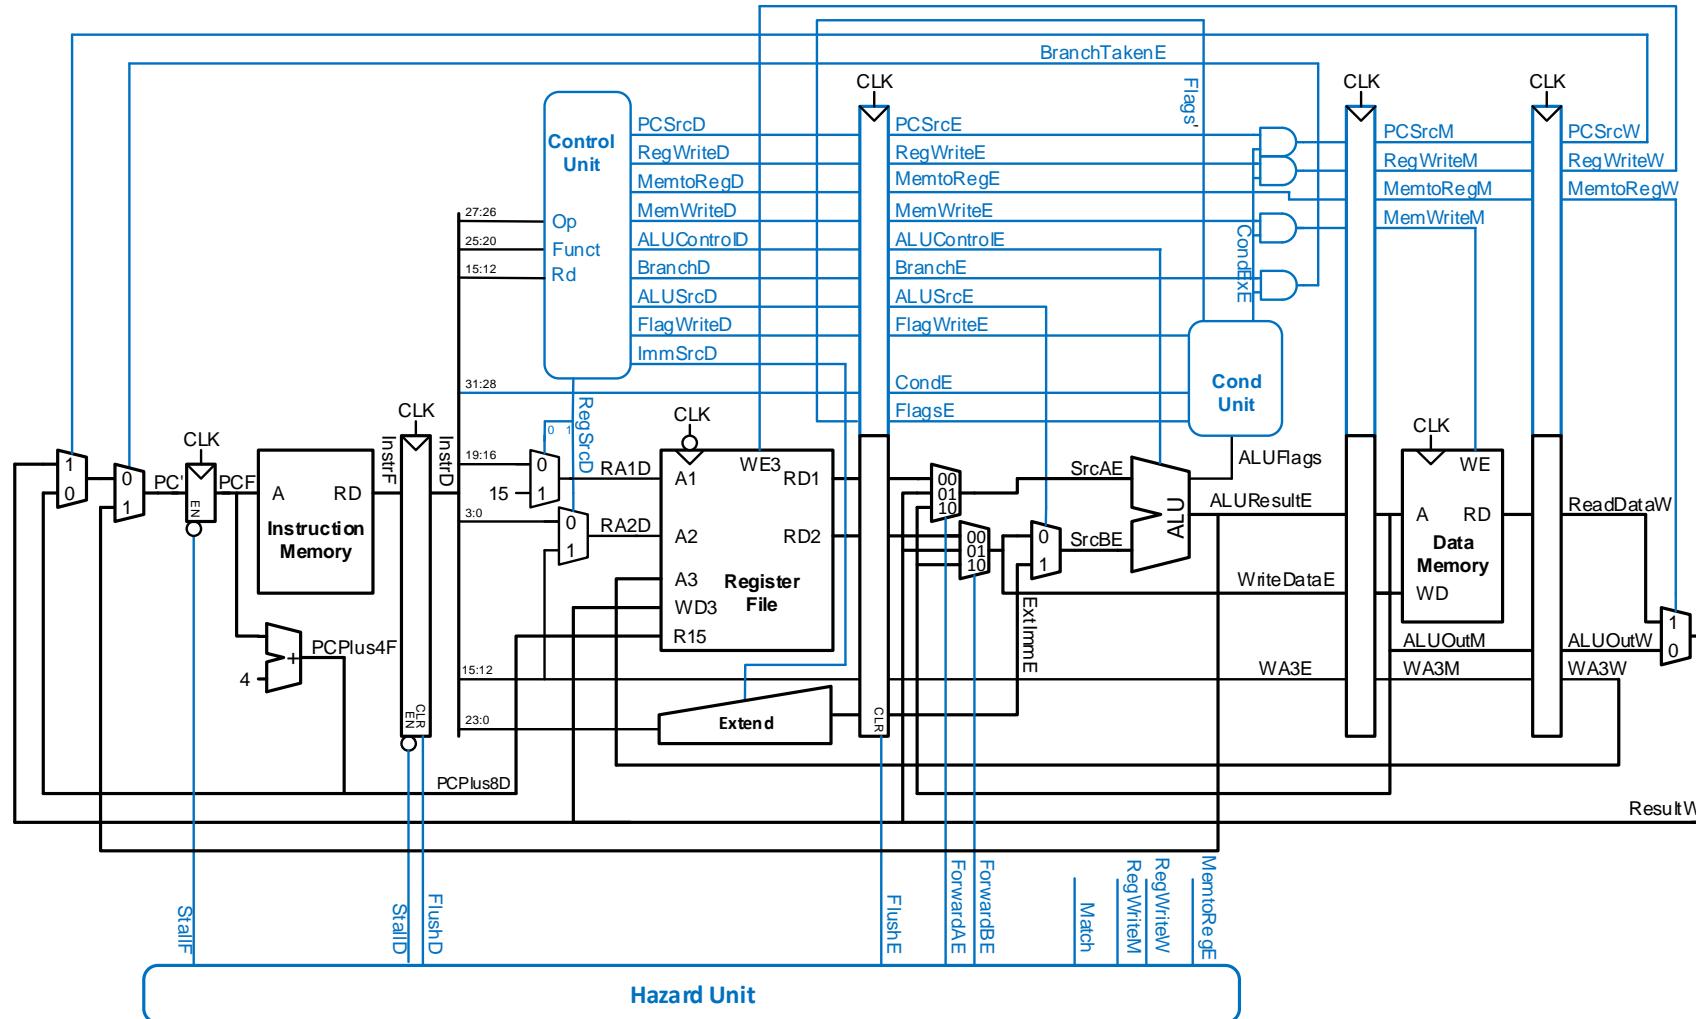
$$FlushE = IdrStallD + BranchTakenE$$

- **Stall Decode** if  $IdrStallD$  (as before)

$$StallD = IdrStallD$$



# ARM Pipelined Processor with Hazard Unit



# Pipelined Performance Example

- **SPECINT2000 benchmark:**
  - 25% loads
  - 10% stores
  - 13% branches
  - 52% R-type
- **Suppose:**
  - 40% of loads used by next instruction
  - 50% of branches mispredicted
- **What is the average CPI?**



# Pipelined Performance Example

- **SPECINT2000 benchmark:**
  - 25% loads
  - 10% stores
  - 13% branches
  - 52% R-type
- **Suppose:**
  - 40% of loads used by next instruction
  - 50% of branches mispredicted
- **What is the average CPI?**
  - Load CPI = 1 when not stalling, 2 when stalling  
So,  $\text{CPI}_{lw} = 1(0.6) + 2(0.4) = 1.4$
  - Branch CPI = 1 when not stalling, 3 when stalling  
So,  $\text{CPI}_{beq} = 1(0.5) + 3(0.5) = 2$

$$\text{Average CPI} = (0.25)(1.4) + (0.1)(1) + (0.13)(2) + (0.52)(1) = 1.23$$



# Pipelined Performance

- Pipelined processor critical path:

$$T_{c3} = \max [$$

$t_{pcq} + t_{\text{mem}} + t_{\text{setup}}$	Fetch
$2(t_{\text{RFread}} + t_{\text{setup}})$	Decode
$t_{pcq} + 2t_{\text{mux}} + t_{\text{ALU}} + t_{\text{setup}}$	Execute
$t_{pcq} + t_{\text{mem}} + t_{\text{setup}}$	Memory
$2(t_{pcq} + t_{\text{mux}} + t_{\text{RFwrite}}) ]$	Writeback



# Pipelined Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	40
Register setup	$t_{\text{setup}}$	50
Multiplexer	$t_{\text{mux}}$	25
ALU	$t_{\text{ALU}}$	120
Memory read	$t_{\text{mem}}$	200
Register file read	$t_{RF\text{read}}$	100
Register file setup	$t_{RF\text{setup}}$	60
Register file write	$t_{RF\text{write}}$	70

**Cycle time:**  $T_{c3} = ?$



# Pipelined Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	40
Register setup	$t_{\text{setup}}$	50
Multiplexer	$t_{\text{mux}}$	25
ALU	$t_{\text{ALU}}$	120
Memory read	$t_{\text{mem}}$	200
Register file read	$t_{RF\text{read}}$	100
Register file setup	$t_{RF\text{setup}}$	60
Register file write	$t_{RF\text{write}}$	70

$$\begin{aligned}\textbf{Cycle time: } T_{c3} &= 2(t_{RF\text{read}} + t_{\text{setup}}) \\ &= 2[100 + 50] \text{ ps} = 300 \text{ ps}\end{aligned}$$



# Pipelined Performance Example

Program with 100 billion instructions

$$\begin{aligned}\text{Execution Time} &= (\# \text{ instructions}) \times \text{CPI} \times T_c \\ &= (100 \times 10^9)(1.23)(300 \times 10^{-12}) \\ &= \mathbf{36.9 \text{ seconds}}\end{aligned}$$



# Processor Performance Comparison

Processor	Execution Time (seconds)	Speedup (single-cycle as baseline)
Single-cycle	84	1
Multicycle	140	0.6
Pipelined	36.9	2.28



# Advanced Microarchitecture

- Deep Pipelining
- Micro-operations
- Branch Prediction
- Superscalar Processors
- Out of Order Processors
- Register Renaming
- SIMD
- Multithreading
- Multiprocessors



# Deep Pipelining

- 10-20 stages typical
- Number of stages limited by:
  - Pipeline hazards
  - Sequencing overhead
  - Power
  - Cost



# Micro-operations

- Decompose more complex instructions into a series of simple instructions called *micro-operations* (*micro-ops* or  $\mu$ -*ops*)
- At run-time, complex instructions are decoded into one or more micro-ops
- Used heavily in CISC (complex instruction set computer) architectures (e.g., x86)
- Used for some ARM instructions, for example:

## Complex Op

LDR R1, [R2], #4

## Micro-op Sequence

LDR R1, [R2]

ADD R2, R2, #4

**Without u-ops, would need 2nd write port on the register file**



